Databricks ML Associate

Questions 301–400 (Fourth 100)

Practice Exam Deck

How to Use This Deck

100 new questions (Q301–Q400), none repeated from Q1–Q300.

Each question slide has four options (A–D).

The next slide gives the correct answer and explains why the right option is right and why each wrong option is wrong.

This set emphasizes areas not yet covered: CLI/REST API, cluster config, registry specifics, model flavors, DLT, profiling, time series, and governance.

Question 301

Which tool lets you automate Databricks operations from a terminal?

A. Databricks CLI

B. Databricks SQL Editor

C. Notebook widgets

D. MLflow UI

Answer 301

Correct: A — the Databricks CLI automates operations.

Wrong:

B is SQL editing.

C is notebook UI.

D is experiment tracking.

Question 302

Which REST API endpoint lists clusters?

A. `GET /api/2.0/clusters/list`

B. `GET /api/2.0/clusters/all`

C. `GET /api/2.0/clusters/describe`

D. `GET /api/2.0/clusters`

Answer 302

Correct: A — `/clusters/list` returns clusters.

Wrong:

B, C, D are not the correct endpoint.

Question 303

Which authentication method is recommended for the Databricks CLI?

A. Personal access token or OAuth

B. Username and password only

C. SSH key only

D. Anonymous access

Answer 303

Correct: A — tokens or OAuth are recommended.

Wrong:

B, C, D are not recommended.

Question 304

What is a personal access token used for?

A. Authenticating API and CLI requests
B. Training models
C. Storing features
D. Registering models

Answer 304

Correct: A — PATs authenticate requests.

Wrong:

B, C, D are unrelated.

Question 305

Which component hosts the MLflow tracking backend in Databricks?

A. The Databricks-managed tracking server

B. A local SQLite database only

C. A notebook

D. A cluster driver only

Answer 305

Correct: A — Databricks hosts a managed tracking server.

Wrong:

B, C, D are not the managed backend.

Question 306

What is a cluster access mode?

A. It controls user isolation and supported workloads
B. It sets the runtime version
C. It sets the number of workers
D. It sets the storage location

Answer 306

Correct: A — access modes define isolation and features.

Wrong:

B, C, D are other settings.

Question 307

Which access mode is required for Unity Catalog-enabled clusters?

A. Shared or single-user with UC support

B. No isolation

C. Legacy only

D. Personal only

Answer 307

Correct: A — UC requires supported access modes.

Wrong:

B, C, D are not the UC-required modes.

Question 308

What does autoscaling do?

A. Adjusts the number of workers based on load
B. Changes the runtime version
C. Changes the access mode
D. Changes the storage

Answer 308

Correct: A — autoscaling adds/removes workers.

Wrong:

B, C, D are unrelated.

Question 309

What is a cluster policy?

A. A rule set that constrains cluster configuration

B. A model version

C. A feature table

D. A notebook

Answer 309

Correct: A — policies constrain cluster settings.

Wrong:

B, C, D are unrelated.

Question 310

What is a single-node cluster best suited for?

A. Small-scale development and testing

B. Large-scale distributed training

C. SQL BI only

D. Streaming only

Answer 310

Correct: A — single-node is for small workloads.

Wrong:

B needs workers.

C is SQL warehouses.

D is not its main use.

Question 311

What is a driver node responsible for in a cluster?

A. Coordinating work and running the Spark driver
B. Storing all data
C. Serving REST requests only
D. Running SQL only

Answer 311

Correct: A — the driver coordinates execution.

Wrong:

B, C, D are not the driver's role.

Question 312

What does a worker node do?

A. Executes tasks assigned by the driver

B. Hosts the Spark driver

C. Serves REST endpoints

D. Stores the model registry

Answer 312

Correct: A — workers execute tasks.

Wrong:

B is the driver.

C, D are other services.

Question 313

Which runtime version should you select to use Delta Live Tables?

A. A DLT-compatible runtime version
B. Any standard runtime
C. A single-node runtime only
D. A SQL warehouse

Answer 313

Correct: A — DLT requires a compatible runtime.

Wrong:

B, C, D are not DLT runtimes.

Question 314

What is a Delta Live Tables pipeline?

A. A declarative pipeline for batch and streaming data

B. A model

C. A notebook

D. A cluster

Answer 314

Correct: A — DLT builds declarative pipelines.

Wrong:

B, C, D are not pipelines.

Question 315

What does DLT's `@dlt.table` decorator do?

A. Defines a table in a DLT pipeline

B. Trains a model

C. Registers a model

D. Deploys an endpoint

Answer 315

Correct: A — it defines a DLT table.

Wrong:

B, C, D are unrelated.

Question 316

What does DLT's `@dlt.view` decorator do?

A. Defines a temporary view in a DLT pipeline

B. Defines a table

C. Defines a model

D. Defines a cluster

Answer 316

Correct: A — it defines a DLT view.

Wrong:

B, C, D are false.

Question 317

What is a data profiling notebook used for?

A. Summarizing data quality and distributions
B. Training models
C. Deploying models
D. Registering models

Answer 317

Correct: A — profiling summarizes data.

Wrong:

B, C, D are unrelated.

Question 318

Which chart best shows the relationship between two numeric variables?

A. Scatter plot
B. Histogram
C. Bar chart
D. Pie chart

Answer 318

Correct: A — scatter plots show two-variable relationships.

Wrong:

B is one variable's distribution.

C is categorical.

D is part-to-whole.

Question 319

Which chart shows the distribution of one numeric variable?

A. Histogram
B. Scatter plot
C. Line chart
D. Heatmap

Answer 319

Correct: A — histograms show distributions.

Wrong:

B is relationships.

C is trends.

D is a matrix view.

Question 320

Which chart is best for trends over time?

A. Line chart
B. Scatter plot
C. Pie chart
D. Bar chart

Answer 320

Correct: A — line charts show trends.

Wrong:

B, C, D are less suited to trends.

Question 321

What is a correlation matrix used for?

A. Showing pairwise correlations between numeric features
B. Training models
C. Serving models
D. Registering models

Answer 321

Correct: A — correlation matrices show pairwise relationships.

Wrong:

B, C, D are unrelated.

Question 322

What is a boxplot useful for?

A. Showing quartiles and outliers
B. Showing correlations
C. Showing trends
D. Showing class balance

Answer 322

Correct: A — boxplots show quartiles and outliers.

Wrong:

B, C, D are other visualizations.

Question 323

Why is a holdout test set important after tuning?

A. It gives an unbiased final estimate

B. It tunes hyperparameters

C. It stores models

D. It registers models

Answer 323

Correct: A — the test set gives a final estimate.

Wrong:

B tuning uses validation data.

C, D are unrelated.

Question 324

What is the risk of tuning on the test set?

A. The test set is no longer an unbiased estimate
B. Training becomes slower
C. Accuracy always drops
D. The model cannot be saved

Answer 324

Correct: A — tuning on test data leaks it.

Wrong:

B, C, D are not the main risk.

Question 325

What is nested cross-validation used for?

A. Tuning and estimating performance without optimistic bias
B. Training one model
C. Deploying a model
D. Registering a model

Answer 325

Correct: A — nested CV separates tuning from evaluation.

Wrong:

B, C, D are unrelated.

Question 326

Which MLflow model flavor supports scikit-learn models?

A. `mlflow.sklearn`

B. `mlflow.tensorflow`

C. `mlflow.pytorch`

D. `mlflow.onnx`

Answer 326

Correct: A — `sklearn` is the scikit-learn flavor.

Wrong:

B, C, D are other flavors.

Question 327

Which MLflow model flavor supports Spark ML models?

A. `mlflow.spark`

B. `mlflow.sklearn`

C. `mlflow.pytorch`

D. `mlflow.tensorflow`

Answer 327

Correct: A — `spark` is the Spark ML flavor.

Wrong:

B, C, D are other flavors.

Question 328

Why is a custom `pyfunc` model useful?

A. It wraps arbitrary prediction logic behind a standard interface

B. It trains faster

C. It stores data

D. It registers models

Answer 328

Correct: A — pyfunc standardizes custom logic.

Wrong:

B, C, D are false.

Question 329

What method must a custom pyfunc model implement?

A. `predict()`

B. `fit()`

C. `transform()`

D. `evaluate()`

Answer 329

Correct: A — pyfunc models implement `predict()`.

Wrong:

B, C, D are not required for prediction.

Question 330

What does `mlflow.pyfunc.log_model()` do?

A. Logs a custom pyfunc model as an MLflow artifact

B. Registers a model

C. Trains a model

D. Deploys a model

Answer 330

Correct: A — it logs a pyfunc model.

Wrong:

B, C, D are other operations.

Question 331

What does model serving require before an endpoint can be created?

A. A registered model version

B. A feature table

C. A notebook

D. A cluster

Answer 331

Correct: A — serving needs a registered version.

Wrong:

B, C, D are not required.

Question 332

What is a serving endpoint's scale-to-zero behavior?

A. It shuts down idle capacity to save cost

B. It deletes the model

C. It retrains the model

D. It changes the schema

Answer 332

Correct: A — scale-to-zero saves cost when idle.

Wrong:

B, C, D are false.

Question 333

What is a feature specification used for in online serving?

A. It tells the endpoint which online table to query for features
B. It trains the model
C. It registers the model
D. It deploys the model

Answer 333

Correct: A — feature specs connect endpoints to online tables.

Wrong:

B, C, D are unrelated.

Question 334

What is online feature serving latency primarily affected by?

A. Online store query latency and model inference time

B. Training time

C. Number of features in the training set only

D. Cluster size

Answer 334

Correct: A — latency depends on online lookup and inference.

Wrong:

B, C, D are not the primary factors.

Question 335

What is a batch prediction table?

A. A table storing offline model predictions

B. A feature table

C. A model

D. A notebook

Answer 335

Correct: A — it stores offline predictions.

Wrong:

B, C, D are unrelated.

Question 336

What is the purpose of a feature freshness policy?

A. To ensure features are updated within an acceptable window

B. To train models

C. To register models

D. To deploy models

Answer 336

Correct: A — freshness policies bound feature staleness.

Wrong:

B, C, D are unrelated.

Question 337

What is the main difference between a view and a table in Unity Catalog?

A. A view is a saved query; a table stores data

B. A view stores data; a table is a query

C. They are identical

D. A view cannot be queried

Answer 337

Correct: A — views are saved queries.

Wrong:

B is reversed.

C, D are false.

Question 338

What is a managed table in Unity Catalog?

A. A table whose data lifecycle is managed by UC

B. A table stored outside UC

C. A view

D. A model

Answer 338

Correct: A — managed tables have UC-managed storage.

Wrong:

B is external.

C, D are other objects.

Question 339

What is an external table in Unity Catalog?

A. A table whose data lives at a user-specified storage path

B. A table managed by UC

C. A view

D. A model

Answer 339

Correct: A — external tables use user storage.

Wrong:

B is managed.

C, D are other objects.

Question 340

Which Unity Catalog privilege allows executing a model?

A. `EXECUTE`

B. `SELECT`

C. `MODIFY`

D. `READ`

Answer 340

Correct: A — `EXECUTE` allows model use.

Wrong:

B is table read.

C is table write.

D is not a UC privilege.

Question 341

What does a model version's description field store?

A. Human-readable notes about the version

B. The training data

C. The model weights

D. The cluster config

Answer 341

Correct: A — descriptions store notes.

Wrong:

B, C, D are stored elsewhere.

Question 342

What does a model version's source run link provide?

A. Traceability back to the originating MLflow run

B. The cluster config

C. The serving endpoint

D. The feature table

Answer 342

Correct: A — source run links provide traceability.

Wrong:

B, C, D are unrelated.

Question 343

What is a model version's stage used for in the workspace registry?

A. Indicating lifecycle state such as Staging or Production

B. Storing metrics

C. Training the model

D. Storing data

Answer 343

Correct: A — stages indicate lifecycle state.

Wrong:

B, C, D are unrelated.

Question 344

Why do Unity Catalog model registries favor aliases over stages?

A. Aliases are more flexible and support multiple deployment references

B. Stages are deprecated in all tools

C. Aliases store data

D. Aliases train models

Answer 344

Correct: A — aliases are flexible deployment references.

Wrong:

B, C, D are false.

Question 345

What is a model version's archive action used for?

A. Retiring a version from active use

B. Training the model

C. Registering a new version

D. Deploying the model

Answer 345

Correct: A — archiving retires versions.

Wrong:

B, C, D are other actions.

Question 346

What is a model version's comment used for?

A. Collaborators leave notes about the version
B. Storing data
C. Training the model
D. Serving the model

Answer 346

Correct: A — comments store collaborator notes.

Wrong:

B, C, D are unrelated.

Question 347

What is a MLflow model version's tag used for?

A. Attaching key-value metadata
B. Storing metrics
C. Storing features
D. Storing the model weights

Answer 347

Correct: A — tags attach metadata.

Wrong:

B, C, D are stored elsewhere.

Question 348

Which MLflow API sets a model version tag?

A. `mlflow.set_model_version_tag()`

B. `mlflow.log_tag()`

C. `mlflow.set_tag()`

D. `mlflow.tag_model()`

Answer 348

Correct: A — it sets a model version tag.

Wrong:

B, C set run tags.

D is not the API.

Question 349

What is a model version's alias change recorded as?

A. A registry activity event

B. A training log

C. A feature table

D. A notebook

Answer 349

Correct: A — alias changes are registry activity.

Wrong:

B, C, D are unrelated.

Question 350

What is a Databricks notebook's default language?

A. Python
B. SQL
C. Scala
D. R

Answer 350

Correct: A — Python is the default.

Wrong:

B, C, D are other supported languages.

Question 351

What is a magic command in a Databricks notebook?

A. A command that changes the notebook's behavior or language

B. A model

C. A table

D. A cluster

Answer 351

Correct: A — magics alter notebook behavior.

Wrong:

B, C, D are unrelated.

Question 352

Which magic command switches a cell to SQL?

A. `%sql`

B. `%python`

C. `%scala`

D. `%md`

Answer 352

Correct: A — `%sql` switches to SQL.

Wrong:

B, C, D switch to other languages or Markdown.

Question 353

Which magic command switches a cell to Markdown?

A. `%md`

B. `%sql`

C. `%python`

D. `%sh`

Answer 353

Correct: A — `%md` switches to Markdown.

Wrong:

B, C, D switch to other languages or shell.

Question 354

Which magic command runs shell commands?

A. `%sh`

B. `%sql`

C. `%python`

D. `%md`

Answer 354

Correct: A — `%sh` runs shell commands.

Wrong:

B, C, D are other magics.

Question 355

What does `dbutils.fs.ls()` do?

A. Lists files in a filesystem path

B. Lists clusters

C. Lists models

D. Lists tables

Answer 355

Correct: A — it lists files.

Wrong:

B, C, D are other utilities.

Question 356

What does `dbutils.notebook.run()` do?

A. Runs another notebook and returns its exit value

B. Trains a model

C. Registers a model

D. Lists files

Answer 356

Correct: A — it runs a notebook and returns its value.

Wrong:

B, C, D are unrelated.

Question 357

What does `dbutils.notebook.exit()` do?

A. Exits a notebook with a value

B. Exits a cluster

C. Exits a model

D. Exits a table

Answer 357

Correct: A — it exits a notebook with a value.

Wrong:

B, C, D are unrelated.

Question 358

What is a notebook task in a Databricks Job?

A. A single notebook executed as part of the job

B. A cluster

C. A model

D. A table

Answer 358

Correct: A — a task runs a notebook.

Wrong:

B, C, D are unrelated.

Question 359

What is a job dependency?

A. A task that runs after another task succeeds

B. A cluster

C. A model

D. A table

Answer 359

Correct: A — dependencies order tasks.

Wrong:

B, C, D are unrelated.

Question 360

What is a job trigger?

A. A schedule or event that starts a job

B. A model

C. A table

D. A cluster

Answer 360

Correct: A — triggers start jobs.

Wrong:

B, C, D are unrelated.

Question 361

What does a job retry policy control?

A. How many times and how often failed tasks are retried

B. The model type

C. The feature table

D. The cluster size

Answer 361

Correct: A — retry policies control retries.

Wrong:

B, C, D are unrelated.

Question 362

What is a cluster autoscaling minimum?

A. The smallest number of workers a cluster can scale to
B. The runtime version
C. The driver size
D. The storage size

Answer 362

Correct: A — minimum workers bound autoscaling.

Wrong:

B, C, D are other settings.

Question 363

What is a cluster autoscaling maximum?

A. The largest number of workers a cluster can scale to
B. The runtime version
C. The driver size
D. The storage size

Answer 363

Correct: A — maximum workers bound autoscaling.

Wrong:

B, C, D are other settings.

Question 364

What is a spot instance in Databricks?

A. A lower-cost, interruptible cloud instance

B. A type of model

C. A feature table

D. A cluster policy

Answer 364

Correct: A — spot instances are cheaper but interruptible.

Wrong:

B, C, D are unrelated.

Question 365

Why might spot instances be risky for long training jobs?

A. They can be reclaimed, interrupting the job

B. They are slower

C. They cannot run Python

D. They cannot use GPUs

Answer 365

Correct: A — spot instances can be interrupted.

Wrong:

B, C, D are false.

Question 366

What is a GPU cluster used for?

A. Deep learning and other compute-intensive workloads
B. SQL BI only
C. File storage
D. Model serving only

Answer 366

Correct: A — GPUs accelerate heavy workloads.

Wrong:

B, C, D are not their main use.

Question 367

What is the purpose of a cluster init script?

A. To run custom setup when a cluster starts
B. To train a model
C. To register a model
D. To deploy a model

Answer 367

Correct: A — init scripts customize cluster startup.

Wrong:

B, C, D are unrelated.

Question 368

What is the risk of long-running init scripts?

A. They slow cluster startup and can cause timeouts
B. They improve accuracy
C. They reduce storage
D. They register models

Answer 368

Correct: A — slow scripts delay startup.

Wrong:

B, C, D are false.

Question 369

What is the purpose of environment variables on a cluster?

A. To configure runtime behavior and credentials
B. To train models
C. To register models
D. To store tables

Answer 369

Correct: A — env vars configure runtime behavior.

Wrong:

B, C, D are unrelated.

Question 370

What does a notebook-scoped library installation do?

A. Installs a library only for the current notebook session
B. Installs a library cluster-wide
C. Installs a model
D. Installs a table

Answer 370

Correct: A — notebook-scoped libraries are session-only.

Wrong:

B is cluster-scoped.

C, D are false.

Question 371

What does a cluster-scoped library installation do?

A. Installs a library for all notebooks on the cluster
B. Installs a library for one notebook only
C. Installs a model
D. Installs a table

Answer 371

Correct: A — cluster-scoped libraries apply cluster-wide.

Wrong:

B is notebook-scoped.

C, D are false.

Question 372

What is a time-series forecasting problem?

A. Predicting future values from ordered historical data
B. Clustering customers
C. Classifying text
D. Reducing dimensions

Answer 372

Correct: A — forecasting predicts future values.

Wrong:

B, C, D are other tasks.

Question 373

What is the role of a lag feature in time-series forecasting?

A. It provides past values as predictors
B. It encodes categories
C. It reduces dimensions
D. It scales features

Answer 373

Correct: A — lag features use past values.

Wrong:

B, C, D are unrelated.

Question 374

What is a rolling window feature?

A. An aggregate over a moving window of past data

B. A category encoding

C. A dimensionality reduction

D. A scaling method

Answer 374

Correct: A — rolling windows aggregate past data.

Wrong:

B, C, D are unrelated.

Question 375

What is the risk of using future data in time-series features?

A. Data leakage and unrealistic performance
B. Faster training
C. Higher accuracy
D. Smaller models

Answer 375

Correct: A — future data causes leakage.

Wrong:

B, C, D are false.

Question 376

What is a recommendation system?

A. A model that suggests items to users
B. A clustering model
C. A forecasting model
D. A dimensionality reduction model

Answer 376

Correct: A — recommenders suggest items.

Wrong:

B, C, D are other tasks.

Question 377

Which Spark ML algorithm is used for collaborative filtering?

A. ALS
B. KMeans
C. LinearRegression
D. RandomForest

Answer 377

Correct: A — ALS is for collaborative filtering.

Wrong:

B, C, D are other algorithms.

Question 378

What does a graph analytics library help with?

A. Analyzing relationships between entities

B. Training image models

C. Encoding text

D. Scaling features

Answer 378

Correct: A — graph analytics studies relationships.

Wrong:

B, C, D are unrelated.

Question 379

What is a graph in graph analytics?

A. A set of nodes and edges

B. A plot

C. A table

D. A model

Answer 379

Correct: A — graphs have nodes and edges.

Wrong:

B, C, D are not graphs.

Question 380

What is a unit test in ML code?

A. A test of a small piece of logic

B. A full training run

C. A model

D. A cluster

Answer 380

Correct: A — unit tests check small logic pieces.

Wrong:

B, C, D are not unit tests.

Question 381

What is a data validation test in an ML pipeline?

A. A check that data meets expected schema and quality rules

B. A model evaluation

C. A cluster test

D. A serving test

Answer 381

Correct: A — data validation enforces schema and quality.

Wrong:

B, C, D are other tests.

Question 382

What is a model validation test?

A. A check that model performance meets thresholds

B. A data schema check

C. A cluster check

D. A serving check

Answer 382

Correct: A — model validation checks performance.

Wrong:

B, C, D are other checks.

Question 383

Why is CI/CD important for ML?

A. It automates testing and deployment of code and models

B. It trains faster

C. It stores data

D. It reduces feature count

Answer 383

Correct: A — CI/CD automates testing and deployment.

Wrong:

B, C, D are false.

Question 384

What is a model deployment pipeline?

A. An automated sequence that validates and deploys models

B. A training notebook

C. A feature table

D. A cluster

Answer 384

Correct: A — deployment pipelines automate release.

Wrong:

B, C, D are unrelated.

Question 385

What is a rollback in model deployment?

A. Reverting to a previous model version

B. Training a new model

C. Registering a model

D. Deleting data

Answer 385

Correct: A — rollback reverts to an earlier version.

Wrong:

B, C, D are unrelated.

Question 386

What is a model alias in Unity Catalog used for during deployment?

A. It points to the version served by an endpoint

B. It stores data

C. It trains the model

D. It deletes the model

Answer 386

Correct: A — aliases point endpoints to versions.

Wrong:

B, C, D are false.

Question 387

What is A/B testing in model deployment?

A. Comparing performance of two model versions in production

B. Training two models

C. Deleting two models

D. Registering two models

Answer 387

Correct: A — A/B testing compares deployed versions.

Wrong:

B, C, D are not A/B testing.

Question 388

How is traffic splitting used for A/B testing?

A. It routes a fraction of requests to each model version

B. It trains models

C. It deletes models

D. It registers models

Answer 388

Correct: A — traffic splitting routes requests.

Wrong:

B, C, D are false.

Question 389

What is a model's inference latency?

A. The time to produce a prediction

B. The training time

C. The storage size

D. The feature count

Answer 389

Correct: A — latency is prediction time.

Wrong:

B, C, D are unrelated.

Question 390

What is throughput in model serving?

A. The number of predictions served per unit time

B. The training time

C. The storage size

D. The feature count

Answer 390

Correct: A — throughput is predictions per time.

Wrong:

B, C, D are unrelated.

Question 391

Why might batch inference be preferred over real-time serving?

A. Predictions are not needed immediately and batch is cheaper

B. Batch is lower latency

C. Batch requires a REST API

D. Batch cannot use models

Answer 391

Correct: A — batch suits non-urgent scoring.

Wrong:

B, C, D are false.

Question 392

What is the role of a feature table in batch inference?

A. It stores features joined to the scoring data
B. It trains the model
C. It serves the model
D. It registers the model

Answer 392

Correct: A — feature tables supply features.

Wrong:

B, C, D are unrelated.

Question 393

What is a model's input schema in serving?

A. The expected structure of incoming data
B. The model weights
C. The training time
D. The cluster size

Answer 393

Correct: A — input schema defines expected data.

Wrong:

B, C, D are unrelated.

Question 394

What happens if serving input violates the model signature?

A. The request may fail validation

B. The model retrains

C. The endpoint deletes itself

D. The data is stored

Answer 394

Correct: A — signature validation can reject input.

Wrong:

B, C, D are false.

Question 395

What is a model's output schema in serving?

A. The structure of predictions returned by the model

B. The model weights

C. The training time

D. The cluster size

Answer 395

Correct: A — output schema defines returned predictions.

Wrong:

B, C, D are unrelated.

Question 396

What is the benefit of a feature store for real-time serving?

A. It provides low-latency online feature lookups

B. It trains models

C. It registers models

D. It stores notebooks

Answer 396

Correct: A — online tables serve low-latency features.

Wrong:

B, C, D are unrelated.

Question 397

What is a common cause of training-serving skew?

A. Features computed differently in training and serving
B. Using too many workers
C. Using too few features
D. Using a large model

Answer 397

Correct: A — inconsistent feature logic causes skew.

Wrong:

B, C, D are not the cause.

Question 398

How can Feature Store reduce training-serving skew?

A. By logging feature metadata with the model for consistent lookups

B. By training more models

C. By deleting features

D. By reducing cluster size

Answer 398

Correct: A — logged metadata enables consistent feature lookups.

Wrong:

B, C, D do not reduce skew.

Question 399

Why is monitoring prediction distributions important?

A. Shifts can indicate drift or data issues
B. It trains models
C. It registers models
D. It deletes models

Answer 399

Correct: A — prediction shifts signal drift.

Wrong:

B, C, D are unrelated.

Question 400

What is the best practice for retiring a model version in production?

A. Move the serving alias to a new version before archiving the old one

B. Delete the model immediately

C. Retrain from scratch

D. Stop all monitoring

Answer 400

Correct: A — move the alias first, then archive.

Wrong:

B causes downtime.

C is unnecessary.

D removes safeguards.