Databricks ML Associate

Questions 701–800 (Eighth 100)

Practice Exam Deck

How to Use This Deck

100 new questions (Q701–Q800), none repeated from Q1–Q700.

Each question slide has four options (A–D).

The next slide gives the correct answer and explains why the right option is right and why each wrong option is wrong.

This set emphasizes Feature Store creation parameters, MLflow project/run tags, serving endpoint CRUD, registry operations, evaluator parameters, Hyperopt + MLflow, DLT modes, and CLI/API patterns.

Question 701

Which parameter identifies the primary key when creating a Feature Store table?

A. `primary_keys`

B. `keys`

C. `index_cols`

D. `row_id`

Answer 701

Correct: A — `primary_keys` defines the key columns.

Wrong:

B, C, D are not the Feature Store parameter.

Question 702

Which parameter identifies the timestamp column when creating a time-series feature table?

A. `timestamp_keys`

B. `event_time`

C. `time_col`

D. `ts_col`

Answer 702

Correct: A — `timestamp_keys` defines time columns.

Wrong:

B, C, D are not the Feature Store parameter.

Question 703

Which parameter gives a feature table a human-readable description?

A. `description`

B. `comment`

C. `metadata`

D. `notes`

Answer 703

Correct: A — `description` sets the table description.

Wrong:

B, C, D are not the Feature Store parameter.

Question 704

Which parameter specifies the database or schema for a feature table?

A. `database_name` or schema name in the table name

B. `catalog`

C. `namespace`

D. `prefix`

Answer 704

Correct: A — the table name includes the schema.

Wrong:

B, C, D are not the Feature Store parameter.

Question 705

What does the `tags` parameter on a feature table do?

A. Attaches key-value metadata to the table

B. Trains the model

C. Serves the model

D. Registers the model

Answer 705

Correct: A — tags attach metadata.

Wrong:

B, C, D are unrelated.

Question 706

What is an MLflow project?

A. A packaged ML code project with an MLproject file

B. A model

C. A feature table

D. A cluster

Answer 706

Correct: A — projects package code with an MLproject file.

Wrong:

B, C, D are unrelated.

Question 707

What file defines an MLflow project?

A. `MLproject`

B. `project.yaml`

C. `mlflow.json`

D. `requirements.txt`

Answer 707

Correct: A — `MLproject` defines a project.

Wrong:

B, C, D are not the project definition file.

Question 708

What does an MLflow project's entry point define?

A. A runnable command and its parameters

B. A model

C. A feature table

D. A cluster

Answer 708

Correct: A — entry points define runnable commands.

Wrong:

B, C, D are unrelated.

Question 709

Which MLflow API runs an MLflow project?

A. `mlflow.run()`

B. `mlflow.start()`

C. `mlflow.execute()`

D. `mlflow.launch()`

Answer 709

Correct: A — `mlflow.run()` runs a project.

Wrong:

B, C, D are not the API.

Question 710

What is a run tag used for in MLflow?

A. Attaching metadata to a run

B. Training a model

C. Serving a model

D. Registering a model

Answer 710

Correct: A — tags attach metadata.

Wrong:

B, C, D are unrelated.

Question 711

Which MLflow API creates a serving endpoint?

A. `mlflow.deployments.create_endpoint()`

B. `mlflow.serve()`

C. `mlflow.endpoint()`

D. `mlflow.deploy()`

Answer 711

Correct: A — `create_endpoint()` creates endpoints.

Wrong:

B, C, D are not the API.

Question 712

Which MLflow API queries a serving endpoint?

A. `mlflow.deployments.predict()`

B. `mlflow.predict()`

C. `mlflow.query()`

D. `mlflow.infer()`

Answer 712

Correct: A — `predict()` queries endpoints.

Wrong:

B, C, D are not the API.

Question 713

Which MLflow API lists serving endpoints?

A. `mlflow.deployments.list_endpoints()`

B. `mlflow.endpoints()`

C. `mlflow.list_serve()`

D. `mlflow.deploy_list()`

Answer 713

Correct: A — `list_endpoints()` lists endpoints.

Wrong:

B, C, D are not the API.

Question 714

Which MLflow API deletes a serving endpoint?

A. `mlflow.deployments.delete_endpoint()`

B. `mlflow.delete_serve()`

C. `mlflow.remove_endpoint()`

D. `mlflow.endpoint_delete()`

Answer 714

Correct: A — `delete_endpoint()` deletes endpoints.

Wrong:

B, C, D are not the API.

Question 715

What does a serving endpoint's traffic configuration control?

A. How requests are routed to served entities
B. The training algorithm
C. The feature table
D. The cluster size

Answer 715

Correct: A — traffic configs route requests.

Wrong:

B, C, D are unrelated.

Question 716

What is a served entity in a serving endpoint?

A. A model version or feature spec the endpoint serves
B. A training job
C. A feature table
D. A cluster

Answer 716

Correct: A — served entities define what is served.

Wrong:

B, C, D are unrelated.

Question 717

Which MLflow API creates a registered model in Unity Catalog?

A. `mlflow.create_model()` or `register_model()`

B. `mlflow.new_model()`

C. `mlflow.add_model()`

D. `mlflow.make_model()`

Answer 717

Correct: A — models are created via those APIs.

Wrong:

B, C, D are not the API.

Question 718

Which MLflow API sets a registered model's description?

A. `mlflow.update_model()`

B. `mlflow.set_model_description()`

C. `mlflow.describe_model()`

D. `mlflow.model_info()`

Answer 718

Correct: A — `update_model()` updates model metadata.

Wrong:

B, C, D are not the API.

Question 719

Which MLflow API transitions a model version's stage?

A. `mlflow.transition_model_version_stage()`

B. `mlflow.set_stage()`

C. `mlflow.move_version()`

D. `mlflow.change_stage()`

Answer 719

Correct: A — that API transitions stages.

Wrong:

B, C, D are not the API.

Question 720

Which MLflow API deletes a model version?

A. `mlflow.delete_model_version()`

B. `mlflow.remove_version()`

C. `mlflow.drop_version()`

D. `mlflow.model_delete()`

Answer 720

Correct: A — `delete_model_version()` deletes versions.

Wrong:

B, C, D are not the API.

Question 721

Which MLflow API updates a model version's description?

A. `mlflow.update_model_version()`

B. `mlflow.set_version_desc()`

C. `mlflow.describe_version()`

D. `mlflow.version_info()`

Answer 721

Correct: A — `update_model_version()` updates metadata.

Wrong:

B, C, D are not the API.

Question 722

What does `setLabelCol()` do in a Spark ML evaluator?

A. Sets the target column for evaluation
B. Sets the metric
C. Sets the prediction column
D. Sets the model

Answer 722

Correct: A — it sets the label column.

Wrong:

B is `setMetricName`.

C is `setPredictionCol`.

D is unrelated.

Question 723

What does `setRawPredictionCol()` do in a Spark ML evaluator?

A. Sets the column containing raw prediction scores

B. Sets the target column

C. Sets the metric

D. Sets the model

Answer 723

Correct: A — it sets the raw prediction column.

Wrong:

B is `setLabelCol`.

C is `setMetricName`.

D is unrelated.

Question 724

What does `setProbabilityCol()` do in a Spark ML evaluator?

A. Sets the column containing predicted probabilities

B. Sets the target column

C. Sets the metric

D. Sets the model

Answer 724

Correct: A — it sets the probability column.

Wrong:

B is `setLabelCol`.

C is `setMetricName`.

D is unrelated.

Question 725

Which Spark ML evaluator computes `areaUnderPR`?

A. `BinaryClassificationEvaluator`

B. `RegressionEvaluator`

C. `ClusteringEvaluator`

D. `MulticlassClassificationEvaluator`

Answer 725

Correct: A — binary evaluator computes `areaUnderPR`.

Wrong:

B, C, D are other evaluators.

Question 726

Which Spark ML evaluator supports `accuracy`?

A. `MulticlassClassificationEvaluator`
B. `RegressionEvaluator`
C. `ClusteringEvaluator`
D. `BinaryClassificationEvaluator`

Answer 726

Correct: A — multiclass evaluator supports accuracy.

Wrong:

B, C, D are other evaluators.

Question 727

Which Spark ML evaluator supports `mae`?

A. `RegressionEvaluator`
B. `BinaryClassificationEvaluator`
C. `MulticlassClassificationEvaluator`
D. `ClusteringEvaluator`

Answer 727

Correct: A — regression evaluator supports MAE.

Wrong:

B, C, D are other evaluators.

Question 728

What is an MLflow autologging benefit?

A. It removes the need to log params/metrics manually
B. It trains faster
C. It serves faster
D. It registers models

Answer 728

Correct: A — autologging automates logging.

Wrong:

B, C, D are false.

Question 729

How does Hyperopt integrate with MLflow?

A. By logging each trial's results to MLflow when enabled
B. By replacing MLflow
C. By serving models
D. By registering models

Answer 729

Correct: A — Hyperopt can log trials to MLflow.

Wrong:

B, C, D are false.

Question 730

Which Hyperopt feature logs trial results to MLflow?

A. `MLflowLogger` or auto-logging integration

B. `SparkTrials` only

C. `STATUS_OK`

D. `fmin`

Answer 730

Correct: A — MLflow logging is enabled via the integration.

Wrong:

B is distribution.

C is a status.

D is the optimizer.

Question 731

What is a Delta Live Tables pipeline's continuous mode?

A. It keeps the pipeline running to process new data
B. It runs once and stops
C. It trains models
D. It serves models

Answer 731

Correct: A — continuous pipelines keep processing new data.

Wrong:

B is triggered mode.

C, D are unrelated.

Question 732

What is a Delta Live Tables pipeline's triggered mode?

A. It processes data and then stops
B. It runs continuously
C. It trains models
D. It serves models

Answer 732

Correct: A — triggered pipelines run and stop.

Wrong:

B is continuous.

C, D are unrelated.

Question 733

What is a Delta Live Tables pipeline's event log used for?

A. Tracking pipeline execution and data quality
B. Training models
C. Serving models
D. Registering models

Answer 733

Correct: A — event logs track pipeline execution.

Wrong:

B, C, D are unrelated.

Question 734

What is a Delta Live Tables pipeline's materialized view?

A. A precomputed query result stored as a table

B. A model

C. A feature table

D. A cluster

Answer 734

Correct: A — materialized views store query results.

Wrong:

B, C, D are unrelated.

Question 735

Which is a valid Databricks CLI command to list jobs?

A. `databricks jobs list`

B. `databricks list jobs`

C. `databricks jobs all`

D. `databricks job ls`

Answer 735

Correct: A — `jobs list` lists jobs.

Wrong:

B, C, D are not the command syntax.

Question 736

Which is a valid Databricks CLI command to run a job?

A. `databricks jobs run-now`

B. `databricks run job`

C. `databricks jobs start`

D. `databricks job execute`

Answer 736

Correct: A — `jobs run-now` starts a job.

Wrong:

B, C, D are not the CLI syntax.

Question 737

Which is a valid Databricks CLI command to list clusters?

A. `databricks clusters list`

B. `databricks list clusters`

C. `databricks clusters all`

D. `databricks cluster ls`

Answer 737

Correct: A — `clusters list` lists clusters.

Wrong:

B, C, D are not the CLI syntax.

Question 738

Which REST API endpoint lists jobs?

A. `GET /api/2.1/jobs/list`

B. `GET /api/2.1/jobs/all`

C. `GET /api/2.1/list/jobs`

D. `GET /api/2.1/jobs`

Answer 738

Correct: A — `/jobs/list` lists jobs.

Wrong:

B, C, D are not the correct endpoint.

Question 739

Which REST API endpoint triggers a job run?

A. `POST /api/2.1/jobs/run-now`

B. `POST /api/2.1/jobs/start`

C. `POST /api/2.1/jobs/execute`

D. `POST /api/2.1/jobs/launch`

Answer 739

Correct: A — `run-now` triggers a run.

Wrong:

B, C, D are not the endpoint.

Question 740

Which REST API endpoint lists clusters?

A. `GET /api/2.0/clusters/list`

B. `GET /api/2.0/clusters/all`

C. `GET /api/2.0/clusters`

D. `GET /api/2.0/list/clusters`

Answer 740

Correct: A — `/clusters/list` lists clusters.

Wrong:

B, C, D are not the endpoint.

Question 741

What is a Databricks personal access token used for?

A. Authenticating API and CLI calls

B. Training models

C. Serving models

D. Registering models

Answer 741

Correct: A — PATs authenticate API/CLI calls.

Wrong:

B, C, D are unrelated.

Question 742

Which is a best practice for personal access tokens?

A. Rotate tokens regularly and store them securely
B. Share tokens publicly
C. Commit tokens to Git
D. Use the same token for all users

Answer 742

Correct: A — rotate and store tokens securely.

Wrong:

B, C, D are unsafe.

Question 743

What is OAuth used for in Databricks?

A. Delegated authentication for users and service principals

B. Training models

C. Serving models

D. Registering models

Answer 743

Correct: A — OAuth supports delegated authentication.

Wrong:

B, C, D are unrelated.

Question 744

What is a Unity Catalog model's fully qualified name format?

A. `catalog.schema.model`

B. `database.model`

C. `workspace.model`

D. `hive.model`

Answer 744

Correct: A — UC models use `catalog.schema.model`.

Wrong:

B, C, D are not UC formats.

Question 745

Which privilege is needed to register a model in Unity Catalog?

A. `USE CATALOG` and `USE SCHEMA` plus create permission

B. `SELECT` only

C. `MODIFY` only

D. `EXECUTE` only

Answer 745

Correct: A — registration needs catalog/schema use and create.

Wrong:

B, C, D are insufficient alone.

Question 746

Which privilege is needed to serve a model from Unity Catalog?

A. `EXECUTE`

B. `SELECT`

C. `MODIFY`

D. `READ`

Answer 746

Correct: A — `EXECUTE` allows model use.

Wrong:

B, C, D are table privileges.

Question 747

What is a model version's URI in Unity Catalog?

A. `models:/catalog.schema.model/version`

B. `runs:/version/model`

C. `models:/catalog/version`

D. `models://version`

Answer 747

Correct: A — UC model URIs use `models:/<name>/<version>`.

Wrong:

B is a run URI.

C, D are malformed.

Question 748

What does `mlflow.pyfunc.load_model()` accept?

A. A model URI such as a registry URI or run URI

B. A DataFrame

C. A cluster

D. A feature table

Answer 748

Correct: A — it loads by model URI.

Wrong:

B, C, D are not model URIs.

Question 749

What does `mlflow.pyfunc.predict()` accept after loading?

A. A pandas or Spark DataFrame of inputs

B. A model

C. A feature table

D. A cluster

Answer 749

Correct: A — `predict()` accepts input DataFrames.

Wrong:

B, C, D are not inputs.

Question 750

What is an MLflow model URI `runs:/...` used for?

A. Loading a model artifact from a run

B. Loading a feature table

C. Loading a cluster

D. Loading a notebook

Answer 750

Correct: A — run URIs load artifacts from a run.

Wrong:

B, C, D are not loaded by run URIs.

Question 751

What is an MLflow run's source name?

A. The file or notebook that started the run

B. The model

C. The feature table

D. The cluster

Answer 751

Correct: A — source name identifies the run origin.

Wrong:

B, C, D are unrelated.

Question 752

What is an MLflow run's user?

A. The user who started the run

B. The model

C. The feature table

D. The cluster

Answer 752

Correct: A — the run user is the initiator.

Wrong:

B, C, D are unrelated.

Question 753

What is an MLflow run's start time?

A. When the run began

B. When the run ended

C. When the model was served

D. When the model was registered

Answer 753

Correct: A — start time is when the run began.

Wrong:

B, C, D are other times.

Question 754

What is an MLflow run's end time?

A. When the run ended

B. When the run began

C. When the model was served

D. When the model was registered

Answer 754

Correct: A — end time is when the run ended.

Wrong:

B, C, D are other times.

Question 755

What is an MLflow experiment's artifact location?

A. Where artifacts for runs in the experiment are stored
B. Where models are registered
C. Where clusters are configured
D. Where notebooks are stored

Answer 755

Correct: A — artifact locations define storage for runs.

Wrong:

B, C, D are unrelated.

Question 756

What is an MLflow experiment's lifecycle stage?

A. Whether the experiment is active or deleted
B. The model's stage
C. The cluster's mode
D. The feature table's mode

Answer 756

Correct: A — lifecycle stages track active/deleted experiments.

Wrong:

B, C, D are unrelated.

Question 757

What is a nested run in MLflow?

A. A run started inside another run

B. A model

C. A feature table

D. A cluster

Answer 757

Correct: A — nested runs belong to a parent run.

Wrong:

B, C, D are unrelated.

Question 758

Which flag starts a nested run?

A. `nested=True`

B. `child=True`

C. `sub=True`

D. `nest=True`

Answer 758

Correct: A — `nested=True` starts a nested run.

Wrong:

B, C, D are not the API flag.

Question 759

What is an MLflow parent run?

A. The outer run containing nested runs

B. A model

C. A feature table

D. A cluster

Answer 759

Correct: A — the parent run contains nested runs.

Wrong:

B, C, D are unrelated.

Question 760

What does `mlflow.get_parent_run()` return?

A. The parent run of the active nested run

B. A model

C. A feature table

D. A cluster

Answer 760

Correct: A — it returns the parent run.

Wrong:

B, C, D are unrelated.

Question 761

Which MLflow API lists nested runs?

A. `mlflow.search_runs()` with a parent filter

B. `mlflow.list_children()`

C. `mlflow.children()`

D. `mlflow.nested_runs()`

Answer 761

Correct: A — search with a parent run filter.

Wrong:

B, C, D are not the API.

Question 762

What is a Feature Store training set's source?

A. A label DataFrame joined with feature lookups
B. A model
C. A cluster
D. A notebook

Answer 762

Correct: A — training sets combine labels and features.

Wrong:

B, C, D are unrelated.

Question 763

What does a training set's label column represent?

A. The target value for each row

B. The feature table

C. The model

D. The cluster

Answer 763

Correct: A — label columns hold targets.

Wrong:

B, C, D are unrelated.

Question 764

What does `exclude_columns` do in a FeatureLookup?

A. Excludes specific columns from the lookup

B. Trains a model

C. Serves a model

D. Registers a model

Answer 764

Correct: A — `exclude_columns` removes columns.

Wrong:

B, C, D are unrelated.

Question 765

What does a FeatureLookup's `lookup_key` specify?

A. The join key used to match features to the label DataFrame

B. The model

C. The cluster

D. The metric

Answer 765

Correct: A — `lookup_key` is the join key.

Wrong:

B, C, D are unrelated.

Question 766

What does a FeatureLookup's `feature_names` specify?

A. The specific feature columns to look up

B. The model

C. The cluster

D. The metric

Answer 766

Correct: A — `feature_names` selects columns.

Wrong:

B, C, D are unrelated.

Question 767

What is a model's input example used for in MLflow?

A. Documenting expected input format for serving
B. Training a model
C. Registering a model
D. Deleting a model

Answer 767

Correct: A — input examples document expected inputs.

Wrong:

B, C, D are unrelated.

Question 768

What does `mlflow.models.infer_signature()` need to work?

A. Sample input data and a model

B. A cluster

C. A feature table

D. A notebook

Answer 768

Correct: A — it needs sample input and model.

Wrong:

B, C, D are unrelated.

Question 769

What is an MLflow model signature's input schema?

A. Column names and types expected at inference
B. The model weights
C. The training data
D. The cluster

Answer 769

Correct: A — input schema defines expected columns/types.

Wrong:

B, C, D are unrelated.

Question 770

What is an MLflow model signature's output schema?

A. Column names and types of predictions
B. The model weights
C. The training data
D. The cluster

Answer 770

Correct: A — output schema defines prediction columns/types.

Wrong:

B, C, D are unrelated.

Question 771

What does a model signature's `params` field describe?

A. Inference-time parameters

B. Training data

C. Model weights

D. Cluster settings

Answer 771

Correct: A — `params` describes inference parameters.

Wrong:

B, C, D are unrelated.

Question 772

Why is a model signature important for pipelines?

A. It catches input/output mismatches early

B. It trains faster

C. It serves faster

D. It registers faster

Answer 772

Correct: A — signatures catch mismatches early.

Wrong:

B, C, D are false.

Question 773

What is a Spark ML `OneHotEncoder` estimator output?

A. Sparse vectors

B. Scalars

C. Strings

D. Buckets

Answer 773

Correct: A — one-hot outputs sparse vectors.

Wrong:

B, C, D are not outputs.

Question 774

What is a Spark ML `StringIndexer` default behavior for unseen labels?

A. It can throw an error or assign a default index

B. It ignores them

C. It drops them

D. It duplicates them

Answer 774

Correct: A — unseen labels can error or get defaults.

Wrong:

B, C, D are not its behavior.

Question 775

What is a Spark ML `VectorIndexer` used for?

A. Indexing categorical columns inside a vector
B. Scaling numeric features
C. Imputing values
D. Assembling vectors

Answer 775

Correct: A — VectorIndexer indexes categoricals in vectors.

Wrong:

B, C, D are other transformers.

Question 776

What is a Spark ML `SQLTransformer` used for?

A. Applying a SQL query as a transformation
B. Training a model
C. Serving a model
D. Registering a model

Answer 776

Correct: A — SQLTransformer applies SQL transformations.

Wrong:

B, C, D are unrelated.

Question 777

What is a Spark ML `Imputer` default strategy?

A. Mean

B. Median

C. Mode

D. Zero

Answer 777

Correct: A — the default strategy is mean.

Wrong:

B, C, D are other strategies.

Question 778

What is a Spark ML `VectorAssembler` output?

A. A single feature vector column

B. A scalar

C. A string

D. A bucket

Answer 778

Correct: A — it outputs a feature vector.

Wrong:

B, C, D are not outputs.

Question 779

What is a Spark ML `Bucketizer` used for?

A. Binning continuous features

B. Scaling features

C. Encoding categories

D. Imputing values

Answer 779

Correct: A — Bucketizer bins continuous features.

Wrong:

B, C, D are other transformers.

Question 780

What is a Spark ML `QuantileDiscretizer` used for?

A. Binning continuous features by quantiles

B. Scaling features

C. Encoding categories

D. Imputing values

Answer 780

Correct: A — quantile discretization uses quantiles.

Wrong:

B, C, D are other transformers.

Question 781

What is a Spark ML `MinMaxScaler` output range?

A. $[0, 1]$
B. $[-1, 1]$
C. $[0, 100]$
D. Unbounded

Answer 781

Correct: A — MinMaxScaler outputs $[0, 1]$ by default.

Wrong:

B, C, D are other ranges.

Question 782

What is a Spark ML `StandardScaler` output?

A. Zero-mean, unit-variance features

B. $[0, 1]$ features

C. $[-1, 1]$ features

D. Unbounded features

Answer 782

Correct: A — StandardScaler z-scores features.

Wrong:

B, C, D are other scales.

Question 783

What is a Spark ML `Normalizer` used for?

A. Scaling each row to unit norm

B. Scaling each column

C. Encoding categories

D. Imputing values

Answer 783

Correct: A — Normalizer scales rows to unit norm.

Wrong:

B, C, D are other transformers.

Question 784

What is a Spark ML `MaxAbsScaler` used for?

A. Scaling by the maximum absolute value

B. Scaling to $[0, 1]$

C. Encoding categories

D. Imputing values

Answer 784

Correct: A — MaxAbsScaler divides by max abs.

Wrong:

B, C, D are other methods.

Question 785

What is a Spark ML `PolynomialExpansion` output?

A. Polynomial features up to a degree
B. A single feature
C. A category
D. A bucket

Answer 785

Correct: A — it outputs polynomial features.

Wrong:

B, C, D are not outputs.

Question 786

What is a Spark ML `DCT` used for?

A. Discrete cosine transform on a vector
B. Scaling
C. Encoding
D. Imputing

Answer 786

Correct: A — DCT applies a cosine transform.

Wrong:

B, C, D are other transformations.

Question 787

What is a Spark ML `ElementwiseProduct` used for?

A. Multiplying each feature by a scalar

B. Scaling to $[0, 1]$

C. Encoding categories

D. Imputing values

Answer 787

Correct: A — ElementwiseProduct scales features.

Wrong:

B, C, D are other methods.

Question 788

What is a Spark ML `Interaction` transformer used for?

A. Creating interaction features between vectors
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 788

Correct: A — Interaction creates interaction features.

Wrong:

B, C, D are other transformers.

Question 789

What is a Spark ML `VectorSlicer` used for?

A. Selecting a subset of features from a vector
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 789

Correct: A — VectorSlicer selects features.

Wrong:

B, C, D are other transformers.

Question 790

What is a Spark ML `ChiSqSelector` used for?

A. Selecting features using a chi-squared test
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 790

Correct: A — ChiSqSelector is a feature selector.

Wrong:

B, C, D are other transformers.

Question 791

What is a Spark ML `VarianceThresholdSelector` used for?

A. Selecting features with variance above a threshold
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 791

Correct: A — it selects features by variance.

Wrong:

B, C, D are other transformers.

Question 792

What is a Spark ML `UnivariateFeatureSelector` used for?

A. Selecting features via univariate statistical tests
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 792

Correct: A — it selects features statistically.

Wrong:

B, C, D are other transformers.

Question 793

What is a Spark ML `RFormula` used for?

A. Specifying a model formula in R syntax
B. Scaling features
C. Encoding categories
D. Imputing values

Answer 793

Correct: A — RFormula uses R-style formulas.

Wrong:

B, C, D are other transformers.

Question 794

What is a Spark ML `DecisionTreeClassifier` used for?

A. Classification using decision trees

B. Clustering

C. Regression

D. Dimensionality reduction

Answer 794

Correct: A — DecisionTreeClassifier is a classifier.

Wrong:

B, C, D are other tasks.

Question 795

What is a Spark ML `DecisionTreeRegressor` used for?

A. Regression using decision trees

B. Classification

C. Clustering

D. Dimensionality reduction

Answer 795

Correct: A — DecisionTreeRegressor is a regressor.

Wrong:

B, C, D are other tasks.

Question 796

What is a Spark ML `GBTClassifier` used for?

A. Classification using gradient-boosted trees

B. Clustering

C. Regression

D. Dimensionality reduction

Answer 796

Correct: A — GBTClassifier is a classifier.

Wrong:

B, C, D are other tasks.

Question 797

What is a Spark ML `GBTRegressor` used for?

A. Regression using gradient-boosted trees

B. Classification

C. Clustering

D. Dimensionality reduction

Answer 797

Correct: A — GBTRegressor is a regressor.

Wrong:

B, C, D are other tasks.

Question 798

What is a Spark ML `RandomForestClassifier` used for?

A. Classification using random forests

B. Clustering

C. Regression

D. Dimensionality reduction

Answer 798

Correct: A — RandomForestClassifier is a classifier.

Wrong:

B, C, D are other tasks.

Question 799

What is a Spark ML `RandomForestRegressor` used for?

A. Regression using random forests

B. Classification

C. Clustering

D. Dimensionality reduction

Answer 799

Correct: A — RandomForestRegressor is a regressor.

Wrong:

B, C, D are other tasks.

Question 800

What is a Spark ML `LinearSVC` used for?

A. Linear support vector classification
B. Clustering
C. Regression
D. Dimensionality reduction

Answer 800

Correct: A — LinearSVC is a linear classifier.

Wrong:

B, C, D are other tasks.