Databricks ML Associate
100 Databricks-Specific Q&A

Exam Blueprint Aligned

Blueprint Weighting (Official)

Databricks Machine Learning: 38%

Model Development: 31%

ML Workflows: 19%

Model Deployment: 12%

Source: Databricks exam guide [citation:2][citation:5][citation:12]

Q1 — MLflow Autologging

An experiment's parameters, metrics, and model artifact must be recorded automatically. Best Databricks approach?

A. Manual notes in a notebook.

B. `mlflow.autolog()` or explicit `log_param`/`log_metric` within a run.

C. Print statements.

D. A spreadsheet.

A1 — Answer: B

Correct: B.

Autologging captures params, metrics, model for supported libraries in one line.

Runs become comparable in MLflow UI by default.

Q2 — UC Model Registry Promotion

Model must be promoted to production with auditable lifecycle. Recommended Databricks approach?

A. Copy the pickle file.
B. Register in Unity Catalog with versions and aliases like `@champion`; grant access via catalog.
C. Email the pickle.
D. Rename the artifact.

A2 — Answer: B

Correct: B.

Registry keeps lineage to producing run; promotion is explicit alias move.

Aliases replaced old stages; apps reference `@champion`.

Q3 — Hyperopt Parallelization

Hyperparameters must be tuned in parallel across a Databricks cluster. Which component?

A. A sequential loop.
B. Hyperopt with `SparkTrials`.
C. Manual tuning.
D. One configuration.

A3 — Answer: B

Correct: B.

`SparkTrials` distributes trials across workers.

TPE search converges faster than grid; trials log to MLflow automatically.

Q4 — Feature Store for Training-Serving Skew

Features in training must match inference. Which Databricks capability?

A. Recompute in both places.
B. Databricks Feature Store (Feature Engineering in Unity Catalog).
C. Shared spreadsheet.
D. Copy the code.

A4 — Answer: B

Correct: B.

One authoritative definition served offline and online.

Point-in-time joins prevent using values unavailable at prediction time.

Q5 — Spark ML Pipeline Purpose

A pipeline includes indexing, assembling, classifier. Why wrap in `Pipeline`?

A. Run each step manually.

B. `Pipeline` fit once → `PipelineModel` applies same transforms at inference.

C. Separate notebooks.

D. SQL only.

A5 — Answer: B

Correct: B.

Fitted preprocessing travels with model.

Prevents applying different encoding at serving.

Q6 — Imbalanced Classification

98% accuracy with 2% positives. What should be done?

A. Ship it.

B. Evaluate precision, recall, F1, PR-AUC; consider class weights/resampling on training only.

C. Increase epochs.

D. Add features.

A6 — Answer: B

Correct: B.

Majority-class predictor scores 98%.

Precision/recall expose minority performance; resample only training split.

Q7 — Deep Learning Multi-GPU

Deep learning model must train on multiple GPUs on Databricks. Which tool?

A. Single GPU only.
B. `TorchDistributor` (or Ray on Databricks).
C. CPU training.
D. Reduce model.

A7 — Answer: B

Correct: B.

`TorchDistributor` runs PyTorch distributed training on Databricks cluster.

Start single-node multi-GPU; scale to multi-node only if needed.

Q8 — Real-Time Serving Autoscaling

Model must serve real-time predictions with autoscaling. Which Databricks component?

A. A notebook.
B. Mosaic AI Model Serving with registered model version, scaling on request volume.
C. A job cluster.
D. A dashboard.

A8 — Answer: B

Correct: B.

Model Serving handles REST endpoints and autoscaling.

Notebook/job cluster are not serving infrastructure.

Q9 — Feature Store Creation Method

Which method creates a feature table in Unity Catalog?

A. `fe.create_table()`.
B. `spark.sql("CREATE TABLE")`.
C. `mlflow.log_artifact()`.
D. `dbutils.fs.mkdirs()`.

A9 — Answer: A

Correct: A.

`FeatureEngineeringClient.create_table()` registers schema, PKs, timestamp keys.

Plain CREATE TABLE is not a feature table.

Q10 — Training Set from Feature Store

Which method creates training set with point-in-time lookups?

A. `fe.create_training_set()`.

B. `spark.read.table()`.

C. `mlflow.log_model()`.

D. `Pipeline.fit()`.

A10 — Answer: A

Correct: A.

`create_training_set()` joins features as-of label timestamp.

Prevents leakage from future feature values.

Q11 — Online vs Offline Feature Tables

Primary difference between online and offline feature tables?

A. Online for training; offline for serving.
B. Online = low-latency inference; offline = batch training/analysis.
C. Online cannot be in UC.
D. Offline no PIT joins.

Correct: B.

Online backed by low-latency store (e.g., DynamoDB).

Offline for batch reads; automatically synced.

Q12 — Best Run via MLflow Client

How to identify best run using MLflow Client API?

A. `client.search_runs()` with filter/order.

B. `client.list_experiments()`.

C. `client.get_run()`.

D. `client.delete_run()`.

A12 — Answer: A

Correct: A.

`search_runs()` supports `metrics.accuracy > 0.9` and ordering.

`get_run()` needs known run_id.

Q13 — Manual Metric Logging

Which function logs a numeric metric to MLflow?

A. `mlflow.log_metric()`.
B. `mlflow.log_param()`.
C. `mlflow.log_artifact()`.
D. `mlflow.set_tag()`.

A13 — Answer: A

Correct: A.

`log_param` = hyperparameters; `log_artifact` = files.

Q14 — MLflow UI Source

Where in MLflow UI to find notebook that created a run?

A. “Source” link for the run.
B. “Artifacts” tab only.
C. Cluster event log.
D. Home page.

A14 — Answer: A

Correct: A.

Source links to notebook and revision used.

Q15 — Register Model UC via Client

Which function registers model in Unity Catalog?

A. `mlflow.register_model()`.
B. `mlflow.log_model()`.
C. `mlflow.save_model()`.
D. `mlflow.export_model()`.

A15 — Answer: A

Correct: A.

`register_model(model_uri, name)` in UC registry.

`log_model` logs artifact to run but does not register.

Q16 — UC vs Workspace Registry Benefit

Key benefit of UC registry over workspace registry?

A. Accessible only to original author.
B. Governed, discoverable, promotable across workspaces at account level.
C. No version numbers needed.
D. Auto-deployed to endpoints.

A16 — Answer: B

Correct: B.

UC = account-level governance and cross-workspace discovery.

Q17 — Promote Code vs Model

When is promoting code preferred over promoting models?

A. Simple sklearn model.

B. Complex training pipeline, small model artifact.

C. Large artifact, expensive training.

D. Only AutoML.

Correct: B.

Promote code when pipeline is complex and training is cheap: rebuild in each environment.

Promote model when training is expensive and artifact small.

Q18 — Model Tags

How to set tag on registered model in UC?

A. `mlflow.set_model_tag()`.
B. `spark.conf.set()`.
C. `dbutils.notebook.exit()`.
D. `display()`.

A18 — Answer: A

Correct: A.

Q19 — Champion Alias

How is challenger promoted to champion using aliases?

A. Delete old model, upload new.

B. Move `@champion` alias to new version.

C. Rename model.

D. Create new experiment.

A19 — Answer: B

Correct: B.

Aliases are mutable pointers; consumers keep using `@champion`.

Q20 — MLOps Best Practice

Which setup reflects sound MLOps strategy?

A. One shared workspace, edit prod directly.
B. Separate dev/staging/prod with versioned code and tracked artifacts.
C. Manual file copies.
D. No prior run metadata.

A20 — Answer: B

Correct: B.

Q21 — ML Runtime Advantage

Primary advantage of Databricks Runtime for ML?

A. Common ML libraries and integrations preinstalled.
B. Eliminates all package conflicts.
C. Removes need for permissions.
D. Skips feature engineering.

A21 — Answer: A

Correct: A.

ML Runtime includes scikit-learn, MLflow, etc.

Q22 — AutoML Output

Useful AutoML output for reviewing/customizing winning approach?

A. Screenshot of leaderboard only.
B. Generated notebooks for data exploration and best trial.
C. Binary with no source.
D. Cluster event log.

A22 — Answer: B

Correct: B.

Q23 — MLflow Nested Runs

Tracking hyperparameter tuning with many child runs. Best pattern?

A. Separate experiment per trial.

B. Child runs inside parent using `start_run(nested=True)`.

C. All trials to one run without nesting.

D. No logging.

A23 — Answer: B

Correct: B.

Q24 — AutoML Constraint

Which constraint can be specified for AutoML?

A. `timeout_minutes`.
B. Random seed only.
C. Only neural networks.
D. Fixed features.

A24 — Answer: A

Correct: A.

Q25 — Feature Store Sync

How are online tables kept synchronized with offline?

A. Manual export/import.
B. Databricks automatically publishes online from offline.
C. Never synced.
D. Only batch jobs.

A25 — Answer: B

Correct: B.

Q26 — UC Registry Lineage

What information does UC registry retain for lineage?

A. Only model name.

B. Run ID, source experiment, model version lineage.

C. Only file size.

D. Only upload timestamp.

A26 — Answer: B

Correct: B.

Q27 — Model Signature

Purpose of model signature in MLflow?

A. Defines expected input/output schema.
B. Encrypts model.
C. Compresses model.
D. Deletes model.

A27 — Answer: A

Correct: A.

Q28 — Tracking Server

Where are MLflow experiments/runs stored in Databricks?

A. Driver local disk.
B. Databricks-managed MLflow tracking service.
C. Git repository.
D. Cluster event log.

A28 — Answer: B

Correct: B.

Q29 — Registry Permission

Permission to promote model version to champion in UC?

A. USE CATALOG only.

B. EXECUTE on model.

C. MANAGE or APPLY TAG on model.

D. SELECT on training table.

A29 — Answer: C

Correct: C.

Q30 — AutoML Explainability

What does AutoML provide for interpretation?

A. SHAP-based feature importance and notebooks.
B. Only leaderboard.
C. Only predictions.
D. Binary.

A30 — Answer: A

Correct: A.

Q31 — Feature Lookup for Inference

Which method retrieves features for inference?

A. `fe.get_online_features()`.

B. `spark.read.table()` only.

C. `mlflow.load_model()`.

D. `dbutils.fs.head()`.

A31 — Answer: A

Correct: A.

Q32 — Model Signature Enforcement

Input doesn't match signature at serving. What happens?

A. Request fails validation.
B. Silently ignores mismatch.
C. Retrains automatically.
D. Updates signature.

A32 — Answer: A

Correct: A.

Q33 — Experiment Artifact Location

Default artifact storage location?

A. Local cluster disk.
B. MLflow experiment's artifact location (DBFS/cloud).
C. Notebook revision history.
D. Driver memory.

A33 — Answer: B

Correct: B.

Q34 — Deployment Readiness

Most important check before production?

A. Training accuracy only.
B. Performance on held-out test set and schema validation.
C. Notebook cell count.
D. Cluster size.

A34 — Answer: B

Correct: B.

Q35 — Feature Table Primary Keys

Purpose of primary keys in UC feature table?

A. Uniquely identify entity for PIT lookups.
B. Encrypt table.
C. Reduce storage.
D. Replace timestamp.

A35 — Answer: A

Correct: A.

Q36 — Run Hierarchy

Correct MLflow structure?

A. Experiments contain runs; runs contain artifacts.
B. Runs contain experiments.
C. Artifacts contain runs.
D. Unrelated.

A36 — Answer: A

Correct: A.

Q37 — Autologging Libraries

Which libraries does `mlflow.autolog()` support?

A. Only scikit-learn.

B. sklearn, XGBoost, LightGBM, Spark ML, and more.

C. Only TensorFlow.

D. Only PyTorch.

A37 — Answer: B

Correct: B.

Q38 — UC Three-Level Namespace

Three-level namespace for UC model?

A. `catalog.schema.model`.
B. `workspace.model.version`.
C. `database.table.column`.
D. `user.model.run`.

A38 — Answer: A

Correct: A.

Q39 — Regression Metrics

RMSE=25, MAE=20, R^2=0.85. Which two are correct?

A. Model captures 85% of variance.

B. RMSE ¿ MAE indicates larger errors exist.

C. Model is underfitting.

D. MAE is 25.

A39 — Answer: A and B

Correct: A, B.

R^2=0.85 → 85% variance explained.

RMSE ¿ MAE → some large errors penalized by squaring.

Q40 — Imbalanced Metrics

Best metric for highly imbalanced binary classification?

A. Accuracy.
B. PR-AUC.
C. MAE.
D. R^2.

A40 — Answer: B

Correct: B.

Q41 — One-Hot Encoding Limitation

When is one-hot encoding least appropriate?

A. Low cardinality.
B. Extremely high cardinality.
C. Linear models.
D. String feature.

A41 — Answer: B

Correct: B.

Q42 — Imputation for Skewed

Preferred imputation for highly skewed distribution?

A. Mean.
B. Median.
C. Mode.
D. Zero.

A42 — Answer: B

Correct: B.

Q43 — Log Transformation

When is log transform appropriate?

A. Normally distributed.
B. Right-skewed, spans orders of magnitude.
C. Categorical.
D. Binary target.

A43 — Answer: B

Correct: B.

Q44 — Estimator vs Transformer

Key difference in Spark ML?

A. Estimator has `fit()`; Transformer has `transform()`.
B. Estimator transforms; Transformer fits.
C. Identical.
D. Transformer cannot be in Pipeline.

A44 — Answer: A

Correct: A.

Q45 — Random Split

Which Spark operation randomly splits DataFrame?

A. `DataFrame.randomSplit()`.

B. `TrainValidationSplit`.

C. `CrossValidator`.

D. `DataFrame.where()`.

A45 — Answer: A

Correct: A.

Q46 — StringIndexer Before OHE

Before one-hot encoding string column, which stage required?

A. `StringIndexer`.
B. `VectorAssembler`.
C. `StandardScaler`.
D. `Imputer`.

A46 — Answer: A

Correct: A.

Q47 — VectorAssembler Purpose

Purpose of `VectorAssembler`?

A. Combine feature columns into vector column.
B. Impute missing.
C. Encode categorical.
D. Split data.

A47 — Answer: A

Correct: A.

Q48 — StandardScaler

Which scaler gives zero mean/unit variance?

A. MinMaxScaler.
B. StandardScaler.
C. MaxAbsScaler.
D. Normalizer.

A48 — Answer: B

Correct: B.

Q49 — Data Leakage

Standardizing before split. Risk?

A. Test statistics leak into train $\rightarrow$ optimistic eval.

B. Underfitting.

C. Features become categorical.

D. Nothing.

A49 — Answer: A

Correct: A.

Q50 — Pandas UDF Benefit

Benefit of vectorized pandas UDFs?

A. Batch processing via Arrow $\rightarrow$ faster.
B. Only single-node.
C. No Python.
D. Auto-handles missing.

A50 — Answer: A

Correct: A.

Q51 — Target Encoding

Most useful for which situation?

A. Low cardinality.
B. High cardinality.
C. Continuous with outliers.
D. Binary targets.

A51 — Answer: B

Correct: B.

Q52 — Overfitting

Good train, poor test. Called?

A. Underfitting.
B. Overfitting.
C. Leakage.
D. Regularization.

A52 — Answer: B

Correct: B.

Q53 — High Variance

High variance corresponds to?

A. Overfitting.
B. Underfitting.
C. Imbalance.
D. Scaling.

A53 — Answer: A

Correct: A.

Q54 — Class Weights

Mitigate imbalance without resampling?

A. Class weights.
B. Remove minority.
C. Duplicate majority.
D. Ignore.

A54 — Answer: A

Correct: A.

Q55 — RMSE vs MAE

Large errors costly. Best regression metric?

A. RMSE.
B. MAE.
C. Accuracy.
D. F1.

A55 — Answer: A

Correct: A.

Q56 — CrossValidator Grid

How is hyperparameter grid specified?

A. `ParamGridBuilder`.
B. Python dict.
C. CSV.
D. SQL.

A56 — Answer: A

Correct: A.

Q57 — Hyperopt Objective

What must objective return?

A. Scalar loss to minimize.
B. DataFrame.
C. Model only.
D. String.

A57 — Answer: A

Correct: A.

Q58 — Hyperopt Uniform

Uniform distribution?

A. `hp.uniform('x', low, high)`.
B. `hp.choice('x', [a,b])`.
C. `hp.quniform('x', low, high, q)`.
D. `hp.loguniform('x', low, high)`.

A58 — Answer: A

Correct: A.

Q59 — StringIndexer handleInvalid

Default for unseen label at transform?

A. Throws error unless `handleInvalid` set.
B. Ignores row.
C. Maps to 0.
D. Maps to most frequent.

A59 — Answer: A

Correct: A.

Q60 — OneHotEncoder Input

Input type for `OneHotEncoder`?

A. Numeric index from StringIndexer.
B. Raw string.
C. Continuous.
D. Vector.

A60 — Answer: A

Correct: A.

Q61 — PipelineModel

Fitting Pipeline produces?

A. PipelineModel.
B. DataFrame.
C. Cluster.
D. Experiment.

A61 — Answer: A

Correct: A.

Q62 — Spark ML Save Format

Spark ML model save format?

A. Directory with metadata and Parquet.
B. CSV.
C. Pickle.
D. JSON.

A62 — Answer: A

Correct: A.

Q63 — MLflow Flavor

What is MLflow model flavor?

A. Convention for saving/loading models from different libraries.
B. Cluster type.
C. Visualization.
D. Data source.

A63 — Answer: A

Correct: A.

Q64 — Imputer Strategies

Spark ML `Imputer` supports?

A. Mean, median, mode.

B. Only mean.

C. Only zero.

D. Only median.

A64 — Answer: A

Correct: A.

Q65 — CV numFolds

`numFolds` controls?

A. Number of k folds.
B. Models to deploy.
C. Features to select.
D. Clusters.

A65 — Answer: A

Correct: A.

Q66 — trainRatio

`trainRatio` controls?

A. Fraction for training.
B. Folds.
C. Trials.
D. Learning rate.

A66 — Answer: A

Correct: A.

Q67 — Binary Metric

Metric in `BinaryClassificationEvaluator`?

A. Area under ROC.
B. RMSE.
C. MAE.
D. R^2.

A67 — Answer: A

Correct: A.

Q68 — Regression Metric

Metric in `RegressionEvaluator`?

A. RMSE.
B. Precision.
C. Recall.
D. F1.

A68 — Answer: A

Correct: A.

Q69 — Multiclass Metric

Metric in `MulticlassClassificationEvaluator`?

A. F1.

B. RMSE.

C. MAE.

D. R^2.

A69 — Answer: A

Correct: A.

Q70 — FeatureHasher

Purpose?

A. Hash categorical to fixed-length vector without dictionary.
B. Impute.
C. Scale.
D. Split.

A70 — Answer: A

Correct: A.

Q71 — SQLTransformer

Purpose?

A. Apply SQL expression within Pipeline.
B. Train.
C. Evaluate.
D. Deploy.

A71 — Answer: A

Correct: A.

Q72 — VectorSlicer

Purpose?

A. Select subset of features from vector.
B. Combine columns.
C. Scale.
D. Impute.

A72 — Answer: A

Correct: A.

Q73 — PolynomialExpansion

Purpose?

A. Generate polynomial/interaction features.
B. Encode categorical.
C. Scale.
D. Split.

A73 — Answer: A

Correct: A.

Q74 — Binarizer

Purpose?

A. Threshold continuous into binary.
B. Encode categorical.
C. Scale.
D. Impute.

A74 — Answer: A

Correct: A.

Q75 — Cross-Validation Trade-off

Main trade-off vs single split?

A. CV faster.
B. CV more stable but more expensive.
C. Split always better.
D. Same cost.

A75 — Answer: B

Correct: B.

Q76 — Hyperopt Parallel

Hyperopt component for cluster parallelism?

A. `SparkTrials`.

B. `Trials`.

C. `fmin`.

D. `hp.uniform`.

A76 — Answer: A

Correct: A.

Q77 — TPE Algorithm

Hyperopt default via `tpe.suggest`?

A. Grid search.

B. Random search.

C. Tree-structured Parzen Estimator.

D. Exhaustive.

A77 — Answer: C

Correct: C.

Q78 — Summary Statistics

Command for summary statistics on Spark DataFrame?

A. `df.summary()`.

B. `df.show()`.

C. `df.printSchema()`.

D. `df.count()`.

A78 — Answer: A

Correct: A.

Q79 — IQR Outlier Removal

IQR method rule?

A. Below Q1 or above Q3.
B. $< Q1 - 1.5\,IQR$ or $> Q3 + 1.5\,IQR$.
C. Beyond mean.
D. $z > 1$.

A79 — Answer: B

Correct: B.

Q80 — Pandas API Refactor

Pandas notebook scale to big data, minimal refactor?

A. Rewrite in Scala.
B. Refactor to pandas API on Spark.
C. Use only SQL.
D. Sample to memory.

A80 — Answer: B

Correct: B.

Q81 — Two Continuous Features

Visualization for two continuous?

A. Scatter plot.
B. Bar chart.
C. Pie chart.
D. Box plot.

A81 — Answer: A

Correct: A.

Q82 — Two Categorical Features

Visualization for two categorical?

A. Grouped bar chart/heatmap.
B. Scatter.
C. Line.
D. Histogram.

A82 — Answer: A

Correct: A.

Q83 — Categorical Imputation

Imputation for categorical?

A. Mean.
B. Median.
C. Mode.
D. Interpolation.

A83 — Answer: C

Correct: C.

Q84 — Standard Deviation Outliers

Remove outliers by standard deviation?

A. Filter beyond $k\sigma$ from mean.

B. Remove above median.

C. Remove top 50%.

D. dropDuplicates.

A84 — Answer: A

Correct: A.

Q85 — Pandas to Spark

Convert Spark DF to pandas API on Spark?

A. `spark_df.pandas_api`.
B. `spark_df.to_pandas()`.
C. `spark_df.collect()`.
D. `spark_df.rdd`.

A85 — Answer: A

Correct: A.

Q86 — toPandas

What does `toPandas()` do?

A. Collects to local pandas on driver.
B. Converts to pandas API on Spark.
C. Saves to disk.
D. Deletes.

A86 — Answer: A

Correct: A.

Q87 — Crosstab

`df.crosstab()` produces?

A. Contingency table between two categorical.
B. Scatter.
C. Correlation matrix.
D. Histogram.

A87 — Answer: A

Correct: A.

Q88 — Null Handling

Drop rows with null?

A. `df.dropna()`.

B. `df.fillna()`.

C. `df.replace()`.

D. `df.filter()`.

A88 — Answer: A

Correct: A.

Q89 — Null Filling

Fill nulls with value?

A. `df.fillna()`.

B. `df.dropna()`.

C. `df.replace()`.

D. `df.filter()`.

A89 — Answer: A

Correct: A.

Q90 — Feature Store Training Set

Create training set from feature table?

A. `fe.create_training_set()`.

B. `spark.read.table()`.

C. `mlflow.log_model()`.

D. `Pipeline.fit()`.

A90 — Answer: A

Correct: A.

Q91 — Serving Endpoint Purpose

Purpose of Databricks serving endpoint?

A. REST API for real-time inference.
B. Store training data.
C. Run SQL.
D. Manage permissions.

A91 — Answer: A

Correct: A.

Q92 — Scale to Zero

Endpoint active only during work hours to minimize cost?

A. Scale to Zero.
B. Autoscaling only.
C. Larger cluster.
D. Versioning.

A92 — Answer: A

Correct: A.

Q93 — Batch vs Real-Time

Difference?

A. Batch large volumes on schedule; real-time individual low-latency.
B. Batch always faster.
C. Real-time cannot use MLflow.
D. Identical.

A93 — Answer: A

Correct: A.

Q94 — Traffic Splitting

Purpose?

A. Route percentage to different versions for A/B testing.
B. Increase training speed.
C. Reduce storage.
D. Impute missing.

A94 — Answer: A

Correct: A.

Q95 — Streaming Inference

Databricks capability for streaming inference?

A. Delta Live Tables with model serving.
B. Only batch.
C. Only real-time.
D. Only AutoML.

A95 — Answer: A

Correct: A.

Q96 — Serving Autoscaling

What does autoscaling do?

A. Adjusts serving instances based on traffic.
B. Trains new models.
C. Changes version.
D. Deletes experiments.

A96 — Answer: A

Correct: A.

Q97 — Deployment Consistency

Why identical feature engineering training/serving?

A. Avoid training-serving skew.
B. Reduce storage.
C. Speed startup.
D. Avoid UC.

A97 — Answer: A

Correct: A.

Q98 — Model Version Deployment

Serve specific model version from UC?

A. Reference by version or alias in endpoint config.
B. Export to CSV.
C. Rename model.
D. Only latest automatically.

A98 — Answer: A

Correct: A.

Q99 — Deployment Monitoring

Key monitoring aspect?

A. Prediction and data drift over time.
B. Only CPU.
C. Ignore post-deployment.
D. Only training accuracy.

A99 — Answer: A

Correct: A.

Q100 — Endpoint Logging

Capture request/response for monitoring?

A. Inference tables.
B. Cluster event logs.
C. Revision history.
D. Git.

A100 — Answer: A

Correct: A.

Exam Day Tips

48 questions / 90 minutes (not 100).

Watch for **"Select two/three"** questions.

Databricks-specific traps:

- UC registry vs workspace registry
- Aliases (`@champion`) not stages
- Feature Store for training-serving skew
- `SparkTrials` for parallel Hyperopt

Master Domain 1 (38%) + Domain 2 (31%) first.