

Responsible & Ethical AI: 50 Advanced Questions

Deep-dive questions: theory, technical mechanisms, and regulation

GARP RAI Exam Preparation

Q1: Rule Consequentialism Collapse - Advanced

QUESTION: A rule consequentialist refuses to kill one patient to save five because a rule against killing maximizes long-run utility (it preserves trust in medicine). Critics say the theory “collapses into deontology.” What is the core of that objection?

OPTIONS:

- ☐ A. Rule consequentialists never consider outcomes, so they were deontologists from the start
- ☒ B. If the rule must be obeyed even where breaking it would yield more utility, rule-adherence works as an independent constraint, which is exactly what deontologists claim
- ☐ C. Deontologists also choose their rules solely by maximizing aggregate utility, so the two share one justification
- ☐ D. Rules only matter case by case, so the theory reduces to act utilitarianism

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Rules are justified by their tendency to maximize utility, but acts are judged by conformity to the rule, not by case-specific outcomes. When conformity overrides a better isolated outcome, the rule behaves like a deontological red line.
- **Why A is wrong:** Rules are selected by their consequences, so outcomes still matter at the rule level
- **Why C is wrong:** Deontologists ground rules in duties and rights, not in utility maximization
- **Why D is wrong:** That describes act utilitarianism, the view rule consequentialism was designed to repair

EXAM TRAP: Rule consequentialism: consequences pick the RULES, rules (not acts) judge conduct, so critics say it ends up deontological!

Q2: Kant vs Contemporary Deontology - Advanced

QUESTION: An AI system could deceive users and produce large net benefits. Which analysis correctly contrasts Kant with most contemporary deontologists?

OPTIONS:

- ☐ A. Kant permits lying when benefits are large; contemporary deontologists forbid it absolutely
- ☐ B. Both ignore consequences entirely, which is what separates them from virtue ethicists
- ☐ C. Contemporary deontologists are act utilitarians who use rights only as tie-breakers
- ☐ D. Kant treats lying as wrong irrespective of consequences; most contemporary deontologists give consequences weight but hold that some red lines cannot be crossed even for good overall outcomes

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Kant places no moral weight on consequences. Contemporary deontologists care about consequences, but rights, rules, and principles act as constraints that go beyond them.
- **Why A is wrong:** This reverses the positions; Kant is the stricter view
- **Why B is wrong:** Most contemporary deontologists do grant consequences moral weight
- **Why C is wrong:** Consequences are one consideration for them, not the only one, so they are not utilitarians

EXAM TRAP: Kant = consequences carry no weight; modern deontologists = consequences count, but RED LINES remain!

Q3: Universalizability Test - Advanced

QUESTION: A firm adopts the maxim “we may mislead customers whenever it is commercially convenient.” Which reasoning applies the categorical imperative most accurately?

OPTIONS:

- ☐ A. It is rejected because average outcomes across firms would be worse in a utility calculation
- ☐ B. It is rejected because a virtuous person would feel ashamed of it
- ☐ C. It is acceptable as long as competitors do not adopt the same maxim
- ☒ D. The maxim cannot be willed as a universal law without contradiction: if everyone misled whenever convenient, trust in statements would collapse and the strategy would defeat itself

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** The test asks whether a maxim can be universalized without contradiction, regardless of consequences. Deception parasitically depends on a background practice of truth-telling that universal deception would destroy.
- **Why A is wrong:** Cost-benefit aggregation is consequentialist, not the Kantian test
- **Why B is wrong:** Asking what a virtuous agent would do is virtue ethics
- **Why C is wrong:** Exempting oneself from a rule that others follow is exactly what universalizability forbids

EXAM TRAP: Categorical imperative = could the maxim be willed as universal law WITHOUT CONTRADICTION?

Q4: Undercutting the Utilitarian Metric - Advanced

QUESTION: Utilitarianism is praised for providing a single, quantifiable, impartial metric. Which statement best explains why this advantage is undercut in practice?

OPTIONS:

- ☐ A. Utilitarianism is backward-looking, so it cannot evaluate future AI deployments
- ☐ B. Utilitarianism requires partiality toward those closest to the decision maker
- ☒ C. What counts as utility (pleasure, happiness, desire satisfaction) is subjective and hard to quantify, and consequences are unpredictable with no clear stopping point for tracing consequences of consequences
- ☐ D. Utility can only be defined as pleasure, which cannot be measured

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** The metric is single in form but contestable in content, and the scope of the calculation is open-ended, so apparent objectivity is fragile.
- **Why A is wrong:** Utilitarianism is forward-looking and focused on end consequences
- **Why B is wrong:** It aims to be impartial and objective, not partial
- **Why D is wrong:** Utility may also be defined as happiness or satisfaction of desires

EXAM TRAP: Utilitarianism's weakness: choice of metric is subjective AND where to stop the calculation is unclear!

Q5: Wrongs Without Harm - Advanced

QUESTION: The module says violating privacy is wrong even when it causes no harm, comparing it to cheating on a spouse who never finds out. Which foundation best supports that claim?

OPTIONS:

- ☒ A. Privacy violations can treat persons merely as means to the violator's ends and infringe moral and legal rights, so they wrong someone independently of tangible harm
- ☐ B. They are wrong only in expectation, because they raise the probability of later harm
- ☐ C. They are wrong only where a statute recognizes a right to privacy
- ☐ D. They are wrong because they reduce aggregate welfare

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** A wrong is a violation of rights or of a person's standing; harm is a separate notion. Treating people instrumentally, as things and not autonomous beings, is wrong even if nothing bad follows.
- **Why B is wrong:** This makes the wrong depend on harm, which the module explicitly denies
- **Why C is wrong:** The right to privacy is a moral right whether or not a law or era recognizes it
- **Why D is wrong:** Aggregate welfare is consequentialist reasoning and again depends on harm

EXAM TRAP: WRONGS can exist without HARM: rights violations and treating people as mere means!

Q6: Virtue Ethics Guidance Problem - Advanced

QUESTION: Critics say virtue ethics gives less guidance than duty-based or consequentialist approaches because virtues can conflict. What reply does the module give?

OPTIONS:

- ☐ A. Virtues never actually conflict once each is properly understood
- ☐ B. Virtue ethics supplies a fixed ranking of virtues that resolves conflicts mechanically
- ☐ C. Guidance is unnecessary because character is fixed in early childhood
- ☒ D. Practical wisdom helps navigate hard cases, and rival theories face parallel problems: duties can conflict and consequences are not always comparable

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Virtue ethicists claim they are no worse off: every framework must weigh competing considerations, and practical wisdom is the faculty for doing so in context.
- **Why A is wrong:** The module accepts that virtues can conflict; its point is that all theories face conflict
- **Why B is wrong:** There is no mechanical ranking; practical wisdom is exercised case by case
- **Why C is wrong:** Virtues are built through practice, habit, and exemplars over a lifetime

EXAM TRAP: Virtue ethics reply = PRACTICAL WISDOM, plus 'you have conflicts too' (duties clash, consequences incomparable)!

Q7: Virtue Ethics and Building AI - Advanced

QUESTION: Some scholars argue that to build ethical AI we must build it “in a virtue ethics way.” What is the argument?

OPTIONS:

- ☒ A. Morality is so complex that it may never be coded top-down, so systems should learn from experience and habituation, the way children acquire morality
- ☐ B. Virtue ethics provides an explicit rule set that can be encoded as hard constraints
- ☐ C. Virtue ethics equates goodness with total utility, which a loss function can optimize
- ☐ D. AI systems have emotions and appetites that must be moderated

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** The virtue approach is bottom-up: character is formed through practice, habit, and modeling exemplars, which parallels learning from experience.
- **Why B is wrong:** Rule sets are the deontological, top-down route the argument says may fail
- **Why C is wrong:** Utility maximization is consequentialist
- **Why D is wrong:** The point is about how morality is learned, not about machine emotions

EXAM TRAP: Virtue-ethics AI = learn from experience and habit (BOTTOM-UP), because top-down coding of morality may be impossible!

Q8: Three Lenses on Privacy Decisions - Advanced

QUESTION: Which mapping of the three frameworks onto decisions about collecting and keeping personal data matches the module?

OPTIONS:

- ☐ A. Consequentialist: the right to privacy. Deontological: worst-case scenarios. Virtue: data minimization
- ☐ B. Consequentialist: what would a responsible company do. Deontological: best/worst cases. Virtue: data minimization
- ☐ C. All three reduce to compliance with the applicable law
- ☒ D. Consequentialist: best/worst cases, data sensitivity, confidence in keeping it safe. Deontological: the right to privacy and principles like data minimization. Virtue: what would a responsible company do

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Consequentialism asks about outcomes and risk, deontology about rights and rules, and virtue ethics about the character of a responsible organization.
- **Why A is wrong:** Rights and principles are deontological, so the roles are swapped
- **Why B is wrong:** Character questions belong to virtue ethics, so the roles are scrambled
- **Why C is wrong:** The module says ethics goes beyond minimal legal requirements

EXAM TRAP: Consequences = risk and sensitivity; Deontology = rights and minimization; Virtue = what a responsible company would do!

Q9: Ethics and Law - Advanced

QUESTION: The module calls ethics “more ambitious than the law” and says it helps distinguish just from unjust laws. Which inference is consistent with this?

OPTIONS:

- ☐ A. Whatever is legal is ethically permissible, because law fully encodes social morality
- ☐ B. Ethics becomes redundant once laws like the GDPR exist
- ☒ C. Abandoning a friend in need is legal yet immoral; because ethics offers standards independent of statutes, it can ground, inform, and criticize law
- ☐ D. Law sets aspirational ideals of the good life while ethics sets minimum requirements

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Laws typically set minimal behavioral requirements so institutions function. Ethics asks what kind of society and life we want and therefore can shape and evaluate laws.
- **Why A is wrong:** Legal does not imply ethical; the friend example shows the gap
- **Why B is wrong:** The module says laws are also needed but do not replace ethics
- **Why D is wrong:** This reverses the roles of law (minimal) and ethics (ambitious)

EXAM TRAP: Law = MINIMUM requirements; ethics = more ambitious, and can judge whether laws are just!

Q10: Moral vs Medical Questions - Advanced

QUESTION: Why does the module use the ventilator and organ procurement from heart-beating bodies with non-functioning brains as a case for formal ethics structures?

OPTIONS:

- ☐ A. Clinicians lacked the technical skill to operate ventilators safely
- ☐ B. The case proved that consequentialism is the only workable framework in medicine
- ☒ C. Whether to take such organs is a moral, not medical, question, so resolving it should not rest solely on clinicians, whose expertise is maintaining health
- ☐ D. Regulators required hospitals to adopt AI ethics codes

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** New technology created dilemmas that medical training does not resolve. This drove ethics frameworks, committees, and codes, and is offered as a parallel for engineers and executives facing AI dilemmas.
- **Why A is wrong:** The issue was moral, not a shortfall of technical skill
- **Why B is wrong:** The module says frameworks are plural and complementary
- **Why D is wrong:** AI ethics codes are not part of the medical case

EXAM TRAP: Ventilator lesson: moral questions are not medical questions, so experts should not decide them alone!

Q11: Nonmaleficence and Risk Allocation - Advanced

QUESTION: Which scenario is **MOST** defensible under the module's account of nonmaleficence, which considers who decides and who bears the brunt of harm?

OPTIONS:

- ☒ A. Very risky clinical research on patients with a serious ailment who stand to benefit if it succeeds
- ☐ B. The same very risky research on healthy volunteers
- ☐ C. The same research on patients whose ailment the research cannot cure
- ☐ D. An algorithm whose risks fall on people who gain nothing, imposed by people facing no risk

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** Risk is more acceptable when imposed on people who stand to benefit. Nonmaleficence does not forbid all harm; it forbids unjustified risk that outweighs benefits.
- **Why B is wrong:** Healthy subjects bear risk without prospect of benefit
- **Why C is wrong:** No possible benefit to those bearing the risk
- **Why D is wrong:** This is the module's own analogy for an unacceptable algorithmic risk

EXAM TRAP: Nonmaleficence is not zero-risk: ask WHO benefits and WHO bears the risk!

Q12: Unjustified Risk Imposition - Advanced

QUESTION: A firm deploys an untested algorithm that could harm users. No harm has yet occurred. Under nonmaleficence:

OPTIONS:

- ☒ A. The firm may already be acting wrongly, because imposing unjustified or unnecessary risk on others is itself covered by the obligation
- ☐ B. No violation exists until measurable harm materializes
- ☐ C. Deployment is acceptable provided expected profit exceeds expected harm to the firm
- ☐ D. The principle applies only to intentional injury

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** The module says nonmaleficence includes a prohibition on imposing risks of harm that are not justified or that outweigh benefits, for example subjecting people to untested harmful algorithms.
- **Why B is wrong:** The obligation reaches risk imposition, not only realized harm
- **Why C is wrong:** Justification must weigh risks to those exposed, not only to the firm
- **Why D is wrong:** The module includes unnecessary and unjustified risk, not only intentional injury

EXAM TRAP: Nonmaleficence covers RISK IMPOSITION, not only realized harm!

Q13: Beneficence Is Not Merely Consequentialist - Advanced

QUESTION: Which statement about beneficence is correct according to the module?

OPTIONS:

- ☐ A. It is purely a consequentialist idea with no place in duty-based ethics
- ☐ B. It is unlimited: each agent must do everything possible to relieve suffering
- ☒ C. Deontological frameworks also include beneficence as a duty above what utility maximization requires, but it is limited by scarce resources, competing obligations, reasonableness, and demandingness
- ☐ D. It is equivalent to nonmaleficence stated positively

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Beneficence requires positive steps to help others, a key deontological duty, but ethics cannot demand unlimited sacrifice.
- **Why A is wrong:** This is the common misunderstanding the module corrects
- **Why B is wrong:** No individual can alleviate all suffering; reasonable constraints apply
- **Why D is wrong:** Nonmaleficence forbids injuring; beneficence commands positive help

EXAM TRAP: Beneficence is ALSO a deontological duty, bounded by demandingness and reasonableness!

Q14: Beneficence vs Autonomy - Advanced

QUESTION: A wellness app decides to override a competent adult's declared choice because it computes that another option would benefit them. What does the module imply?

OPTIONS:

- ☒ A. Autonomy usually trumps unwanted benefits, with exceptions such as public health worries or doubts about the person's decision-making capacity
- ☐ B. Beneficence always outranks autonomy when the benefit is large
- ☐ C. Autonomy is absolute and can never be overridden, even for children
- ☐ D. Surrogate decision makers are needed for every adult decision

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** People have a right to decide what is best for them. Only public-health concerns or worries about capacity typically justify overriding, and then a surrogate acts in best interests.
- **Why B is wrong:** Unwanted benefits that disrespect agency are precluded by the module
- **Why C is wrong:** The module allows surrogates for those lacking capacity (children, unconscious patients)
- **Why D is wrong:** Surrogates are for people lacking capacity, not for every adult

EXAM TRAP: Autonomy usually TRUMPS unwanted beneficence; exceptions: public health and lack of capacity!

Q15: Autonomy and Manipulative Design - Advanced

QUESTION: Which system **MOST** clearly contravenes autonomy as defined in the module (informed, un-coerced choice; technology should advance the user's goals, not third parties')?

OPTIONS:

- ☐ A. A decision aid that shows its recommendation and lets users freely override it
- ☒ B. A recommender that exploits psychological weaknesses to steer purchases benefiting the platform, even though users technically click accept
- ☐ C. A consent flow that discloses data uses and allows withdrawal at any time
- ☐ D. A guardian-facing tool that helps a legal guardian act in a child's best interests

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Coercion, deception, manipulation, and undue influence all undermine autonomy. Nominal clicking does not make manipulated choices autonomous.
- **Why A is wrong:** Overridable recommendations preserve choice
- **Why C is wrong:** Meaningful consent and withdrawal support autonomy
- **Why D is wrong:** Surrogate decision-making for those lacking capacity is a recognized exception

EXAM TRAP: Autonomy is undermined by manipulation even when the user 'agrees'!

Q16: Conceptions of Justice - Advanced

QUESTION: A firm remedies harm from a biased hiring tool by meeting rejected applicants, acknowledging the harm, and compensating them. Which conception of justice is this?

OPTIONS:

- ☐ A. Procedural justice: fair, impartial processes
- ☒ B. Restorative justice: repairing harms by reconciling victims and offenders
- ☐ C. Distributive justice: equitable allocation of benefits and burdens
- ☐ D. Interactional justice: respect and fairness in dealings between individuals, with no repair element

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Restorative justice aims to repair harm. The scenario centers on acknowledgment and remedy after harm has occurred.
- **Why A is wrong:** Procedural concerns the fairness of the decision process itself
- **Why C is wrong:** Distributive concerns allocation, not repair
- **Why D is wrong:** Interactional concerns respectful treatment; the key here is repair

EXAM TRAP: Restorative = REPAIR harm; procedural = fair process; distributive = allocation; interactional = respect; social = just institutions!

Q17: Competing Principles of Justice - Advanced

QUESTION: One triage algorithm allocates scarce resources by need, another by merit. Both claim to be just. Which statement is most accurate?

OPTIONS:

- ☒ A. They reflect competing principles of justice (need, meritocracy, egalitarianism, utilitarianism) that can conflict, so legitimacy depends on decisions being justifiable to all, not on proving one principle correct
- ☐ B. Only need-based allocation counts as justice; merit-based allocation is discrimination
- ☐ C. Justice is purely procedural, so allocation principles are irrelevant
- ☐ D. Broad agreement exists on which goods justice must distribute

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** The module notes disagreement about what a just distribution is and which goods justice concerns. It notes broad agreement on discrimination for morally irrelevant traits.
- **Why B is wrong:** Merit is a recognized justice principle; it is not automatically discrimination
- **Why C is wrong:** Distributive justice concerns allocation and is a distinct conception
- **Why D is wrong:** The module says people disagree on the goods (resources, opportunities, power, welfare)

EXAM TRAP: Justice principles COMPETE; legitimacy requires decisions justifiable to all!

Q18: Local and Model-Agnostic Explainability - Advanced

QUESTION: A regulator wants to know why one applicant was denied by a black-box ensemble the provider will not restructure. Which type of explanation fits best?

OPTIONS:

- ☐ A. A global explanation using rules tailored to decision-tree architecture
- ☒ B. A local, model-agnostic explanation of that specific decision
- ☐ C. A global, model-agnostic summary of behavior across all inputs
- ☐ D. No explanation can exist for a black-box model

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Local explainability targets a single prediction. Model-agnostic methods work with any architecture, while model-specific techniques depend on the model type.
- **Why A is wrong:** Global, and tree-specific rules do not apply to an ensemble
- **Why C is wrong:** Global summaries do not explain one case
- **Why D is wrong:** Explanation methods exist, and auditing tools are being developed for black boxes

EXAM TRAP: Single decision = LOCAL; any architecture = MODEL-AGNOSTIC!

Q19: Counterfactual Explanations and Audience - Advanced

QUESTION: Which statement best captures what counterfactual explanations offer and the caveat about explanations generally?

OPTIONS:

- ☐ A. They describe overall behavior across all inputs; one explanation fits every audience
- ☐ B. They prove the model is unbiased; the audience is irrelevant
- ☒ C. They state conditions under which the outcome would have differed (for example, \$10,000 more savings), giving actionable insight; what counts as an explanation depends on to whom it is owed
- ☐ D. They disclose the training dataset; only experts need explanations

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Counterfactuals contrast the actual decision with a hypothetical one. The module stresses explanations for experts differ from those for regulators or ordinary citizens.
- **Why A is wrong:** That describes global explainability
- **Why B is wrong:** They illuminate variable influence, not the absence of bias
- **Why D is wrong:** They concern decision conditions, not disclosure of training data

EXAM TRAP: Counterfactual = 'what would have changed the decision?'; explanations depend on the AUDIENCE!

Q20: Bias Is Not Unfairness - Advanced

QUESTION: Which statement is consistent with the module's view of algorithmic bias?

OPTIONS:

- ☒ A. Some biases are justifiable, and a statistically or legally unbiased algorithm can still be unethical (for example, one that charges exorbitant fees to everyone)
- ☐ B. Any deviation from parity is unethical, so an unbiased algorithm is always ethical
- ☐ C. A well-designed AI should be entirely objective and value-free
- ☐ D. Bias arises only from prejudiced designers

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** Bias is deviation from a desired norm or intended function. The ethical concern is unjustifiable disadvantage; every algorithm embeds values, so the goal is avoiding problematic bias, not all bias.
- **Why B is wrong:** The fee example shows unbiased is not the same as ethical
- **Why C is wrong:** The module says no algorithm is objective; each reflects embedded values
- **Why D is wrong:** Problematic biases often occur unintentionally due to complexity

EXAM TRAP: Not all bias is problematic, and not all unbiased AI is ethical!

Q21: Problem Specification Bias - Advanced

QUESTION: A bank wants zero lending risk and its algorithm ends up serving only affluent applicants. Which diagnosis and remedy fit best?

OPTIONS:

- ☐ A. Historical bias; add more decades of lending data
- ☒ B. Bias in problem specification: the goal is operationalized in a way that misses the real-world objective, so the target should be modified to accommodate reasonable risk
- ☐ C. Sampling bias; oversample low-income applicants only
- ☐ D. Deployment bias; retrain loan officers

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** The flaw lies in the goal itself, not the data or its use. Zero-risk targets exclude everyone who cannot guarantee repayment.
- **Why A is wrong:** More historical data would reinforce the same target
- **Why C is wrong:** The issue is the objective, not the sample
- **Why D is wrong:** The problem exists before deployment

EXAM TRAP: Bias can be born in the GOAL: fix the target, not just the data!

Q22: Proxy Targets: The Healthcare Example - Advanced

QUESTION: An algorithm meant to identify the sickest patients used healthcare expenditure as a proxy and favored wealthier patients. Which mechanism explains this?

OPTIONS:

- ☐ A. Selective labels: outcomes were unobserved for patients not treated
- ☐ B. Automation bias: clinicians over-trusted the tool
- ☐ C. Sampling bias: too few patients were sampled
- ☒ D. The proxy diverges from the real objective (illness) because spending reflects access to care and wealth, so optimizing it selects the richest rather than the sickest

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** This is problem-specification bias through a flawed proxy: the measurable stand-in encodes unequal spending patterns.
- **Why A is wrong:** Selective labels concerns missing counterfactual outcomes, not the choice of target
- **Why B is wrong:** Automation bias is a deployment issue about deference
- **Why C is wrong:** The failure exists even if the sample is large and random

EXAM TRAP: Cost is not need: a bad PROXY target creates biased outcomes!

Q23: Historical Bias Beyond Discrimination - Advanced

QUESTION: A model trained on 2010-2019 data forecasts office vacancy and commuter rail volume for 2025-2030. Which statement fits the module's discussion of historical data?

OPTIONS:

- ☐ A. Historical bias exists only when data reflect sex or race disparities
- ☐ B. Historical data are reliable if they are large and accurate for their period
- ☒ C. Historical data can mislead even without demographic bias: patterns that held in the past may not hold, producing inaccurate or malfunctioning models
- ☐ D. The problem is deployment bias created by human deference

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** The module extends historical bias beyond sex and race to mismatch between past and future real-world trends, illustrated by pre-pandemic versus later patterns.
- **Why A is wrong:** The module explicitly says it also leads to inaccurate models more broadly
- **Why B is wrong:** Accuracy for the past period does not guarantee validity for a changed future
- **Why D is wrong:** No human deference is involved in this example

EXAM TRAP: Historical bias includes DATA THAT NO LONGER MATCHES THE WORLD!

Q24: Selective Labels vs Sampling Bias - Advanced

QUESTION: Which statement correctly distinguishes selective labels from sampling bias?

OPTIONS:

- ☐ A. Selective labels: outcomes are observed only for people selected by past decisions, so counterfactuals are missing; sampling bias: the sample is not random, so trends may not generalize
- ☐ B. They are the same problem described in two fields
- ☐ C. Selective labels: data drawn from too few people; sampling bias: missing outcomes for those denied
- ☐ D. Selective labels arise in deployment; sampling bias arises in problem specification

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** A lender sees repayment only for approved borrowers, so it never learns whether rejected ones would have repaid. Sampling bias, such as drug data mostly from young men, limits generalization.
- **Why B is wrong:** The module calls them related but distinct
- **Why C is wrong:** This swaps the definitions
- **Why D is wrong:** Both are data biases, not deployment or problem-specification biases

EXAM TRAP: Selective labels = missing COUNTERFACTUAL outcomes; sampling bias = NON-RANDOM sample!

Q25: Popularity Feedback Loops - Advanced

QUESTION: A search engine for financial products favors popular items, giving them more exposure and thereby more popularity, regardless of quality. Which category is this?

OPTIONS:

-  Bias in problem specification
-  Bias in modeling, validation, and algorithm design, which can arise even when specification and data are sound
-  Selective labels bias
-  Bias in deployment caused by human deference

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Choices about optimization functions, subgroups, and presentation can introduce bias. Here the design creates a self-reinforcing exposure loop.
- **Why A is wrong:** The goal (helping users choose) is not the flawed part; the design is
- **Why C is wrong:** Selective labels concerns missing outcomes for those not selected
- **Why D is wrong:** No human deference is described

EXAM TRAP: Design choices can create FEEDBACK LOOPS even with good goals and data!

Q26: Deployment Bias and Deference - Advanced

QUESTION: A bank warns employees not to rely solely on a loan-risk algorithm still under testing, yet many follow it uncritically. Which statement is accurate?

OPTIONS:

- ☐ A. Not a bias issue, because the algorithm itself is unbiased
- ☐ B. It is historical bias, fixable by retraining on more data
- ☐ C. It is caused by algorithms being more accurate than people
- ☒ D. This is deployment bias: an unbiased algorithm can still yield problematic outcomes because deference shields people from blame, so implementation changes behavior

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** The module lists hypotheses: convenience, perceived objectivity, self-doubt, and sharing of responsibility. Deferring lowers the risk of harsher judgment when errors occur.
- **Why A is wrong:** The module says biases can emerge at deployment despite an unbiased algorithm
- **Why B is wrong:** Retraining does not change the incentive to defer
- **Why C is wrong:** The account rests on incentives, not on proven accuracy

EXAM TRAP: Deployment bias: unintended INCENTIVE EFFECTS make people defer to the algorithm!

Q27: The Fairness Impossibility Result - Advanced

QUESTION: According to the module, which set of fairness criteria cannot be satisfied simultaneously when base rates differ across groups?

OPTIONS:

- ☒ A. Procedural fairness, distributive justice, and restorative justice
- ☐ B. Group fairness and privacy, unless the data are synthetic
- ☐ C. Individual fairness and explainability
- ☐ D. Demographic parity, predictive rate parity, and equal opportunity

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Automating fairness becomes feasible only when base rates are equal, which is rare. Cited work includes Kleinberg, Mullainathan, and Raghavan (2016) and Chouldechova (2017).
- **Why A is wrong:** Those are conceptions of justice, not statistical criteria
- **Why B is wrong:** Privacy and synthetic data are unrelated to the result
- **Why C is wrong:** The result concerns three statistical fairness criteria

EXAM TRAP: Unequal base rates = cannot satisfy demographic parity, predictive parity, and equal opportunity at once!

Q28: Can Fairness Auditing Be Automated? - Advanced

QUESTION: The module argues that AI cannot solve the fairness problems it creates. Which chain of reasoning is correct?

OPTIONS:

- ☒ A. Fairness cannot be automated because it involves moral trade-offs, so auditing for fairness cannot be fully automated either; tools help but rarely solve bias on their own
- ☐ B. Fairness is mathematical, but tools are too slow
- ☐ C. Auditing tools are unreliable, so audits should be abandoned
- ☐ D. AI-based auditors avoid the problem because they lack values

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** Fairness is a moral or ethical judgment, not a mathematical one, and involves imperfect compromises that change with circumstances. Audits also need a fresh, diverse human look.
- **Why B is wrong:** The module says fairness is not purely mathematical
- **Why C is wrong:** Auditing is called likely the best way to find and correct bias
- **Why D is wrong:** The module says every algorithm reflects values

EXAM TRAP: Fairness = moral judgment, so tools help but the audit cannot be fully automated!

Q29: Risk Scores and Unequal Base Rates - Advanced

QUESTION: Most people committing a certain crime are men. An AI scores a specific man and woman. Which conclusion follows from the module?

OPTIONS:

- ☐ A. Equal scores are always fair because individuals should be treated alike
- ☐ B. A higher score for the man is always fair because it is statistically justified
- ☒ C. Whatever scoring rule is chosen, one person can claim unfair treatment, showing why fairness is a moral trade-off rather than a computation
- ☐ D. Fairness disputes vanish once the model is accurate

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** A higher score for the man invites his complaint; a similar score invites the woman's complaint because her group's base rate is lower.
- **Why A is wrong:** Equal scores draw the woman's objection
- **Why B is wrong:** A higher score draws the man's objection
- **Why D is wrong:** Accuracy does not resolve conflicts among fairness criteria

EXAM TRAP: With unequal base rates, every choice leaves someone with a fairness complaint!

Q30: Procedural Fairness and Redress - Advanced

QUESTION: In the justice-system analogy, why can errors be justifiable under procedural fairness? Which AI analogue follows?

OPTIONS:

- ☐ A. Outcomes are always correct when the process is fair
- ☐ B. Only outcome fairness matters, so process is secondary
- ☐ C. Errors are never justifiable, so committees are unnecessary
- ☒ D. A fair, impartial process with means of redress makes errors defensible and correctable; firms need structures such as ethics committees weighing all three frameworks

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Mistakes can occur even with fair procedures, but evidence-based process and mechanisms like new trials allow redress. The challenge is to build corporate structures that deliver both procedural and outcome fairness.
- **Why A is wrong:** Mistakes still happen with fair procedures
- **Why B is wrong:** The module says fairness is not only about outcome
- **Why C is wrong:** The module holds that committees can support both kinds of fairness

EXAM TRAP: Procedural fairness = impartial process plus ability to RIGHT WRONGS!

Q31: Synthetic Data Trade-offs - Advanced

QUESTION: Which statement about synthetic data as a bias-mitigation tool is accurate?

OPTIONS:

- ☐ A. It always outperforms real data, so validation is unnecessary
- ☒ B. It avoids privacy risk to real data subjects and can be designed to remove bias, but it can be less precise (especially visual data) and differ from real data in non-obvious ways
- ☐ C. It is bias-free automatically because it is fabricated
- ☐ D. It eliminates privacy and quality concerns entirely

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Synthetic data are fabricated datasets. The advantage comes from design choices, and the disadvantage is fidelity.
- **Why A is wrong:** The module says precision can be lower
- **Why C is wrong:** Bias must be removed by design, not automatically
- **Why D is wrong:** Quality concerns remain

EXAM TRAP: Synthetic data: privacy and bias-control benefits, but FIDELITY risk!

Q32: Self-Driving Trade-offs - Advanced

QUESTION: A self-driving system is tuned to be especially sensitive to children, slightly reducing adult detection accuracy. Which pairing is correct?

OPTIONS:

- ☒ A. The aim is avoiding problematic biases, not all bias; a deontological check is that protected categories are not disfavored (for example, better detection of larger men) and safety minimums hold
- ☐ B. Any accuracy trade-off is unethical because bias must be eliminated
- ☐ C. A consequentialist check: ensure no protected category is disfavored
- ☐ D. A virtue check: compute accident risk for each version

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** Trade-offs must be justifiable to stakeholders. Consequentialists compute accident risk, deontologists guard protected categories and minimums, and virtue ethicists focus on responsible processes.
- **Why B is wrong:** The module says the goal is not to avoid bias in general
- **Why C is wrong:** Protected categories and safety minimums are deontological
- **Why D is wrong:** Risk calculation is consequentialist

EXAM TRAP: Avoid PROBLEMATIC bias; consequences = risk math, deontology = protected groups and minimums, virtue = responsible process!

Q33: Bias vs Legal Discrimination - Advanced

QUESTION: An algorithm disfavors applicants on a characteristic that is irrelevant to the task but not legally protected in the jurisdiction. What is the best assessment?

OPTIONS:

- ☐ A. It is neither, since only protected classes matter
- ☐ B. It is automatically illegal discrimination everywhere
- ☒ C. It is a problematic bias ethically, because it deviates from function based on a morally irrelevant trait, although whether it is legal discrimination depends on jurisdiction and protected classes
- ☐ D. It is acceptable because it is statistically accurate

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Legal discrimination is likely when disfavoring people by race, sex, ethnicity, age, or another protected class. Ethical fairness concerns irrelevant characteristics more broadly.
- **Why A is wrong:** Ethics reaches beyond legal minimums
- **Why B is wrong:** Discrimination law varies by jurisdiction
- **Why D is wrong:** Statistical unbiasedness does not guarantee ethical acceptability

EXAM TRAP: Ethics goes beyond law: legal discrimination varies by JURISDICTION!

Q34: Access-Based Privacy - Advanced

QUESTION: The module defines privacy as others having no access to a personal data point. Under this definition, which event is a privacy loss?

OPTIONS:

- ☐ A. Only a public leak of the data
- ☐ B. Only use of the data for discrimination
- ☐ C. Only access that violates a statute
- ☐ D. An employer gains access to a personal data point, even if it never misuses it

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Privacy is proportional to lack of access, and it protects from abuses of power because more knowledge makes interference easier.
- **Why A is wrong:** Loss occurs on access, not only on publication
- **Why B is wrong:** Misuse is a possible harm, not the definition
- **Why C is wrong:** The moral right holds whether or not the law recognizes it

EXAM TRAP: Privacy = lack of ACCESS; loss does not require misuse!

Q35: Personal Data as a Toxic Asset - Advanced

QUESTION: Which argument best supports data minimization on the module's terms?

OPTIONS:

- ☐ A. Minimization mainly lowers storage costs
- ☐ B. Minimization applies only to regulated data
- ☒ C. Data are cheap to mine and valuable but sensitive and hard to keep safe; each data point is a potential lawsuit or fine, and risk grows with amount, storage time, analysis, and sharing
- ☐ D. Data are risk-free if encrypted

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Data are a toxic yet valuable asset. Collect only when the benefit to data subjects outweighs disadvantages.
- **Why A is wrong:** Cost is secondary to risk and rights
- **Why B is wrong:** The principle is ethical and general
- **Why D is wrong:** Security cannot eliminate risk given attacker advantages

EXAM TRAP: More data collected, kept, analyzed, and shared = more risk!

Q36: Attacker-Defender Asymmetry - Advanced

QUESTION: In cyberspace, defenders are at a disadvantage. What follows for a firm holding excess personal data?

OPTIONS:

- ☐ A. Attacks are avoidable with enough encryption
- ☐ B. Only breaches that occur are ethically relevant
- ☒ C. If an organization cannot keep data safe it puts itself at risk by collecting it, and imposing unjustified breach risk on data subjects can itself be a wrong
- ☐ D. Risk is borne only by the firm

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Attackers choose time and method; defenders must guard against all. Risk imposition can amount to a wrong even if the attempt fails.
- **Why A is wrong:** A resourced, motivated attacker will eventually succeed
- **Why B is wrong:** Imposing risk can be wrong without a completed harm
- **Why D is wrong:** Risk also falls on data subjects and society

EXAM TRAP: Defenders must stop EVERY attack; collecting unnecessary data imposes risk!

Q37: Contextual Integrity - Advanced

QUESTION: A bank stores customer data with strong encryption but sells it to a marketing company. Which principle is violated most directly?

OPTIONS:

- ☐ A. Data deletion: retention beyond need
- ☐ B. Right to be forgotten: removal from search results
- ☐ C. Data security: the encryption was inadequate
- ☐ D. Contextual integrity: data given in one context should not be moved to another that violates its norms

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** The customer provided data for a transaction. Moving it to marketing violates the norms of that context even when the system is secure.
- **Why A is wrong:** Deletion concerns timing, not repurposing
- **Why B is wrong:** That right concerns inadequate, irrelevant, or excessive information in searches and directories
- **Why C is wrong:** The scenario stipulates strong encryption

EXAM TRAP: Secure is not the same as appropriate: contextual integrity governs USE across contexts!

Q38: Right to Be Forgotten - Advanced

QUESTION: When does the right to be forgotten apply, according to the module?

OPTIONS:

- ☐ A Whenever a person finds information embarrassing
- ☒ B When information is inadequate (for example incorrect), irrelevant or no longer relevant, or excessive relative to its purposes, making it less accessible in searches and directories
- ☐ C Only for criminal records
- ☐ D Only when data was obtained unlawfully

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** It concerns removal from internet searches and other directories under specified conditions.
- **Why A is wrong:** Embarrassment alone is not the stated criterion
- **Why C is wrong:** It is not limited to criminal records
- **Why D is wrong:** Lawful data may still become irrelevant

EXAM TRAP: Right to be forgotten = inadequate, irrelevant, or excessive data!

Q39: Control Over Data - Advanced

QUESTION: Which element of control over data goes beyond the raw personal data supplied by individuals?

OPTIONS:

- ☐ A. Ability to withdraw consent
- ☐ B. Ability to delete their data
- ☒ C. Access to the inferences made about them
- ☐ D. Ability to transfer their data

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Control includes denying or withdrawing consent, accessing data and inferences, transfer, and deletion. Inferences are derived rather than supplied.
- **Why A is wrong:** Withdrawal concerns consent to raw collection and processing
- **Why B is wrong:** Deletion concerns stored data
- **Why D is wrong:** Portability concerns moving data

EXAM TRAP: Control includes access to **INFERENCES**, not only raw data!

Q40: Backpropagation and Interpretability - Advanced

QUESTION: Why does training by backpropagation contribute to the interpretability problem?

OPTIONS:

- ☐ A. It randomizes weights so outputs are unpredictable
- ☒ B. It adjusts weights to minimize output error, improving performance without revealing how decisions are reached, and scale (millions or billions of parameters) makes tracing difficult
- ☐ C. It keeps model weights secret from developers
- ☐ D. It prevents networks from using proxies

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Effectiveness at reducing error does not provide insight into reasoning. Auditing methods are being developed to look into the black box.
- **Why A is wrong:** Backpropagation is systematic error minimization
- **Why C is wrong:** The problem persists even when weights are visible
- **Why D is wrong:** Proxies are common, and unknown proxies hamper governance

EXAM TRAP: Backpropagation optimizes error, NOT understanding!

Q41: Unpredictability, RCTs, and Audits - Advanced

QUESTION: Why does the module suggest both randomized controlled trials and periodic audits for unpredictable AI?

OPTIONS:

- ☒ A. Emergent behavior and unforeseen human interactions mean pre-deployment trials alone cannot capture all failure modes, so audits are a complementary check on safety and accuracy
- ☐ B. RCTs eliminate unpredictability, so audits are redundant
- ☐ C. Audits replace RCTs because trials are unethical
- ☐ D. Unpredictability arises only from user misuse

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** Neural networks exhibit emergent properties not coded explicitly and can surprise experts. Misuse is also hard to predict.
- **Why B is wrong:** RCTs minimize but do not eliminate risk
- **Why C is wrong:** The module calls audits complementary to RCTs
- **Why D is wrong:** Unpredictability also arises from emergent capabilities and errors

EXAM TRAP: Emergent behavior: RCTs AND periodic audits!

Q42: LLMs and Truth-Tracking - Advanced

QUESTION: Why are LLM outputs risky even when fluent?

OPTIONS:

- ☐ A. They deliberately deceive users
- ☐ B. They rely on verified knowledge bases
- ☐ C. They are unregulated only in the EU
- ☒ D. They make statistical inferences designed to be plausible, and plausible is not the same as truthful, creating risks of misinformation, physical safety issues, and libel

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** LLMs are not based on an understanding of truth or knowledge of the world, and false responses can still be convincing.
- **Why A is wrong:** The issue is design, not intent
- **Why B is wrong:** The module says they are not truth-tracking
- **Why C is wrong:** The concern is global

EXAM TRAP: Plausible does not mean true!

Q43: LLMs and GDPR Rights - Advanced

QUESTION: Why does the module question whether LLM providers can comply with the GDPR?

OPTIONS:

- ☐ A. The GDPR does not apply to scraped data
- ☒ B. Data subjects may ask what data a company holds, modify it, and delete it, and it is unclear providers can meet such requests for data scraped and absorbed into models
- ☐ C. LLMs never process personal data
- ☐ D. The GDPR bans language models

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** Scraping raises privacy and copyright concerns, and lawsuits were pending at the time of writing.
- **Why A is wrong:** The GDPR applies to personal data generally
- **Why C is wrong:** The module raises scraped personal data as an issue
- **Why D is wrong:** The GDPR is not a ban on any technology

EXAM TRAP: GDPR rights (access, correct, delete) are hard to honor for scraped LLM training data!

Q44: GDPR Reach and Breach Notification - Advanced

QUESTION: A US company with no EU office offers a service to people in the EEA and suffers a breach. What applies?

OPTIONS:

- ☐ A. The GDPR does not apply outside the EEA
- ☐ B. Notification is due within 45 days
- ☐ C. Notification is due only if individuals sue
- ☒ D. The GDPR applies because the company offers services to EEA data subjects; the controller must notify the supervisory authority within 72 hours

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** Extraterritorial reach has made the GDPR globally influential. Controllers must notify within 72 hours.
- **Why A is wrong:** It applies based on offering services to EEA data subjects
- **Why B is wrong:** 45 days is the appendix figure for the Illinois BIPA
- **Why C is wrong:** The duty is legal and does not depend on suits

EXAM TRAP: GDPR: extraterritorial, 72-hour breach notice!

Q45: Legitimate Interest - Advanced

QUESTION: Under the GDPR, when may the right to object to processing not apply?

OPTIONS:

- ☐ A. Whenever the controller is outside the EEA
- ☐ B. Whenever data are anonymized
- ☒ C. When there is a legitimate interest to carry out the service, whose scope is debated and being clarified by rulings
- ☐ D. Never; the right is absolute

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** The right to object covers marketing or purposes unrelated to the service, but does not apply where legitimate interest exists.
- **Why A is wrong:** Location does not remove GDPR duties when serving EEA data subjects
- **Why B is wrong:** Anonymization is unrelated to this exception
- **Why D is wrong:** The module lists legitimate interest as a limit

EXAM TRAP: Right to object has a LEGITIMATE-INTEREST limit!

Q46: EU AI Act: Prohibited vs Regulated - Advanced

QUESTION: Which system is among those prohibited under the EU AI Act, as opposed to merely regulated?

OPTIONS:

- ☐ A. An AI system that screens job applicants
- ☐ B. A chatbot that interacts with customers
- ☒ C. Untargeted scraping of facial images from the internet or CCTV to build facial recognition databases
- ☐ D. A generator of deepfake content

ANSWER: C

DETAILED EXPLANATION:

- **Why C is correct:** Prohibited systems include biometric categorization using sensitive traits, untargeted facial scraping, workplace and educational emotion recognition, social scoring, manipulative systems, and exploiting vulnerabilities.
- **Why A is wrong:** Job screening carries obligations such as preventing undue harm to opportunities
- **Why B is wrong:** Chatbots require notification
- **Why D is wrong:** Deepfakes require labeling

EXAM TRAP: Facial scraping banned; job-screening AI, chatbots and deepfakes regulated!

Q47: EU AI Act Penalties and Timeline - Advanced

QUESTION: Which pairing matches the module's account of AI Act penalties and timing?

OPTIONS:

- ☒ A. Fines range from EUR 35 million or 7% of global turnover down to EUR 7.5 million or 1.5%; prohibitions applied from February 2, 2025 and most provisions from August 2, 2026
- ☐ B. Fines are capped at EUR 1 million; effective January 1, 2020
- ☐ C. Fines equal 72 hours of turnover; all provisions applied on August 1, 2024
- ☐ D. There are no fines; the AI Office only issues guidance

ANSWER: A

DETAILED EXPLANATION:

- **Why A is correct:** The Act entered into force August 1, 2024 and created the European AI Office for compliance and enforcement.
- **Why B is wrong:** This misstates both figures and dates
- **Why C is wrong:** 72 hours is the GDPR breach window
- **Why D is wrong:** The AI Office enforces binding rules

EXAM TRAP: AI Act: up to EUR 35M or 7%; bans Feb 2, 2025; most provisions Aug 2, 2026!

Q48: High-Risk Obligations - Advanced

QUESTION: Which allocation of high-risk AI obligations is accurate?

OPTIONS:

- ☐ A. Only chatbot operators must assess rights impact
- ☐ B. Providers need not keep records if they publish source code
- ☐ C. Only public agencies must document training
- ☒ D. Banks, insurers, and deployers of high-risk AI do rights impact assessments; providers keep records on datasets, training methods, and oversight measures

ANSWER: D

DETAILED EXPLANATION:

- **Why D is correct:** The Act imposes assessments on essential-service organizations and record-keeping duties on providers.
- **Why A is wrong:** Chatbots have notification duties instead
- **Why B is wrong:** No source-code exemption is described
- **Why C is wrong:** The duties are not limited to public agencies

EXAM TRAP: High-risk: impact assessments plus thorough records!

Q49: DMA vs DSA - Advanced

QUESTION: A platform ranks its own products above rivals; another very large platform must let users opt out of personalized recommendations. Which laws apply?

OPTIONS:

- ☐ A. Both are DMA obligations
- ☒ B. The DMA bans self-preferencing by gatekeepers; the DSA requires very large platforms to allow opt-out of personalized content
- ☐ C. The GDPR bans self-preferencing; the AI Act covers opt-outs
- ☐ D. Both are DSA obligations on content moderation only

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** The DMA targets gatekeepers (self-preferencing, data access, no tying, messaging interoperability). The DSA covers illegal content, transparency, disinformation, and algorithmic personalization.
- **Why A is wrong:** Opt-out of personalization is a DSA provision
- **Why C is wrong:** The GDPR concerns personal data and the AI Act concerns AI risk
- **Why D is wrong:** Self-preferencing is a DMA matter

EXAM TRAP: DMA = gatekeepers and competition; DSA = platform content and transparency!

Q50: US vs China Approaches - Advanced

QUESTION: Which comparison of AI governance in the US and China is accurate?

OPTIONS:

- ☐ A. The US has a comprehensive federal AI law; China relies on voluntary guidance
- ☒ B. The US has no federal AI law but uses an executive order and the voluntary NIST AI RMF (govern, map, measure, manage); China has binding rules such as the 2023 Generative AI Measures, PIPL, and algorithm recommendation provisions
- ☐ C. Both rely solely on voluntary frameworks
- ☐ D. NIST AI RMF is mandatory certification

ANSWER: B

DETAILED EXPLANATION:

- **Why B is correct:** The RMF is flexible and non-prescriptive. China's measures impose duties on generative AI providers and mandate data minimization, consent, and transparency.
- **Why A is wrong:** The module says the US has no federal AI law
- **Why C is wrong:** China has enacted binding rules
- **Why D is wrong:** The RMF is voluntary

EXAM TRAP: US: EO plus voluntary RMF; China: binding rules!