

GARP Risk & AI — Q351–Q450

Unofficial Exam Prep

Unofficial

October 3, 2026

Question

A global bank establishes a data governance committee. Which of the following best describes the committee's primary role?

- A. Operationally entering data into the warehouse.
- B. Providing oversight on data strategy, policies, quality, access, and compliance, without performing operational data work.
- C. Building machine learning models.
- D. Marketing data products to external clients.

Answer: B**Explanation**

The data governance committee provides oversight — approving policies, resolving escalations, and ensuring compliance. It does not build models, enter data, or market products. Options A, C, and D describe operational roles.

Question

What is the board's most important responsibility regarding data governance in a financial institution?

- A. Writing data-cleaning code.
- B. Approving the data governance framework and overseeing alignment with strategic objectives, regulatory obligations, and risk appetite.
- C. Performing ETL jobs.
- D. Selecting the firm's cloud vendor.

Answer: B

Explanation

The board approves the framework and provides strategic oversight. Options A, C, and D are operational or IT-level decisions.

Question

Which of the following is NOT typically considered a dimension of data quality?

- A. Accuracy.
- B. Consistency.
- C. Integrity.
- D. Profitability.

Answer: D

Explanation

Data quality dimensions include accuracy, consistency, integrity, completeness, timeliness, and validity. Profitability is a business metric, not a data quality dimension.

Question

A model is deployed using vendor-provided data, but the vendor never disclosed that the data was scraped without permission from a third-party website. Which governance failure has occurred?

- A. Model accuracy failure.
- B. Data provenance failure, exposing the bank to legal and regulatory risk.
- C. Training-time failure.
- D. Data storage failure.

Answer: B**Explanation**

Provenance failures involve unclear legal rights to data. Regulators may require the model to be withdrawn or destroyed. Options A, C, and D are secondary or incorrect.

Question

A dataset contains customer names, addresses, and publicly available property valuations. Which classification is most appropriate?

- A. Public for all three fields.
- B. Confidential or restricted for names and addresses; public for property valuations.
- C. Restricted for all three fields.
- D. Internal for all three fields.

Answer: B

Explanation

Names and addresses are personally identifiable and typically classified as confidential or restricted; property valuations are public. Options A, C, and D misclassify one or more fields.

Question

A data scientist needs to know the exact business definition of a column named "ltv_ratio". Where is this information most reliably documented?

- A. In the raw data file.
- B. In the data dictionary or metadata repository.
- C. In the model's source code.
- D. In the marketing brochure.

Answer: B

Explanation

Column definitions are documented in the data dictionary or metadata repository. Options A, C, and D do not reliably contain this information.

Question

Which access control model uses attributes of the user, the resource, and the environment to make authorisation decisions?

- A. Role-Based Access Control (RBAC).
- B. Attribute-Based Access Control (ABAC).
- C. Mandatory Access Control (MAC).
- D. Discretionary Access Control (DAC).

Answer: B

Explanation

ABAC uses attributes of user, resource, and environment. RBAC uses roles; MAC uses security labels; DAC lets owners set permissions.

Question

Which of the following is NOT a data protection principle commonly found in privacy regulations?

- A. Data minimisation.
- B. Purpose limitation.
- C. Storage limitation.
- D. Unlimited retention.

Answer: D

Explanation

Modern privacy regulations require storage limitation; data must not be kept longer than necessary. Unlimited retention violates this principle.

Question

A data retention policy should specify:

- A. That all data is kept indefinitely.
- B. How long each category of data is retained and the process for secure deletion.
- C. Only the retention of marketing data.
- D. Only regulatory retention.

Answer: B

Explanation

Retention policies define retention periods and deletion procedures. Option A is contrary to privacy principles. Options C and D are incomplete.

Question

Which is the primary purpose of encrypting sensitive data both at rest and in transit?

- A. To increase storage capacity.
- B. To protect confidentiality and integrity against unauthorised access.
- C. To reduce processing time.
- D. To comply with marketing standards.

Answer: B

Explanation

Encryption protects data confidentiality and integrity. Options A, C, and D are unrelated.

Question

Which of the following is typically part of a model governance policy?

- A. The definition of a model and non-model.
- B. The model risk tiering methodology.
- C. Validation and documentation requirements.
- D. All of the above.

Answer: D

Explanation

Model governance policies define models, tiers, validation, documentation, monitoring, and reporting. All three options are standard components.

Question

Which of the following best captures the purpose of model documentation in an AI governance framework?

- A. To increase model accuracy.
- B. To provide sufficient detail for a party unfamiliar with the model to understand its operation, limitations, and key assumptions.
- C. To reduce training time.
- D. To market the model internally.

Answer: B

Explanation

SR 11-7 emphasises documentation sufficient for unfamiliar parties. Options A, C, and D are incorrect.

Question

Which item is essential in a model inventory?

- A. Only the model name.
- B. Model owner, purpose, risk tier, validation status, and interdependencies.
- C. The model's full source code.
- D. The model's marketing brochure.

Answer: B

Explanation

The inventory tracks ownership, purpose, risk tier, validation, and dependencies. Options A, C, and D are insufficient or inappropriate.

Question

What does a model landscape primarily provide that a model inventory does not?

- A. A list of all models.
- B. A map of interactions and interdependencies, highlighting cascade risk and critical nodes.
- C. Model code.
- D. Model marketing material.

Answer: B

Explanation

Landscape captures interactions; inventory lists models. Options A, C, and D describe inventory or unrelated content.

Question

Which of the following is a standard factor in AI/ML model risk tiering?

- A. Materiality of model use.
- B. Complexity of methodology.
- C. Exposure (financial, regulatory, reputational).
- D. All of the above.

Answer: D

Explanation

Risk tiering considers materiality, complexity, and exposure. All three are essential inputs.

Question

What is the primary purpose of independent model validation?

- A. To approve the model without further review.
- B. To critically and objectively assess the model's conceptual soundness, data quality, implementation, and performance, by parties independent of the developers.
- C. To rewrite the model code.
- D. To market the model.

Answer: B

Explanation

Independent validation provides objective challenge. Options A, C, and D are contrary or unrelated.

Question

Which of the following best describes when model validation should occur?

- A. Only at model retirement.
- B. Before deployment, periodically during use, and after material changes.
- C. Only after a serious incident.
- D. Once, at design time.

Answer: B

Explanation

Validation is continuous: initial, periodic, and change-based. Options A, C, and D are too narrow.

Question

The “use test” in model validation is best described as:

- A. A quantitative accuracy metric.
- B. A qualitative assessment of whether the model is appropriately used for its intended purpose in practice.
- C. A test of model speed.
- D. A code-review technique.

Answer: B

Explanation

The use test is qualitative and focused on real-world use. Options A, C, and D are different.

Question

White-box testing examines:

- A. Only inputs and outputs.
- B. Internal logic, code, and structure with full visibility.
- C. Only user interface behaviour.
- D. Only documentation.

Answer: B

Explanation

White-box testing requires access to internal code and logic. Options A, C, and D describe black-box or documentation-focused approaches.

Question

Black-box testing is particularly useful when:

- A. Full source code is available.
- B. The model is vendor-supplied and internal logic is not accessible; behaviour is tested via inputs and outputs.
- C. There are no inputs.
- D. Only documentation matters.

Answer: B

Explanation

Black-box testing is essential for proprietary or opaque models. Options A, C, and D misstate its use.

Question

Why must model use limits be documented and enforced?

- A. Because they are optional.
- B. Because using a model outside its validated domain can produce materially incorrect results and unexpected risk.
- C. Because they speed up training.
- D. Because they reduce storage.

Answer: B

Explanation

Use limits protect against out-of-scope use. Options A, C, and D are incorrect.

Question

A material change to a production model triggers:

- A. Immediate deployment without review.
- B. A change management process including impact assessment, documentation, validation, and approval.
- C. Retirement of the model.
- D. No action.

Answer: B

Explanation

Material changes go through structured change management. Options A, C, and D violate governance.

Question

What should happen when a model is retired?

- A. It is silently deleted.
- B. It is formally decommissioned with documented rationale, inventory update, and archival of documentation.
- C. It continues to be used without governance.
- D. Nothing; retirement is not a governance concept.

Answer: B

Explanation

Retirement is formal with documentation and inventory update. Options A, C, and D violate governance.

Question

Model risk reporting to the board typically includes:

- A. Inventory summary, validation status, key issues, and remediation plans.
- B. The full source code.
- C. Marketing plans.
- D. Only the number of models.

Answer: A

Explanation

Board reporting summarises inventory, validation, issues, and remediation. Options B, C, and D are inappropriate or incomplete.

Question

Who typically sets the model risk appetite?

- A. Model developers.
- B. The board of directors.
- C. External auditors.
- D. Only the CRO.

Answer: B

Explanation

The board sets model risk appetite. Options A, C, and D are incorrect.

Question

In the Three Lines of Defence model, which line is responsible for model development and daily performance?

- A. First line.
- B. Second line.
- C. Third line.
- D. External auditors.

Answer: A

Explanation

First line = business/model owners/developers. Second line = independent validation. Third line = internal audit.

Question

Model owners are typically:

- A. External vendors.
- B. Business-unit staff using the model, accountable for its performance and use.
- C. External auditors.
- D. Regulators.

Answer: B

Explanation

Model owners are typically in the business unit using the model. Options A, C, and D are external or regulators.

Question

Ongoing model monitoring should detect:

- A. Only training errors.
- B. Performance degradation, data drift, concept drift, and other changes that could affect model reliability.
- C. Only model speed.
- D. Only storage issues.

Answer: B

Explanation

Monitoring detects degradation and drift. Options A, C, and D are too narrow.

Question

Which statement correctly distinguishes data drift from concept drift?

- A. They are the same.
- B. Data drift = change in input distribution; concept drift = change in relationship between inputs and target.
- C. Data drift = change in labels only.
- D. Concept drift = change in storage.

Answer: B

Explanation

Data drift: input distribution shift. Concept drift: feature-target relationship shift. Options A, C, and D misstate.

Question

Which of the following best describes periodic model revalidation?

- A. Never needed after initial validation.
- B. A scheduled reassessment confirming assumptions remain valid, data appropriate, and performance acceptable.
- C. Only performed at model retirement.
- D. Only if regulators ask.

Answer: B

Explanation

Periodic revalidation is scheduled and confirms ongoing fitness. Options A, C, and D are incorrect.

Question

Which of the following is the **FIRST** step in the model development process?

- A. Deploy the model.
- B. Define objectives and scope.
- C. Retire the model.
- D. Audit the model.

Answer: B

Explanation

Development starts with clear objectives and scope. Options A, C, and D are subsequent or unrelated stages.

Question

During data collection and preprocessing, analysts should:

- A. Skip documentation.
- B. Identify sources, clean data, handle missing values, and document any exclusions with rationale.
- C. Ignore outliers.
- D. Use all data without inspection.

Answer: B

Explanation

Data preparation requires cleaning, missing-value handling, and documentation. Options A, C, and D skip essential steps.

Question

Feature engineering typically includes:

- A. Combining variables, deriving new features, transformations, and dimensionality reduction.
- B. Only removing features.
- C. Only scaling.
- D. Only imputation.

Answer: A

Explanation

Feature engineering is broad: transformations, new features, combinations, reductions. Options B, C, and D are partial.

Question

Model selection should be based on:

- A. Personal preference.
- B. A comprehensive evaluation of alternatives, considering performance, interpretability, complexity, and business context.
- C. Speed alone.
- D. Cost alone.

Answer: B

Explanation

Model selection balances multiple criteria. Options A, C, and D are oversimplified.

Question

During model training, cross-validation is used to:

- A. Increase training time.
- B. Estimate generalisation performance and tune hyperparameters.
- C. Skip validation.
- D. Remove features.

Answer: B

Explanation

Cross-validation estimates generalisation and supports hyperparameter tuning. Options A, C, and D are incorrect.

Question

Developer testing before deployment is designed to:

- A. Skip validation.
- B. Evaluate performance before production release; ensure accuracy, robustness, and stability.
- C. Replace independent validation.
- D. Reduce code size.

Answer: B

Explanation

Developer testing precedes and complements independent validation. Options A, C, and D misstate its purpose.

Question

Which of the following is NOT a standard testing stage?

- A. Unit tests.
- B. Integration tests.
- C. Performance tests.
- D. Marketing tests.

Answer: D

Explanation

Standard testing: unit, component, integration, system, performance, regression, acceptance, adversarial. Marketing is not a testing stage.

Question

What is adversarial testing in the context of AI models?

- A. Standard unit testing.
- B. Actively probing the model with crafted inputs to discover unintended behaviour or vulnerabilities (red teaming).
- C. Testing hardware.
- D. Testing marketing materials.

Answer: B

Explanation

Adversarial testing = red teaming with crafted inputs. Options A, C, and D are unrelated.

Question

Regression testing ensures:

- A. New functionality works as expected.
- B. The model continues to function correctly after changes elsewhere in the system.
- C. Faster performance.
- D. Lower cost.

Answer: B

Explanation

Regression tests verify existing functionality after modifications. Options A, C, and D describe different goals.

Question

Acceptance testing determines whether:

- A. The model meets agreed-upon requirements and is ready for deployment.
- B. The model is the fastest available.
- C. The model is cheapest.
- D. The code is clean.

Answer: A

Explanation

Acceptance testing verifies readiness for deployment. Options B, C, and D are not the criterion.

Question

Which is a primary benefit of an AI ethics framework?

- A. Increased training speed.
- B. Clear guidance for evaluating ethical dilemmas, enhanced trust, and reduced regulatory non-compliance risk.
- C. Reduced model size.
- D. Eliminated bias.

Answer: B

Explanation

Ethics frameworks provide guidance and reduce regulatory risk. Options A, C, and D are incorrect or overstated.

Question

Which of the following best captures the concept of Responsible AI?

- A. Maximising model accuracy without regard for consequences.
- B. Ethical, transparent, accountable AI aligned with societal values and mitigating potential harms.
- C. Deploying AI as fast as possible.
- D. Ignoring regulations.

Answer: B

Explanation

Responsible AI encompasses ethics, transparency, accountability, and harm mitigation. Options A, C, and D contradict the concept.

Question

Which is a governance challenge for AI?

- A. Power asymmetries between large tech firms and regulators.
- B. Lack of interpretability.
- C. Unpredictability of AI systems.
- D. All of the above.

Answer: D

Explanation

Power asymmetries, opacity, and unpredictability are all governance challenges.

Question

Compared to medical ethics, what is missing from AI ethics?

- A. Nothing.
- B. Mature professional structures: licensing, mandatory training, ethics committees, and enforceable codes.
- C. Only training.
- D. Only licensing.

Answer: B

Explanation

AI ethics lacks the mature professional structures of medical ethics. Options A, C, and D are too narrow.

Question

The EU AI Act regulates AI based on:

- A. A single uniform standard.
- B. A risk-tiered approach: unacceptable, high, limited, minimal.
- C. Self-certification only.
- D. No requirements.

Answer: B

Explanation

The AI Act uses a risk-tiered approach. Options A, C, and D misstate the structure.

Question

Which best describes the US federal AI regulatory landscape?

- A. A single comprehensive AI law.
- B. No comprehensive federal AI law; a patchwork of state laws, executive orders, agency guidance, and frameworks like NIST AI RMF.
- C. Only state-level AI regulation.
- D. A total prohibition on AI in finance.

Answer: B

Explanation

The US lacks comprehensive federal AI law; regulation is a patchwork plus frameworks. Options A, C, and D are inaccurate.

Question

Which best describes China's AI regulation?

- A. No AI regulation.
- B. Targeted regulations including the Provisional Administrative Measures for Generative AI Services and PIPL.
- C. Only medical AI regulation.
- D. Complete industry self-regulation.

Answer: B

Explanation

China has adopted targeted AI regulations. Options A, C, and D are incorrect.

Question

Which statement best captures the relationship between ethics and law in AI governance?

- A. They are identical.
- B. Law sets minimum standards; ethics goes beyond, guiding morally appropriate behaviour even where law is silent.
- C. Ethics is subordinate to law.
- D. Ethics replaces law.

Answer: B

Explanation

Ethics extends beyond law, offering broader moral guidance. Options A, C, and D are incorrect.

Question

Which statement correctly distinguishes consequentialism, deontology, and virtue ethics?

- A. All three are identical.
- B. Consequentialism focuses on outcomes; deontology on duties; virtue ethics on character.
- C. All three focus only on outcomes.
- D. All three ignore outcomes.

Answer: B

Explanation

The three frameworks differ in focus. Options A, C, and D misstate them.

Question

Which statement about justice is correct?

- A. Different distributive justice concepts can conflict.
- B. Impartiality is not required for procedural justice.
- C. Autonomy is secondary to social justice.
- D. Meritocracy is not a distributive justice concept.

Answer: A

Explanation

Distributive justice concepts can conflict. Options B, C, and D are incorrect.

Question

Which of the following is a GenAI-specific risk category?

- A. Functional risks (hallucinations, unpredictability, adversarial inputs).
- B. Data and information risks.
- C. Organizational and strategic risks.
- D. All of the above.

Answer: D

Explanation

GenAI risk profile includes functional, data/information, organizational, strategic, and societal risks. All three listed here apply.

Question

Which is the best mitigation for hallucination risk in a legal-advice GenAI application?

- A. Increase temperature.
- B. Ground the model with retrieval-augmented generation (RAG) using authoritative sources, and require human review for critical outputs.
- C. Use a smaller model.
- D. Disable the model.

Answer: B

Explanation

RAG grounds output; human review catches remaining errors. Options A, C, and D are inadequate or extreme.

Question

What is a key risk of relying on third-party GenAI providers?

- A. Providers may update or deprecate models without notice, changing behaviour unpredictably.
- B. Providers never change their models.
- C. Providers guarantee identical behaviour forever.
- D. Providers remove all risks.

Answer: A

Explanation

Vendor updates create unpredictability. Options B, C, and D are incorrect.

Question

Which is an internal governance lever for GenAI?

- A. Assigning clear ownership and accountability across the GenAI lifecycle.
- B. Tracking external regulatory changes.
- C. Relying only on external standards.
- D. Outsourcing all governance.

Answer: A

Explanation

Internal levers include ownership, review processes, escalation, training. Options B and C are external; D is contrary to governance.

Question

Which is an external governance constraint for GenAI?

- A. Evolving regulatory requirements.
- B. Contractual terms with model providers.
- C. Emerging standards (e.g., NIST AI RMF, ISO/IEC 42001).
- D. All of the above.

Answer: D

Explanation

All three are external governance constraints on GenAI deployments.

Question

Why is continuous monitoring particularly important for GenAI systems?

- A. Because outputs are non-deterministic, providers may update models, and drift can be subtle.
- B. Because GenAI requires no monitoring.
- C. Because GenAI never changes.
- D. Because GenAI cannot be validated.

Answer: A

Explanation

Non-determinism, vendor updates, and subtle drift require ongoing monitoring. Options B, C, and D are incorrect.

Question

Which is the primary purpose of human-in-the-loop (HITL) design?

- A. To increase model speed.
- B. To include human review and verification at critical points, mitigating errors and ensuring oversight.
- C. To replace model validation.
- D. To eliminate human roles.

Answer: B

Explanation

HITL adds human oversight at critical points. Options A, C, and D contradict the purpose.

Question

Which of the following makes GenAI governance harder than traditional ML governance?

- A. Open-ended unstructured output, non-determinism, and vendor dependence.
- B. Simplicity of GenAI systems.
- C. Absence of regulatory exposure.
- D. Universal open-sourcing.

Answer: A

Explanation

Open-ended output, non-determinism, and vendor dependence are the key challenges.

Question

When registering a GenAI system in a model inventory, which additional aspects should be captured beyond traditional models?

- A. The underlying foundation model, prompt templates, and any retrieval/knowledge sources.
- B. Only the API provider.
- C. Only the temperature parameter.
- D. Only the compute costs.

Answer: A

Explanation

GenAI inventory requires capturing foundation model, prompts, and retrieval data. Options B, C, and D are partial.

Question

Which is a unique challenge when validating GenAI models?

- A. Same challenges as traditional models.
- B. Open-ended output, hallucinations, multi-modality, and difficulty establishing ground truth.
- C. Only speed.
- D. Only cost.

Answer: B

Explanation

GenAI validation faces open-ended output, hallucination, multi-modality, and ground-truth issues. Options A, C, and D underestimate the challenge.

Question

Why does the AI/ML paradigm shift governance emphasis toward data?

- A. Data is irrelevant to AI.
- B. AI/ML models often learn directly from data, making data quality, provenance, and bias central to model behaviour and risk.
- C. Models do not depend on data.
- D. Governance is unchanged.

Answer: B

Explanation

AI/ML models learn from data; data governance is therefore central. Options A, C, and D are incorrect.

Question

Which attribute is most relevant when classifying AI/ML models for ethical review?

- A. Model accuracy only.
- B. Data bias, explainability, and recalibration needs, among others.
- C. Model speed.
- D. Compute cost.

Answer: B

Explanation

Ethical review focuses on bias, explainability, and recalibration. Options A, C, and D are secondary.

Question

A model landscape captures interdependencies because:

- A. Models operate in isolation.
- B. Outputs of one model are often inputs to others, so failures can cascade across an organisation.
- C. Dependencies are irrelevant.
- D. Only audit cares about them.

Answer: B

Explanation

Interdependencies create cascading risk. Options A, C, and D are incorrect.

Question

Aggregate model risk can be influenced by:

- A. Only individual model quality.
- B. Model interdependencies, shared assumptions, and correlated methodology.
- C. Only data size.
- D. Only IT infrastructure.

Answer: B

Explanation

Aggregate risk arises from interdependencies and shared assumptions. Options A, C, and D miss this.

Question

Why can aggregating data across sources introduce model risk?

- A. Aggregation always improves quality.
- B. Inconsistent definitions, timeframes, or measurement scales can distort features and lead to biased models.
- C. Aggregation eliminates risk.
- D. Aggregation is always safe.

Answer: B

Explanation

Aggregation can introduce inconsistencies. Options A, C, and D are incorrect.

Question

Which is an example of a production feedback loop creating model risk?

- A. A model trained on past decisions then used to make new decisions, which become the next training data.
- B. A model without historical data.
- C. A model that is never deployed.
- D. A model that is retired.

Answer: A

Explanation

Feedback loops reinforce biases. Options B, C, and D are not feedback loops.

Question

Why is data and model lineage important for AI governance?

- A. To track data and model transformations end-to-end, enabling reproducibility, auditability, and root-cause analysis.
- B. To speed up training.
- C. To reduce costs.
- D. To remove features.

Answer: A

Explanation

Lineage supports reproducibility and audit. Options B, C, and D are secondary.

Question

Data stewards are responsible for:

- A. Only fixing typos.
- B. Managing data quality, access, definitions, and use within a governance framework.
- C. Building models.
- D. Marketing data.

Answer: B

Explanation

Data stewards manage quality, access, and definitions. Options A, C, and D are too narrow or unrelated.

Question

Data custodians are primarily responsible for:

- A. Business ownership of data.
- B. Technical storage, security, and access controls for data.
- C. Building models.
- D. Marketing data.

Answer: B

Explanation

Custodians handle technical storage and access. Stewards handle business data quality; owners provide governance. Options A, C, and D are not custody roles.

Question

Data governance reporting to the board typically includes:

- A. Only storage costs.
- B. Data quality metrics, incidents, compliance status, and key risks.
- C. Only IT staff lists.
- D. Only marketing spend.

Answer: B

Explanation

Board-level reporting covers quality, incidents, compliance, and risk. Options A, C, and D are irrelevant.

Question

Which of the following is a trigger for ad-hoc model revalidation?

- A. Only scheduled review.
- B. Significant changes in inputs, environment, or performance deviations.
- C. Only cost increases.
- D. Only headcount changes.

Answer: B

Explanation

Ad-hoc revalidation is triggered by material changes or performance issues. Options A, C, and D are not standard triggers.

Question

Strong model documentation should include:

- A. Only the model name.
- B. Purpose, methodology, assumptions, data, limitations, validation results, and usage instructions.
- C. Only the code.
- D. Only the developer's email.

Answer: B

Explanation

Documentation is comprehensive: purpose, methodology, assumptions, data, limits, validation, usage. Options A, C, and D are insufficient.

Question

Recalibration of a model is appropriate when:

- A. The model's structure is wrong.
- B. The core structure remains valid but parameters need updating in light of new data.
- C. The model is being retired.
- D. The model has never been used.

Answer: B

Explanation

Recalibration updates parameters; structural issues require redevelopment. Options A, C, and D describe different actions.

Question

A model risk appetite statement typically includes:

- A. Detailed code.
- B. The organisation's tolerance for model risk, escalation triggers, and consequences.
- C. Only model inventory.
- D. Only validation reports.

Answer: B

Explanation

Appetite statements define tolerance and triggers. Options A, C, and D are unrelated.

Question

Which of the following is NOT typically included in a model risk taxonomy?

- A. Conceptual risk.
- B. Data risk.
- C. Implementation risk.
- D. Employee satisfaction risk.

Answer: D

Explanation

Model risk taxonomies focus on conceptual, data, and implementation risk (plus usage risk). Employee satisfaction is not standard.

Question

A model risk register typically tracks:

- A. Only model code.
- B. Identified risks, owners, mitigation actions, deadlines, and status.
- C. Only training data.
- D. Only IT systems.

Answer: B

Explanation

A risk register tracks identified risks, owners, actions, deadlines, and status. Options A, C, and D are unrelated.

Question

A model risk committee is typically responsible for:

- A. Approving models and overseeing model risk governance at the enterprise level.
- B. Entering data.
- C. Training models.
- D. Marketing models.

Answer: A

Explanation

The model risk committee provides enterprise oversight and approvals. Options B, C, and D are operational roles.

Question

Effective model risk training programs should include:

- A. Only technical content.
- B. Technical content plus governance, ethics, and regulatory expectations.
- C. Only ethics.
- D. Only regulation.

Answer: B

Explanation

Model risk training should be comprehensive: technical, governance, ethics, regulation.

Question

A strong model risk culture is characterised by:

- A. Silence about risks.
- B. Awareness, accountability, open communication, and continuous learning around model risk.
- C. Only senior leaders caring about risk.
- D. Ignoring model limitations.

Answer: B

Explanation

A strong culture values awareness, accountability, openness, and learning. Options A, C, and D contradict this.

Question

Model risk reporting to governance bodies typically occurs:

- A. Only when problems arise.
- B. On a scheduled basis (e.g., quarterly) and ad hoc for significant issues.
- C. Never.
- D. Only at year-end.

Answer: B

Explanation

Reporting is scheduled and ad hoc as needed. Options A, C, and D are insufficient.

Question

Which of the following is a lawful basis for processing personal data under GDPR?

- A. Consent.
- B. Contract performance.
- C. Legal obligation.
- D. All of the above.

Answer: D

Explanation

GDPR Art. 6 specifies lawful bases: consent, contract, legal obligation, vital interests, public task, legitimate interests. Options A, B, and C are each valid bases.

Question

Which is a data subject right under GDPR?

- A. Right to access.
- B. Right to rectification.
- C. Right to erasure.
- D. All of the above.

Answer: D

Explanation

GDPR grants multiple rights: access, rectification, erasure, restriction, portability, objection. Options A, B, and C are each among them.

Question

Privacy by design requires:

- A. Adding privacy features at the end.
- B. Embedding privacy protections into the design and architecture of systems from the outset.
- C. Ignoring privacy until a complaint.
- D. Only privacy policies.

Answer: B

Explanation

Privacy by design embeds protections from the start. Options A, C, and D are inconsistent with the principle.

Question

What is the difference between anonymisation and pseudonymisation?

- A. They are the same.
- B. Anonymisation irreversibly removes identifiability; pseudonymisation replaces identifiers with a key that allows re-identification.
- C. Pseudonymisation is irreversible.
- D. Anonymisation always fails.

Answer: B

Explanation

Anonymisation is irreversible; pseudonymisation can be reversed with the key. Options A, C, and D misstate the distinction.

Question

What is the purpose of differential privacy?

- A. To remove all data.
- B. To allow statistical analysis of a dataset while protecting individual records, by adding calibrated noise.
- C. To encrypt the data.
- D. To compress data.

Answer: B

Explanation

Differential privacy adds noise to protect individual records while preserving aggregate statistics. Options A, C, and D are different techniques.

Question

Which of the following is typically part of an ethical AI review?

- A. Only model accuracy.
- B. Assessing fairness, transparency, accountability, and potential harms.
- C. Only speed.
- D. Only cost.

Answer: B

Explanation

Ethical AI reviews assess fairness, transparency, accountability, and harms. Options A, C, and D are insufficient.

Question

An AI ethics committee typically:

- A. Writes code.
- B. Provides guidance on ethical issues, reviews high-risk AI use cases, and ensures alignment with values.
- C. Markets AI products.
- D. Trains models.

Answer: B

Explanation

AI ethics committees provide guidance and review. Options A, C, and D are operational or unrelated.

Question

Transparency reports in AI typically include:

- A. Only marketing content.
- B. Model usage, performance, fairness metrics, and incidents.
- C. Only costs.
- D. Only IT infrastructure.

Answer: B

Explanation

Transparency reports cover usage, performance, fairness, incidents. Options A, C, and D are irrelevant or too narrow.

Question

Why is stakeholder engagement important in responsible AI?

- A. To increase marketing.
- B. To incorporate diverse perspectives, understand impacts, and build trust with users, regulators, and affected parties.
- C. To reduce training costs.
- D. To skip governance.

Answer: B

Explanation

Engagement brings diverse views and builds trust. Options A, C, and D are not the purpose.

Question

An algorithmic impact assessment (AIA) typically evaluates:

- A. Only model accuracy.
- B. Risks and impacts on individuals, groups, and society, especially for high-risk use cases.
- C. Only cost.
- D. Only speed.

Answer: B

Explanation

AIAs evaluate broader risks and impacts. Options A, C, and D are too narrow.

Question

Which of the following model types is typically most interpretable?

- A. Deep neural network.
- B. Gradient-boosted ensemble.
- C. Logistic regression.
- D. Random forest.

Answer: C

Explanation

Logistic regression coefficients are directly interpretable. Deep nets, boosting, and RF are less interpretable. Options A, B, and D are less interpretable.

Question

Which is a limitation of impurity-based feature importance for tree ensembles?

- A. It is always accurate.
- B. It can be biased toward high-cardinality or correlated features.
- C. It cannot be computed.
- D. It replaces SHAP.

Answer: B

Explanation

Impurity-based importance is biased; permutation importance or SHAP are alternatives. Options A, C, and D are incorrect.

Question

Partial dependence plots show:

- A. Raw data.
- B. The marginal effect of a feature on the model's prediction, averaged over other features.
- C. Model code.
- D. Only training errors.

Answer: B

Explanation

PDPs display marginal feature effects. Options A, C, and D are unrelated.

Question

Which property does SHAP have that many other feature attribution methods lack?

- A. Non-determinism.
- B. Consistency — if a feature's contribution increases, its SHAP value does not decrease.
- C. Randomness.
- D. Irreproducibility.

Answer: B

Explanation

SHAP satisfies consistency, along with local accuracy and missingness. Options A, C, and D are not properties of SHAP.

Question

Which is a limitation of LIME?

- A. It is always globally faithful.
- B. It provides local approximations that may not be stable across similar instances and are not guaranteed to be globally faithful.
- C. It has no applications.
- D. It replaces model training.

Answer: B

Explanation

LIME is local and can be unstable. Options A, C, and D misstate.

Question

Counterfactual explanations are particularly useful because:

- A. They describe model code.
- B. They provide actionable, human-understandable explanations of what would need to change for a different decision.
- C. They replace SHAP.
- D. They measure accuracy.

Answer: B

Explanation

Counterfactuals provide actionable insights. Options A, C, and D are incorrect.

Question

Model calibration refers to:

- A. Ensuring predicted probabilities match observed frequencies.
- B. Model accuracy.
- C. Model speed.
- D. Model size.

Answer: A

Explanation

Calibration ensures probabilistic predictions are reliable. Options B, C, and D are different concepts.

Question

When choosing a classification threshold, the analyst should consider:

- A. Only the default threshold of 0.5.
- B. The relative costs and benefits of false positives versus false negatives, aligned with business objectives.
- C. Only accuracy.
- D. Only precision.

Answer: B

Explanation

Threshold choice depends on cost-benefit asymmetry. Options A, C, and D are oversimplifications.

Question

When data drift is detected in a production model, which of the following is generally the appropriate response?

- A. Ignore it.
- B. Investigate cause, consider recalibration or retraining, and update governance documentation accordingly.
- C. Delete the model.
- D. Change the target.

Answer: B

Explanation

Drift triggers investigation and possible recalibration/retraining. Options A, C, and D are not standard.

Question

Which of the following best summarises the guiding principle of Responsible AI governance in financial services?

- A. Maximise model complexity.
- B. Balance performance with governance, validation, monitoring, fairness, explainability, and ethical alignment with regulatory and business requirements.
- C. Deploy all models without validation.
- D. Ignore risks.

Answer: B

Explanation

Responsible AI balances performance with governance, validation, monitoring, fairness, explainability, and ethics. Options A, C, and D contradict the principle.