

GARP Risk & AI — Q551–Q650

Unofficial Exam Prep

Unofficial

October 3, 2026

Question

How does model risk differ from technology risk in an AI governance framework?

- A. They are identical.
- B. Model risk focuses on potential adverse consequences from model errors, misuse, or limitations; technology risk focuses on IT infrastructure, systems, and operational failures.
- C. Model risk applies only to AI; technology risk applies only to statistics.
- D. Neither is regulated.

Answer: B

Explanation

Model risk relates to model-based decisions; technology risk covers systems and infrastructure. Options A, C, and D misstate the distinction.

Question

Which of the following is the primary purpose of an AI risk register?

- A. List every AI model in the firm.
- B. Track identified AI risks, their owners, mitigation actions, and remediation status.
- C. Replace the model inventory.
- D. Document marketing strategy.

Answer: B

Explanation

The AI risk register tracks identified risks, owners, mitigations, and status. Options A, C, and D are different artifacts or unrelated.

Question

Why does an organisation need an AI risk appetite?

- A. To replace regulation.
- B. To define the level and types of AI risk the organisation is willing to accept, guiding investment and control decisions.
- C. To eliminate all AI risk.
- D. To market AI products.

Answer: B

Explanation

AI risk appetite defines acceptable risk to guide decisions. Options A, C, and D misstate its role.

Question

What is the key purpose of enterprise-wide AI governance?

- A. Local optimisation of one team.
- B. Consistent policies, oversight, and controls applied across all AI initiatives, aligned to overall business strategy and risk appetite.
- C. Eliminating AI use.
- D. Outsourcing AI to vendors.

Answer: B

Explanation

Enterprise AI governance provides consistent policies and oversight. Options A, C, and D are inconsistent with the concept.

Question

Which event should trigger an escalation in an AI governance framework?

- A. Routine model retraining.
- B. Material breach of AI risk tolerance, such as a serious fairness violation or a regulatory finding.
- C. Minor UI changes.
- D. Documentation updates.

Answer: B

Explanation

Material breaches require escalation. Options A, C, and D are routine activities.

Question

How often should AI risk reporting occur to senior management?

- A. Only annually.
- B. Regularly (e.g., quarterly) and ad hoc for significant events.
- C. Never.
- D. Only at year-end.

Answer: B

Explanation

Reporting is regular and ad hoc. Options A, C, and D are insufficient.

Question

How does an AI portfolio differ from a model inventory?

- A. They are identical.
- B. A portfolio is a strategic view of AI capabilities and initiatives; the inventory is an operational registry of deployed models.
- C. Portfolio is only for vendors.
- D. Inventory is only for retired models.

Answer: B

Explanation

Portfolio is strategic; inventory is operational. Options A, C, and D misstate.

Question

Which is NOT typically included in board-level AI reporting?

- A. Aggregate AI risk exposure.
- B. Fairness and bias metrics.
- C. Detailed source code of every model.
- D. Material incidents and remediation plans.

Answer: C

Explanation

Board reports are strategic and aggregate; source code is not appropriate. Options A, B, and D are typically included.

Question

What is systemic AI risk?

- A. A minor operational issue.
- B. Risk that a widespread AI failure or correlated AI behaviour could affect multiple firms or the broader financial system.
- C. Only internal risk.
- D. Irrelevant to regulation.

Answer: B

Explanation

Systemic AI risk extends beyond single-firm boundaries. Options A, C, and D understate or dismiss.

Question

Which is a common category in AI risk taxonomies?

- A. Technical risk.
- B. Ethical risk.
- C. Operational risk.
- D. All of the above.

Answer: D

Explanation

AI risk taxonomies span technical, ethical, operational, and other categories.

Question

Which is the most relevant model risk for an AI-based credit scoring system?

- A. Colour of the user interface.
- B. Potential unfairness, inaccuracy, and adverse regulatory consequences from biased or poorly performing predictions.
- C. Storage cost only.
- D. Marketing effectiveness only.

Answer: B

Explanation

Credit scoring model risk includes fairness, accuracy, and regulatory risk. Options A, C, and D are secondary or irrelevant.

Question

For an AI fraud detection model, which is the most critical model risk?

- A. Overfitting to historical fraud patterns, missing novel fraud schemes.
- B. Storage cost.
- C. User interface design.
- D. Training time.

Answer: A

Explanation

Overfitting to past patterns can miss novel fraud. Options B, C, and D are less critical.

Question

Which is a key risk for an AI anti-money-laundering (AML) model?

- A. Low alert volume.
- B. High false-positive rate overwhelming compliance teams and potential missed true positives, both creating regulatory and operational risk.
- C. Storage cost.
- D. Marketing.

Answer: B

Explanation

AML risk includes excessive false positives and missed true positives. Options A, C, and D are secondary.

Question

Which is a typical risk for an AI-enhanced market risk model?

- A. Training on non-stationary data leading to unstable predictions.
- B. Data storage.
- C. Logo design.
- D. Employee satisfaction.

Answer: A

Explanation

Non-stationarity of financial markets is a key challenge. Options B, C, and D are unrelated.

Question

Which is a specific concern for AI underwriting in insurance?

- A. Ethical and legal issues from using alternative data sources that may be considered sensitive or discriminatory.
- B. Faster approvals.
- C. Lower costs.
- D. Better customer experience.

Answer: A

Explanation

Alternative data raises ethical and legal concerns. Options B, C, and D are benefits, not concerns.

Question

Which is a recognised fairness metric used in lending AI?

- A. Demographic parity.
- B. Storage cost parity.
- C. Latency parity.
- D. Colour parity.

Answer: A

Explanation

Demographic parity is a standard fairness metric. Options B, C, and D are not.

Question

Which is a key risk of algorithmic trading models?

- A. Flash-crash amplification, unintended feedback loops, and market impact from correlated strategies.
- B. Storage capacity.
- C. Logo design.
- D. Employee benefits.

Answer: A

Explanation

Algorithmic trading risks include flash crashes and correlated behaviour. Options B, C, and D are unrelated.

Question

Which is the most significant risk from deploying a GenAI chatbot in banking?

- A. Generating incorrect or non-compliant advice to customers, creating legal and reputational risk.
- B. Faster response times.
- C. Cost savings.
- D. Customer satisfaction.

Answer: A

Explanation

Advice quality and compliance are central concerns. Options B, C, and D are benefits.

Question

Which is a key governance concern when using alternative data in financial AI?

- A. Legal right to use the data, privacy, and potential for discriminatory impact.
- B. Storage cost.
- C. Model speed.
- D. Logo design.

Answer: A

Explanation

Legal rights, privacy, and fairness are key governance concerns. Options B, C, and D are secondary.

Question

What is a specific AI risk in AML/KYC?

- A. Customer friction from over-flagging and regulatory exposure from under-flagging, requiring careful threshold calibration.
- B. Storage cost.
- C. Logo design.
- D. Training time.

Answer: A

Explanation

Threshold calibration affects both customer experience and regulatory compliance. Options B, C, and D are less critical.

Question

Which is the correct order of model lifecycle stages?

- A. Development, validation, deployment, monitoring, retirement.
- B. Deployment, development, monitoring, validation, retirement.
- C. Monitoring, development, retirement, validation, deployment.
- D. Retirement, development, deployment, monitoring, validation.

Answer: A

Explanation

Development → validation → deployment → monitoring → retirement. Options B, C, and D reorder incorrectly.

Question

Why must validation occur before deployment?

- A. Regulatory convenience.
- B. To independently confirm the model meets requirements, performs as expected, and is fit for purpose before being used in production.
- C. To reduce cost.
- D. To increase speed.

Answer: B

Explanation

Pre-deployment validation confirms fitness. Options A, C, and D are secondary.

Question

Which activity is part of post-deployment monitoring?

- A. Continuous tracking of model performance, data drift, and input integrity, with alerts for threshold breaches.
- B. Only training.
- C. Only validation.
- D. Only code review.

Answer: A

Explanation

Post-deployment monitoring tracks performance and drift. Options B, C, and D are pre-deployment.

Question

Recalibration is preferred to redevelopment when:

- A. The model structure is wrong.
- B. The model structure remains valid but parameters need updating using new data.
- C. The model is retired.
- D. The model has never been validated.

Answer: B

Explanation

Recalibration updates parameters; structural issues require redevelopment. Options A, C, and D are different situations.

Question

What is model deprecation?

- A. Immediate retirement.
- B. The formal marking of a model as no longer recommended for use, typically while a replacement is being prepared.
- C. Model creation.
- D. Model validation.

Answer: B

Explanation

Deprecation flags models for phase-out. Options A, C, and D describe different actions.

Question

What does a model change impact assessment evaluate?

- A. Only cost.
- B. How a proposed change affects model outputs, downstream users, risk profile, and regulatory compliance.
- C. Only speed.
- D. Only storage.

Answer: B

Explanation

Impact assessment covers outputs, downstream, risk, and compliance. Options A, C, and D are too narrow.

Question

Why is model versioning important?

- A. To reduce storage.
- B. To track distinct model iterations, enabling reproducibility, audit, and controlled rollback if needed.
- C. To reduce cost.
- D. To speed up training.

Answer: B

Explanation

Versioning supports reproducibility, audit, and rollback. Options A, C, and D are secondary.

Question

What is the champion/challenger framework?

- A. A marketing strategy.
- B. A framework where the current production model (champion) is compared against alternative models (challengers) to evaluate whether a switch would improve performance.
- C. A training technique.
- D. A validation report.

Answer: B

Explanation

Champion/challenger compares production models to alternatives. Options A, C, and D are unrelated.

Question

What is shadow deployment of a model?

- A. Deploying without monitoring.
- B. Running a new model alongside production without using its outputs, to evaluate performance in real conditions.
- C. Deploying to production only.
- D. Retiring a model.

Answer: B

Explanation

Shadow deployment evaluates models in parallel without affecting production. Options A, C, and D misstate.

Question

What is A/B testing in the AI deployment context?

- A. Randomly assigning users to two model versions and comparing outcomes.
- B. Retiring both models.
- C. Doubling training data.
- D. Storing data twice.

Answer: A

Explanation

A/B testing compares model versions via controlled experiments. Options B, C, and D are unrelated.

Question

Which is a key element of an AI incident response plan?

- A. Ignoring incidents.
- B. Defined escalation paths, roles, playbooks, and communication plans for handling AI-related incidents.
- C. Only post-mortems.
- D. Only IT involvement.

Answer: B

Explanation

An incident response plan includes escalation, roles, playbooks, and communication. Options A, C, and D are insufficient.

Question

An AI model begins producing biased outputs. Which is the most likely root cause?

- A. Data drift or newly biased data.
- B. Server failure.
- C. Power outage.
- D. UI redesign.

Answer: A

Explanation

Bias in outputs typically stems from data or model development. Options B, C, and D are unrelated.

Question

What is the purpose of a post-incident review for AI incidents?

- A. Assign blame.
- B. Learn from the incident, improve controls and processes, and prevent recurrence.
- C. End investigation.
- D. Skip documentation.

Answer: B

Explanation

Post-incident reviews focus on learning and prevention. Options A, C, and D are contrary.

Question

Why is communication important during an AI incident?

- A. It is not.
- B. Clear communication with regulators, customers, and stakeholders maintains trust and meets regulatory obligations.
- C. Only internal communication matters.
- D. Only external communication matters.

Answer: B

Explanation

Communication with all stakeholders is essential. Options A, C, and D are incomplete or wrong.

Question

How are AI incidents typically classified?

- A. By colour.
- B. By severity and impact, considering factors such as harm to customers, financial loss, and regulatory exposure.
- C. Randomly.
- D. By alphabet.

Answer: B

Explanation

Incidents are classified by severity and impact. Options A, C, and D are not standard.

Question

When a serious AI incident affects customers, what is typically required?

- A. No action.
- B. Prompt notification to regulators where required and remediation plans.
- C. Only internal review.
- D. Silence.

Answer: B

Explanation

Serious incidents may trigger regulatory notification. Options A, C, and D violate governance.

Question

Which metric helps track AI incident management effectiveness?

- A. Total storage used.
- B. Time to detection, time to remediation, recurrence rate, and impact.
- C. Marketing reach.
- D. UI clicks.

Answer: B

Explanation

Detection, remediation, recurrence, and impact are key metrics. Options A, C, and D are unrelated.

Question

What is red teaming in AI?

- A. A marketing exercise.
- B. Adversarial testing by a dedicated team to uncover vulnerabilities and unintended behaviours before deployment.
- C. A financial audit.
- D. A user survey.

Answer: B

Explanation

Red teaming proactively tests for vulnerabilities. Options A, C, and D misstate.

Question

Why is logging data important for AI incident management?

- A. To increase storage.
- B. To reconstruct events, identify root causes, and support regulatory reporting.
- C. To slow down systems.
- D. To increase costs.

Answer: B

Explanation

Logging supports root-cause analysis and reporting. Options A, C, and D are incorrect or secondary.

Question

What is the purpose of continuous improvement in AI governance?

- A. To increase costs.
- B. To iterate on governance frameworks and controls based on incidents, audits, and evolving best practices.
- C. To avoid change.
- D. To slow down deployment.

Answer: B

Explanation

Continuous improvement ensures governance keeps pace with change. Options A, C, and D are contrary.

Question

A typical AI governance maturity model includes levels such as:

- A. Ad hoc, developing, defined, managed, optimised.
- B. Red, blue, green.
- C. Only one level.
- D. No levels.

Answer: A

Explanation

Maturity models typically progress from ad hoc to optimised. Options B, C, and D are not standard.

Question

Which is a common gap in AI governance frameworks?

- A. Extensive controls.
- B. Missing governance for GenAI, alternative data, or third-party models.
- C. Complete coverage.
- D. Perfect alignment.

Answer: B

Explanation

Common gaps include GenAI, alternative data, and vendor models. Options A, C, and D are not gaps.

Question

What is the purpose of an AI governance roadmap?

- A. A marketing document.
- B. A phased plan outlining how the organisation will build, enhance, and sustain AI governance capabilities over time.
- C. A single policy.
- D. A code repository.

Answer: B

Explanation

A roadmap is a phased plan. Options A, C, and D misstate.

Question

Which is an appropriate KPI for AI governance?

- A. Number of employees.
- B. Percentage of models validated on schedule.
- C. Office size.
- D. Marketing spend.

Answer: B

Explanation

Validation timeliness is a governance KPI. Options A, C, and D are unrelated.

Question

Which audience requires AI governance training?

- A. Only data scientists.
- B. A broad set including executives, developers, validators, business users, compliance, and internal audit.
- C. Only IT.
- D. Only external auditors.

Answer: B

Explanation

AI governance training spans many roles. Options A, C, and D are too narrow.

Question

Why is data collection for AI governance metrics important?

- A. To meet marketing goals.
- B. To enable measurement, benchmarking, and continuous improvement of AI governance practices.
- C. To increase costs.
- D. To slow operations.

Answer: B

Explanation

Data supports measurement and improvement. Options A, C, and D are incorrect.

Question

What is the purpose of AI governance certifications (e.g., ISO/IEC 42001)?

- A. Marketing only.
- B. To demonstrate externally that the organisation meets recognised standards for AI governance.
- C. To replace regulations.
- D. To lower costs.

Answer: B

Explanation

Certifications demonstrate conformance with standards. Options A, C, and D are not the primary purpose.

Question

Why is board AI literacy important?

- A. To increase headcount.
- B. To enable effective oversight of AI strategy, risk, and governance at the highest level.
- C. To reduce costs.
- D. To speed deployment.

Answer: B

Explanation

Board AI literacy enables effective oversight. Options A, C, and D are not the primary reason.

Question

How should an AI ethics committee integrate with the model risk function?

- A. No integration is needed.
- B. Through coordinated review of high-risk AI use cases, sharing of insights on fairness and ethics, and joint escalation to senior management.
- C. Only through annual meetings.
- D. Only through external advisors.

Answer: B

Explanation

Coordination between ethics and model risk is essential. Options A, C, and D are insufficient.

Question

Why must AI governance frameworks evolve?

- A. To increase cost.
- B. Because AI technologies, use cases, and regulations change rapidly, requiring ongoing updates to remain effective.
- C. To slow operations.
- D. To reduce transparency.

Answer: B

Explanation

Rapid change in AI and regulation requires evolving governance. Options A, C, and D contradict this.

Question

How can GenAI undermine human autonomy?

- A. By providing choices.
- B. By generating persuasive, tailored content that manipulates users without their awareness.
- C. By explaining decisions.
- D. By slowing systems.

Answer: B

Explanation

Manipulative content undermines autonomy. Options A, C, and D respect or are unrelated to autonomy.

Question

How does GenAI amplify bias?

- A. By generating content that reflects biases present in training data, at scale.
- B. By reducing accuracy.
- C. By slowing training.
- D. By increasing costs.

Answer: A

Explanation

GenAI can amplify training-data bias at scale. Options B, C, and D are unrelated.

Question

Why is GenAI a challenge for truth and accountability?

- A. It only outputs truth.
- B. It generates plausible but sometimes incorrect content, making accountability for misinformation difficult.
- C. It has no training data.
- D. It never hallucinates.

Answer: B

Explanation

Hallucinations and plausible output complicate accountability. Options A, C, and D are incorrect.

Question

What consent issues arise with GenAI training data?

- A. None.
- B. Data subjects whose content was used in training may not have consented, raising ethical and legal concerns.
- C. Consent is automatically granted.
- D. Consent is not required.

Answer: B

Explanation

Training data consent is a critical concern. Options A, C, and D are incorrect.

Question

Why is attribution important in GenAI?

- A. It isn't.
- B. To trace sources of content and information, supporting transparency and possibly legal rights such as copyright.
- C. To increase cost.
- D. To slow inference.

Answer: B

Explanation

Attribution supports transparency and copyright compliance. Options A, C, and D are incorrect.

Question

What is a copyright concern with GenAI?

- A. None.
- B. Training on copyrighted material without permission and generating content that may infringe copyright.
- C. GenAI is always copyright-clean.
- D. Copyright does not apply.

Answer: B

Explanation

Copyright concerns are significant. Options A, C, and D are incorrect.

Question

Which is a safety concern with GenAI in critical applications?

- A. Faster response.
- B. Generating unsafe or harmful content that could cause physical, psychological, or financial harm.
- C. Lower costs.
- D. Higher accuracy.

Answer: B

Explanation

Harmful content generation is a safety concern. Options A, C, and D are benefits.

Question

Which consumer protection issue arises from GenAI?

- A. Increasing prices only.
- B. Misleading or inaccurate information provided to consumers, potentially leading to harm.
- C. Faster service.
- D. Lower fees.

Answer: B

Explanation

Misleading information is a consumer protection issue. Options A, C, and D are unrelated.

Question

What regulatory concern arises when GenAI is used for financial advice?

- A. Lower costs.
- B. Potential non-compliance with advice regulations if unvetted or unsuitable advice is provided.
- C. Faster response.
- D. Higher throughput.

Answer: B

Explanation

Regulatory concerns exist around advice suitability. Options A, C, and D are benefits, not concerns.

Question

How can GenAI facilitate fraud?

- A. By making it easier to create convincing phishing content and deepfakes at scale.
- B. By improving accuracy.
- C. By reducing costs.
- D. By speeding up transactions.

Answer: A

Explanation

GenAI enables scalable, convincing fraud content. Options B, C, and D are unrelated.

Question

What are the key risks of agentic AI?

- A. Unintended autonomous actions, difficulty in oversight, and cascading effects from chains of decisions.
- B. Faster decisions.
- C. Lower cost.
- D. Simplicity.

Answer: A

Explanation

Agentic AI risks relate to autonomy and oversight. Options B, C, and D are benefits or inaccurate.

Question

What risk is specific to multi-agent AI systems?

- A. Faster training.
- B. Emergent behaviour arising from interactions between agents, which may be difficult to predict or control.
- C. Lower costs.
- D. Higher accuracy.

Answer: B

Explanation

Emergent multi-agent behaviour is a specific risk. Options A, C, and D are benefits.

Question

How does AI increase cyber risk?

- A. By reducing attack surface.
- B. By enabling more sophisticated, scalable attacks, and by introducing new vulnerabilities (e.g., model extraction, adversarial inputs).
- C. By eliminating threats.
- D. By lowering costs.

Answer: B

Explanation

AI introduces new attack vectors and scales threats. Options A, C, and D are incorrect.

Question

What is AI supply chain risk?

- A. Risk of low prices.
- B. Risk arising from dependencies on external AI providers, datasets, and libraries, including security, availability, and continuity concerns.
- C. Lower costs.
- D. Higher speed.

Answer: B

Explanation

Supply chain risk involves external dependencies. Options A, C, and D are benefits or unrelated.

Question

How should third-party AI risk be managed?

- A. Ignore it.
- B. Through due diligence, contractual protections, ongoing monitoring, and contingency planning for provider failures.
- C. Only contracts.
- D. Only internal audits.

Answer: B

Explanation

Third-party risk requires due diligence, contracts, monitoring, and contingency planning. Options A, C, and D are insufficient.

Question

Why is AI concentration risk a concern?

- A. Because it reduces costs.
- B. Because reliance on a small number of AI providers or models creates correlated failure risk across the ecosystem.
- C. Because it increases accuracy.
- D. Because it improves speed.

Answer: B

Explanation

Concentration amplifies correlated risk. Options A, C, and D are benefits.

Question

Which is a reputational risk from AI?

- A. Faster delivery.
- B. Public backlash from biased, unethical, or incompetent AI use, resulting in loss of trust and brand value.
- C. Lower cost.
- D. Higher accuracy.

Answer: B

Explanation

Reputational risk stems from public perception of AI failures. Options A, C, and D are benefits.

Question

Which is a legal risk from AI?

- A. Lower cost.
- B. Liability for harm caused by AI decisions, non-compliance with regulations, and IP violations.
- C. Faster deployment.
- D. Higher accuracy.

Answer: B

Explanation

Legal risks include liability, non-compliance, and IP issues. Options A, C, and D are benefits.

Question

What environmental risk does AI present?

- A. None.
- B. Significant compute-related energy consumption and carbon emissions from training and inference of large models.
- C. Lower energy use.
- D. Positive environmental impact only.

Answer: B

Explanation

AI compute consumption creates environmental risk. Options A, C, and D are incorrect.

Question

Why is regulatory fragmentation a challenge for AI governance?

- A. Because it speeds up compliance.
- B. Because different jurisdictions impose different rules, complicating compliance for globally operating firms.
- C. Because it reduces cost.
- D. Because it simplifies everything.

Answer: B

Explanation

Fragmented regulation complicates global compliance. Options A, C, and D are incorrect.

Question

Which is the first line's responsibility for AI models?

- A. Independent validation.
- B. Design, develop, own, and operate AI models; monitor performance and own data quality.
- C. Set policy.
- D. Audit.

Answer: B

Explanation

First line owns AI models and operations. Options A, C, and D belong to second or third lines.

Question

Which is the second line's responsibility for AI models?

- A. Build AI models.
- B. Independent validation, policy, and challenge of first-line assumptions and controls.
- C. Use models for decisions.
- D. Market models.

Answer: B

Explanation

Second line validates, sets policy, and challenges. Options A, C, and D belong to first line.

Question

Which is the third line's responsibility for AI models?

- A. Build models.
- B. Provide independent assurance on the effectiveness of AI governance and controls.
- C. Set risk appetite.
- D. Manage model data.

Answer: B

Explanation

Third line provides independent assurance. Options A, C, and D belong to first or second lines.

Question

Why must AI validation be independent?

- A. To speed up processes.
- B. To avoid conflicts of interest and ensure objective assessment of model quality and risk.
- C. To reduce costs.
- D. To bypass policies.

Answer: B

Explanation

Independence ensures objective validation. Options A, C, and D are incorrect.

Question

What does effective challenge mean in the AI governance context?

- A. Automatic approval.
- B. Critical, objective, and informed review of AI models and their assumptions by qualified parties, independent from developers.
- C. Marketing endorsement.
- D. Rubber-stamping.

Answer: B

Explanation

Effective challenge is critical review. Options A, C, and D are inconsistent.

Question

How should AI risk appetite cascade through the organisation?

- A. It shouldn't.
- B. From board-set appetite to business-unit tolerances and model-level thresholds, all aligned.
- C. Only at business-unit level.
- D. Only at IT level.

Answer: B

Explanation

Appetite cascades from board to units to models. Options A, C, and D are too narrow.

Question

Who owns AI risk in an organisation?

- A. The IT department.
- B. Business units that use AI, supported by risk, compliance, and other functions.
- C. Only the CRO.
- D. Only external vendors.

Answer: B

Explanation

AI risk ownership sits with business users supported by risk functions. Options A, C, and D are incomplete.

Question

What is the purpose of AI escalation paths?

- A. To slow down decisions.
- B. To ensure that AI incidents, breaches, and issues are reported and addressed at the appropriate level promptly.
- C. To reduce transparency.
- D. To bypass governance.

Answer: B

Explanation

Escalation paths ensure timely, appropriate responses. Options A, C, and D are contrary.

Question

What does AI risk aggregation mean?

- A. Averaging individual risks.
- B. Combining individual AI risks into an enterprise-wide view, accounting for interdependencies and correlations.
- C. Ignoring enterprise risk.
- D. Removing individual risks.

Answer: B

Explanation

Aggregation considers interdependencies. Options A, C, and D misstate.

Question

Which is a good metric for AI governance success?

- A. Number of meetings.
- B. Reduction in AI incidents, improvements in validation timeliness, and increasing compliance with governance standards.
- C. Number of documents.
- D. Cost of IT.

Answer: B

Explanation

Governance success metrics track incidents, timeliness, and compliance. Options A, C, and D are not meaningful metrics.

Question

Under GDPR, data subjects can request:

- A. Only access to data.
- B. Access, rectification, erasure, restriction, portability, and objection.
- C. Only rectification.
- D. Nothing.

Answer: B

Explanation

GDPR grants multiple rights. Options A, C, and D understate or misstate.

Question

What is the difference between a data controller and a data processor under GDPR?

- A. They are the same.
- B. Controller determines the purposes and means of processing; processor processes data on behalf of the controller.
- C. Controller only stores data.
- D. Processor only collects data.

Answer: B

Explanation

Controller decides; processor acts on behalf. Options A, C, and D misstate.

Question

What is the DPO's role under GDPR?

- A. Marketing.
- B. Independent oversight of data protection compliance, advising on GDPR obligations, and liaising with the supervisory authority.
- C. Sales.
- D. IT support.

Answer: B

Explanation

The DPO provides independent compliance oversight. Options A, C, and D are unrelated.

Question

What is a DPIA?

- A. A marketing report.
- B. A Data Protection Impact Assessment required for processing likely to result in high risk to data subjects, documenting risks and mitigations.
- C. A code review.
- D. A financial audit.

Answer: B

Explanation

DPIA documents high-risk processing and mitigations. Options A, C, and D misstate.

Question

Which AI systems are classified as high-risk under the EU AI Act?

- A. All chatbots.
- B. Those affecting safety, fundamental rights, or critical infrastructure, such as biometrics, employment, credit scoring, and law enforcement.
- C. Only medical devices.
- D. Only toys.

Answer: B

Explanation

High-risk systems affect safety, rights, or critical infrastructure. Options A, C, and D are narrower or incorrect.

Question

What transparency obligations does the EU AI Act impose?

- A. None.
- B. Notifying users when interacting with AI systems, disclosing deepfakes, and providing information about AI-generated content.
- C. Only financial disclosure.
- D. Only marketing disclosure.

Answer: B

Explanation

The AI Act imposes transparency obligations. Options A, C, and D are incorrect.

Question

What does the EU AI Act require of general-purpose AI (GPAI) models?

- A. Nothing.
- B. Documentation, transparency on training data, copyright compliance, and safety measures, with stricter obligations for models with systemic risk.
- C. Only marketing disclosure.
- D. Only to register.

Answer: B

Explanation

GPAI models have documentation, transparency, and safety obligations. Options A, C, and D understate or misstate.

Question

How is the EU AI Act being implemented over time?

- A. All at once.
- B. In phases: prohibitions first, then high-risk rules, then remaining provisions, with transitional periods for existing systems.
- C. Only in 2030.
- D. Never.

Answer: B

Explanation

The AI Act phases in obligations. Options A, C, and D are inaccurate.

Question

What rights does CCPA grant California consumers?

- A. No rights.
- B. Right to know, delete, opt out of sale, correct, and non-discrimination.
- C. Only access.
- D. Only deletion.

Answer: B

Explanation

CCPA grants multiple rights. Options A, C, and D are incomplete.

Question

What is the purpose of China's Personal Information Protection Law (PIPL)?

- A. To ban personal data.
- B. To protect personal data, requiring consent, data minimisation, transparency, and specific rules for cross-border transfer.
- C. Only for government data.
- D. Only for healthcare data.

Answer: B

Explanation

PIPL protects personal data with consent, minimisation, and transfer rules. Options A, C, and D are inaccurate.

Question

What evidence should model validation produce?

- A. None.
- B. Documented assessment of conceptual soundness, data appropriateness, implementation, performance, and limitations, with findings and recommendations.
- C. Only accuracy metrics.
- D. Only source code.

Answer: B

Explanation

Validation produces documented, comprehensive evidence. Options A, C, and D are insufficient.

Question

Why does validation include benchmarking against challenger models?

- A. To confuse stakeholders.
- B. To compare performance against alternatives, identifying whether the model is the best choice for its purpose.
- C. To reduce cost.
- D. To speed up training.

Answer: B

Explanation

Benchmarking compares models for suitability. Options A, C, and D are incorrect.

Question

What does sensitivity analysis in model validation examine?

- A. Only cost.
- B. How model outputs respond to changes in input assumptions, parameters, or data, evaluating robustness.
- C. Only speed.
- D. Only storage.

Answer: B

Explanation

Sensitivity analysis tests robustness to changes. Options A, C, and D are unrelated.

Question

What is the purpose of stress testing AI models?

- A. To make them faster.
- B. To evaluate their behaviour under extreme, adverse, or unlikely scenarios.
- C. To reduce cost.
- D. To simplify them.

Answer: B

Explanation

Stress tests evaluate behaviour under extremes. Options A, C, and D are unrelated.

Question

What does backtesting involve for AI models?

- A. Forward-looking simulation only.
- B. Testing the model against historical data to evaluate how it would have performed.
- C. Marketing analysis.
- D. Cost analysis.

Answer: B

Explanation

Backtesting evaluates historical performance. Options A, C, and D are unrelated.

Question

Outcome analysis in model validation is used to:

- A. Reduce costs.
- B. Compare model predictions against realised outcomes to detect bias, degradation, or incorrect assumptions.
- C. Speed up training.
- D. Increase accuracy.

Answer: B

Explanation

Outcome analysis compares predictions to reality. Options A, C, and D are unrelated.

Question

What evidence should the use test provide?

- A. None.
- B. Qualitative evidence that the model is used appropriately, understood by users, and aligned with business decisions.
- C. Only accuracy metrics.
- D. Only cost data.

Answer: B

Explanation

The use test produces qualitative evidence. Options A, C, and D are insufficient.

Question

How should model use restrictions be enforced?

- A. Through documentation only.
- B. Through documentation, access controls, system limits, and monitoring.
- C. By ignoring them.
- D. By removing them.

Answer: B

Explanation

Use restrictions require documentation, controls, limits, and monitoring. Options A, C, and D are insufficient or contrary.

Question

What documentation should accompany model validation?

- A. Nothing.
- B. A validation report with methodology, findings, limitations, and recommendations, plus supporting analyses.
- C. Only the model code.
- D. Only marketing material.

Answer: B

Explanation

Validation documentation is comprehensive. Options A, C, and D are insufficient or inappropriate.

Question

Which best summarises the guiding principle for responsible AI governance?

- A. Maximise complexity.
- B. Balance performance with governance, validation, monitoring, fairness, explainability, and ethical alignment, consistent with regulatory and business expectations.
- C. Deploy without validation.
- D. Ignore risks.

Answer: B

Explanation

Responsible AI balances performance and governance across all dimensions. Options A, C, and D contradict the principle.