

GARP Risk & AI — Q551–Q650

Unofficial Exam Prep

Unofficial

October 3, 2026

Question

How does an AI portfolio differ from a model inventory?

- A. The portfolio and inventory are functionally isomorphic artefacts, both maintaining versioned metadata schemas, ownership attestations, and validation-state enumerations across the model estate.
- B. The portfolio is a strategic capability map spanning build-versus-buy decisions, capital allocation, and value-realisation roadmaps, whereas the inventory is an operational registry of deployed artefacts with lineage, drift telemetry, and lifecycle-state metadata.
- C. The portfolio is scoped exclusively to third-party procured solutions governed by vendor risk schedules, while the inventory is restricted to internally developed models that have entered post-retirement archival states.
- D. The inventory is a forward-looking capability roadmap with investment theses, while the portfolio is a real-time inference-call ledger capturing latency percentiles, throughput counters, and token-consumption metrics.

Answer: B

Explanation

Portfolio is strategic; inventory is operational. Options A, C, and D misstate the distinction.

Question

Which is NOT typically included in board-level AI reporting?

- A. Aggregate AI risk exposure, including concentration indices, interdependency graphs, and enterprise risk-appetite utilisation across business units.
- B. Fairness and bias metrics, including demographic parity ratios, equalised-odds differentials, and disparate-impact analyses stratified by protected class.
- C. Detailed source code of every model, including line-by-line comments, hyperparameter configurations, gradient histograms, and raw training-data samples.
- D. Material incidents and remediation plans, including root-cause analyses, regulatory notifications, and remediation timelines with accountable owners.

Answer: C

Explanation

Board reports are strategic and aggregate; source code is not appropriate. Options A, B, and D are typically included.

Question

Which is a common category in AI risk taxonomies?

- A. Technical risk, encompassing adversarial robustness, distributional shift, data-provenance integrity, and reproducibility variance.
- B. Ethical risk, encompassing allocative harm, representational harm, transparency deficits, and human-oversight sufficiency.
- C. Operational risk, encompassing third-party concentration, model-drift telemetry, process-failure modes, and resource-exhaustion scenarios.
- D. All of the above are common categories in AI risk taxonomies, alongside legal, regulatory, and reputational dimensions.

Answer: D

Explanation

AI risk taxonomies span technical, ethical, operational, and other categories; all are common.

Question

Which is the most relevant model risk for an AI-based credit scoring system?

- A. The colour scheme and layout of the user interface may not align with brand guidelines, leading to inconsistent customer experience and reduced engagement metrics.
- B. Potential unfairness, inaccuracy, and adverse regulatory consequences from biased or poorly performing predictions, including proxy discrimination, explainability gaps, and ECOA/FCRA non-compliance.
- C. The storage cost of retaining historical training data may exceed budgeted cloud infrastructure allocations over multi-year horizons, impacting FinOps unit economics.
- D. Marketing effectiveness may decline if the model's output is not integrated with campaign management platforms, customer segmentation tools, and CRM workflows.

Answer: B

Explanation

Credit scoring model risk includes fairness, accuracy, and regulatory risk. Options A, C, and D are secondary or irrelevant.

Question

For an AI fraud detection model, which is the most critical model risk?

- A. Overfitting to historical fraud patterns, missing novel fraud schemes, and failing to adapt to adversarial concept drift under non-stationary distributions.
- B. Storage cost of retaining transaction-level features and model artefacts across multiple retraining cycles, impacting data-lake tiering strategies.
- C. User interface design that may not clearly present fraud alerts to analysts, slowing investigation workflows and increasing mean-time-to-resolution.
- D. Training time may exceed operational windows, delaying model refresh cycles and reducing responsiveness to emerging fraud typologies.

Answer: A

Explanation

Overfitting to past patterns can miss novel fraud. Options B, C, and D are less critical.

Question

Which is a key risk for an AI anti-money-laundering (AML) model?

- A. Low alert volume may indicate under-detection, reducing investigative workload but increasing regulatory exposure under BSA/AML regimes.
- B. High false-positive rate overwhelming compliance teams and potential missed true positives, both creating regulatory and operational risk.
- C. Storage cost of retaining alert histories and case management records across multiple jurisdictions, impacting data-retention compliance.
- D. Marketing effectiveness may decline if AML alerts are not integrated with customer relationship management platforms and case-management systems.

Answer: B

Explanation

AML risk includes excessive false positives and missed true positives. Options A, C, and D are secondary.

Question

Which is a typical risk for an AI-enhanced market risk model?

- A. Training on non-stationary data leading to unstable predictions and regime-change failures under structural breaks.
- B. Data storage costs may escalate if tick-level market data is retained for extended backtesting periods, impacting data-lake economics.
- C. Logo design may not comply with brand guidelines, creating inconsistent visual identity across risk reports and dashboards.
- D. Employee satisfaction may decline if model outputs are not integrated with existing workflow tools and trading-desk platforms.

Answer: A

Explanation

Non-stationarity of financial markets is a key challenge. Options B, C, and D are unrelated.

Question

Which is a specific concern for AI underwriting in insurance?

- A. Ethical and legal issues from using alternative data sources that may be considered sensitive or discriminatory under unfair-discrimination statutes.
- B. Faster approvals may reduce customer wait times but could increase adverse selection if not properly controlled via risk-based pricing.
- C. Lower costs may improve competitiveness but could reduce investment in human oversight and quality assurance.
- D. Better customer experience may increase retention but could mask underlying model deficiencies and bias.

Answer: A

Explanation

Alternative data raises ethical and legal concerns. Options B, C, and D are benefits, not concerns.

Question

Which is a recognised fairness metric used in lending AI?

- A. Demographic parity, measured as the ratio of positive prediction rates between protected and non-protected groups.
- B. Storage cost parity, measured as the ratio of data storage costs between protected and non-protected groups.
- C. Latency parity, measured as the ratio of inference times between protected and non-protected groups.
- D. Colour parity, measured as the ratio of UI colour contrast ratios between protected and non-protected groups.

Answer: A

Explanation

Demographic parity is a standard fairness metric. Options B, C, and D are not.

Question

Which is a key risk of algorithmic trading models?

- A. Flash-crash amplification, unintended feedback loops, and market impact from correlated strategies under stressed liquidity conditions.
- B. Storage capacity may be insufficient to retain tick-level data for regulatory reporting and backtesting, impacting MiFID II compliance.
- C. Logo design may not comply with exchange branding requirements, creating regulatory filing issues and venue-certification delays.
- D. Employee benefits may not be competitive, leading to talent attrition in quantitative teams and loss of institutional knowledge.

Answer: A

Explanation

Algorithmic trading risks include flash crashes and correlated behaviour. Options B, C, and D are unrelated.

Question

Which is the most significant risk from deploying a GenAI chatbot in banking?

- A. Generating incorrect or non-compliant advice to customers, creating legal and reputational risk under consumer-protection regimes.
- B. Faster response times may reduce customer wait times but could increase volume of unverified interactions and hallucination exposure.
- C. Cost savings may improve margins but could reduce investment in human oversight and quality assurance.
- D. Customer satisfaction may increase but could mask underlying model deficiencies and bias.

Answer: A

Explanation

Advice quality and compliance are central concerns. Options B, C, and D are benefits.

Question

Which is a key governance concern when using alternative data in financial AI?

- A. Legal right to use the data, privacy, and potential for discriminatory impact through proxy variables and disparate-impact mechanisms.
- B. Storage cost may exceed budgeted cloud infrastructure allocations over multi-year horizons, impacting FinOps unit economics.
- C. Model speed may decline if alternative data requires extensive preprocessing and feature engineering pipelines.
- D. Logo design may not comply with brand guidelines, creating inconsistent visual identity across risk reports.

Answer: A

Explanation

Legal rights, privacy, and fairness are key governance concerns. Options B, C, and D are secondary.

Question

What is a specific AI risk in AML/KYC?

- A. Customer friction from over-flagging and regulatory exposure from under-flagging, requiring careful threshold calibration under BSA/AML regimes.
- B. Storage cost of retaining KYC documents and alert histories across multiple jurisdictions, impacting data-retention compliance.
- C. Logo design may not comply with brand guidelines, creating inconsistent visual identity across compliance reports.
- D. Training time may exceed operational windows, delaying model refresh cycles and reducing responsiveness to emerging typologies.

Answer: A

Explanation

Threshold calibration affects both customer experience and regulatory compliance. Options B, C, and D are less critical.

Question

Which is the correct order of model lifecycle stages?

- A. Development, validation, deployment, monitoring, retirement.
- B. Deployment, development, monitoring, validation, retirement.
- C. Monitoring, development, retirement, validation, deployment.
- D. Retirement, development, deployment, monitoring, validation.

Answer: A

Explanation

Development → validation → deployment → monitoring → retirement. Options B, C, and D reorder incorrectly.

Question

Why must validation occur before deployment?

- A. Regulatory convenience, as validation timelines are aligned with regulatory reporting cycles and supervisory review schedules.
- B. To independently confirm the model meets requirements, performs as expected, and is fit for purpose before being used in production.
- C. To reduce cost, as validation is cheaper than post-deployment remediation and incident response.
- D. To increase speed, as validation accelerates deployment timelines through early defect detection.

Answer: B

Explanation

Pre-deployment validation confirms fitness. Options A, C, and D are secondary.

Question

Which activity is part of post-deployment monitoring?

- A. Continuous tracking of model performance, data drift, and input integrity, with alerts for threshold breaches and automated escalation.
- B. Only training, as monitoring is limited to retraining cycles and periodic model refreshes.
- C. Only validation, as monitoring is limited to annual validation reviews and independent assessment.
- D. Only code review, as monitoring is limited to static analysis and vulnerability scanning.

Answer: A

Explanation

Post-deployment monitoring tracks performance and drift. Options B, C, and D are pre-deployment.

Question

Recalibration is preferred to redevelopment when:

- A. The model structure is wrong, requiring fundamental redesign and re-engineering of feature pipelines.
- B. The model structure remains valid but parameters need updating using new data under stable distributional assumptions.
- C. The model is retired, as recalibration is only applicable to post-retirement archival models.
- D. The model has never been validated, as recalibration replaces validation and independent review.

Answer: B

Explanation

Recalibration updates parameters; structural issues require redevelopment. Options A, C, and D are different situations.

Question

What is model deprecation?

- A. Immediate retirement, as deprecation is synonymous with decommissioning and archival.
- B. The formal marking of a model as no longer recommended for use, typically while a replacement is being prepared.
- C. Model creation, as deprecation is the initial stage of model development and experimentation.
- D. Model validation, as deprecation is a validation outcome and independent review artefact.

Answer: B

Explanation

Deprecation flags models for phase-out. Options A, C, and D describe different actions.

Question

What does a model change impact assessment evaluate?

- A. Only cost, as impact assessment is limited to financial considerations and FinOps unit economics.
- B. How a proposed change affects model outputs, downstream users, risk profile, and regulatory compliance.
- C. Only speed, as impact assessment is limited to latency considerations and inference-time optimisation.
- D. Only storage, as impact assessment is limited to infrastructure considerations and data-lake tiering.

Answer: B

Explanation

Impact assessment covers outputs, downstream, risk, and compliance. Options A, C, and D are too narrow.

Question

Why is model versioning important?

- A. To reduce storage, as versioning eliminates redundant model artefacts and deduplicates parameter checkpoints.
- B. To track distinct model iterations, enabling reproducibility, audit, and controlled rollback if needed.
- C. To reduce cost, as versioning eliminates redundant training runs and optimises compute allocation.
- D. To speed up training, as versioning accelerates model refresh cycles and hyperparameter search.

Answer: B

Explanation

Versioning supports reproducibility, audit, and rollback. Options A, C, and D are secondary.

Question

What is the champion/challenger framework?

- A. A marketing strategy, as champion/challenger is used for brand positioning and customer-segment targeting.
- B. A framework where the current production model (champion) is compared against alternative models (challengers) to evaluate whether a switch would improve performance.
- C. A training technique, as champion/challenger is used for hyperparameter tuning and neural architecture search.
- D. A validation report, as champion/challenger is a documentation artefact and independent review deliverable.

Answer: B

Explanation

Champion/challenger compares production models to alternatives. Options A, C, and D are unrelated.

Question

What is shadow deployment of a model?

- A. Deploying without monitoring, as shadow deployment bypasses monitoring requirements and observability instrumentation.
- B. Running a new model alongside production without using its outputs, to evaluate performance in real conditions.
- C. Deploying to production only, as shadow deployment is a production deployment mode with canary routing.
- D. Retiring a model, as shadow deployment is a decommissioning technique and archival strategy.

Answer: B

Explanation

Shadow deployment evaluates models in parallel without affecting production. Options A, C, and D misstate.

Question

What is A/B testing in the AI deployment context?

- A. Randomly assigning users to two model versions and comparing outcomes under controlled experimental conditions.
- B. Retiring both models, as A/B testing is a decommissioning technique and archival strategy.
- C. Doubling training data, as A/B testing is a data augmentation technique and synthetic-data generation method.
- D. Storing data twice, as A/B testing is a storage redundancy technique and disaster-recovery strategy.

Answer: A

Explanation

A/B testing compares model versions via controlled experiments. Options B, C, and D are unrelated.

Question

Which is a key element of an AI incident response plan?

- A. Ignoring incidents, as incident response is limited to post-mortem documentation and retrospective analysis.
- B. Defined escalation paths, roles, playbooks, and communication plans for handling AI-related incidents.
- C. Only post-mortems, as incident response is limited to retrospective analysis and lessons-learned exercises.
- D. Only IT involvement, as incident response is limited to infrastructure teams and site-reliability engineers.

Answer: B

Explanation

An incident response plan includes escalation, roles, playbooks, and communication. Options A, C, and D are insufficient.

Question

An AI model begins producing biased outputs. Which is the most likely root cause?

- A. Data drift or newly biased data, as the model's input distribution has shifted under non-stationary conditions.
- B. Server failure, as hardware faults can introduce bias into model outputs through corrupted tensor operations.
- C. Power outage, as interrupted training can introduce bias into model weights through checkpoint corruption.
- D. UI redesign, as changes to the user interface can introduce bias into model inputs through altered data-capture flows.

Answer: A

Explanation

Bias in outputs typically stems from data or model development. Options B, C, and D are unrelated.

Question

What is the purpose of a post-incident review for AI incidents?

- A. Assign blame, as post-incident reviews are primarily disciplinary processes and accountability mechanisms.
- B. Learn from the incident, improve controls and processes, and prevent recurrence.
- C. End investigation, as post-incident reviews are terminal documentation exercises and closure artefacts.
- D. Skip documentation, as post-incident reviews are informal debriefs and verbal knowledge-transfer sessions.

Answer: B

Explanation

Post-incident reviews focus on learning and prevention. Options A, C, and D are contrary.

Question

Why is communication important during an AI incident?

- A. It is not important, as incident response is purely technical and does not require stakeholder engagement.
- B. Clear communication with regulators, customers, and stakeholders maintains trust and meets regulatory obligations.
- C. Only internal communication matters, as external parties are not affected by AI incidents.
- D. Only external communication matters, as internal teams are not affected by AI incidents.

Answer: B

Explanation

Communication with all stakeholders is essential. Options A, C, and D are incomplete or wrong.

Question

How are AI incidents typically classified?

- A. By colour, as incident classification uses a colour-coded severity scale and RAG status indicators.
- B. By severity and impact, considering factors such as harm to customers, financial loss, and regulatory exposure.
- C. Randomly, as incident classification is not standardised and varies by team discretion.
- D. By alphabet, as incident classification uses alphabetical ordering and lexical sorting.

Answer: B

Explanation

Incidents are classified by severity and impact. Options A, C, and D are not standard.

Question

When a serious AI incident affects customers, what is typically required?

- A. No action, as AI incidents are not subject to regulatory notification and supervisory reporting.
- B. Prompt notification to regulators where required and remediation plans.
- C. Only internal review, as regulatory notification is optional and discretionary.
- D. Silence, as public disclosure is prohibited and confidentiality is paramount.

Answer: B

Explanation

Serious incidents may trigger regulatory notification. Options A, C, and D violate governance.

Question

Which metric helps track AI incident management effectiveness?

- A. Total storage used, as storage metrics correlate with incident management effectiveness and data-retention efficiency.
- B. Time to detection, time to remediation, recurrence rate, and impact.
- C. Marketing reach, as customer awareness correlates with incident management effectiveness and brand perception.
- D. UI clicks, as user engagement correlates with incident management effectiveness and interface usability.

Answer: B

Explanation

Detection, remediation, recurrence, and impact are key metrics. Options A, C, and D are unrelated.

Question

What is red teaming in AI?

- A. A marketing exercise, as red teaming is used for brand positioning and competitive differentiation.
- B. Adversarial testing by a dedicated team to uncover vulnerabilities and unintended behaviours before deployment.
- C. A financial audit, as red teaming is used for financial controls testing and SOX compliance.
- D. A user survey, as red teaming is used for customer feedback collection and satisfaction measurement.

Answer: B

Explanation

Red teaming proactively tests for vulnerabilities. Options A, C, and D misstate.

Question

Why is logging data important for AI incident management?

- A. To increase storage, as logging is primarily a storage optimization technique and data-lake tiering strategy.
- B. To reconstruct events, identify root causes, and support regulatory reporting.
- C. To slow down systems, as logging introduces latency for diagnostic purposes and observability instrumentation.
- D. To increase costs, as logging is primarily a cost optimization technique and FinOps unit-economics lever.

Answer: B

Explanation

Logging supports root-cause analysis and reporting. Options A, C, and D are incorrect or secondary.

Question

What is the purpose of continuous improvement in AI governance?

- A. To increase costs, as continuous improvement is primarily a cost optimization technique and FinOps unit-economics lever.
- B. To iterate on governance frameworks and controls based on incidents, audits, and evolving best practices.
- C. To avoid change, as continuous improvement is primarily a stability technique and change-management control.
- D. To slow down deployment, as continuous improvement is primarily a risk mitigation technique and deployment-gating mechanism.

Answer: B

Explanation

Continuous improvement ensures governance keeps pace with change. Options A, C, and D are contrary.

Question

A typical AI governance maturity model includes levels such as:

- A. Ad hoc, developing, defined, managed, optimised.
- B. Red, blue, green, as maturity models use colour-coded levels and RAG status indicators.
- C. Only one level, as maturity models are binary and non-hierarchical.
- D. No levels, as maturity models are non-hierarchical and network-based.

Answer: A

Explanation

Maturity models typically progress from ad hoc to optimised. Options B, C, and D are not standard.

Question

Which is a common gap in AI governance frameworks?

- A. Extensive controls, as frameworks often over-control low-risk use cases and create unnecessary friction.
- B. Missing governance for GenAI, alternative data, or third-party models.
- C. Complete coverage, as frameworks often over-cover all risk dimensions and create unnecessary complexity.
- D. Perfect alignment, as frameworks often over-align with business objectives and create unnecessary constraints.

Answer: B

Explanation

Common gaps include GenAI, alternative data, and vendor models. Options A, C, and D are not gaps.

Question

What is the purpose of an AI governance roadmap?

- A. A marketing document, as roadmaps are used for brand positioning and competitive differentiation.
- B. A phased plan outlining how the organisation will build, enhance, and sustain AI governance capabilities over time.
- C. A single policy, as roadmaps are policy documents and compliance artefacts.
- D. A code repository, as roadmaps are technical artefacts and infrastructure-as-code modules.

Answer: B

Explanation

A roadmap is a phased plan. Options A, C, and D misstate.

Question

Which is an appropriate KPI for AI governance?

- A. Number of employees, as headcount correlates with governance effectiveness and organisational capacity.
- B. Percentage of models validated on schedule.
- C. Office size, as physical footprint correlates with governance effectiveness and operational scale.
- D. Marketing spend, as brand awareness correlates with governance effectiveness and market positioning.

Answer: B

Explanation

Validation timeliness is a governance KPI. Options A, C, and D are unrelated.

Question

Which audience requires AI governance training?

- A. Only data scientists, as governance training is limited to technical roles and model-development functions.
- B. A broad set including executives, developers, validators, business users, compliance, and internal audit.
- C. Only IT, as governance training is limited to infrastructure teams and site-reliability engineers.
- D. Only external auditors, as governance training is limited to assurance providers and independent reviewers.

Answer: B

Explanation

AI governance training spans many roles. Options A, C, and D are too narrow.

Question

Why is data collection for AI governance metrics important?

- A. To meet marketing goals, as governance metrics are used for brand positioning and competitive differentiation.
- B. To enable measurement, benchmarking, and continuous improvement of AI governance practices.
- C. To increase costs, as governance metrics are primarily a cost optimization technique and FinOps unit-economics lever.
- D. To slow operations, as governance metrics are primarily a risk mitigation technique and deployment-gating mechanism.

Answer: B

Explanation

Data supports measurement and improvement. Options A, C, and D are incorrect.

Question

What is the purpose of AI governance certifications (e.g., ISO/IEC 42001)?

- A. Marketing only, as certifications are used for brand positioning and competitive differentiation.
- B. To demonstrate externally that the organisation meets recognised standards for AI governance.
- C. To replace regulations, as certifications are legally binding substitutes and safe-harbour mechanisms.
- D. To lower costs, as certifications are primarily a cost optimization technique and FinOps unit-economics lever.

Answer: B

Explanation

Certifications demonstrate conformance with standards. Options A, C, and D are not the primary purpose.

Question

Why is board AI literacy important?

- A. To increase headcount, as board literacy correlates with hiring needs and organisational capacity.
- B. To enable effective oversight of AI strategy, risk, and governance at the highest level.
- C. To reduce costs, as board literacy correlates with cost efficiency and FinOps unit economics.
- D. To speed deployment, as board literacy correlates with deployment velocity and time-to-market.

Answer: B

Explanation

Board AI literacy enables effective oversight. Options A, C, and D are not the primary reason.

Question

How should an AI ethics committee integrate with the model risk function?

- A. No integration is needed, as ethics and model risk are independent domains and separate governance tracks.
- B. Through coordinated review of high-risk AI use cases, sharing of insights on fairness and ethics, and joint escalation to senior management.
- C. Only through annual meetings, as integration is limited to periodic reviews and annual attestations.
- D. Only through external advisors, as integration is outsourced to third parties and independent consultants.

Answer: B

Explanation

Coordination between ethics and model risk is essential. Options A, C, and D are insufficient.

Question

Why must AI governance frameworks evolve?

- A. To increase cost, as evolution is primarily a cost optimization technique and FinOps unit-economics lever.
- B. Because AI technologies, use cases, and regulations change rapidly, requiring ongoing updates to remain effective.
- C. To slow operations, as evolution is primarily a risk mitigation technique and deployment-gating mechanism.
- D. To reduce transparency, as evolution is primarily a confidentiality technique and information-security control.

Answer: B

Explanation

Rapid change in AI and regulation requires evolving governance. Options A, C, and D contradict this.

Question

How can GenAI undermine human autonomy?

- A. By providing choices, as GenAI expands the option set available to users and increases decision-space dimensionality.
- B. By generating persuasive, tailored content that manipulates users without their awareness.
- C. By explaining decisions, as GenAI provides transparency into decision processes and interpretability artefacts.
- D. By slowing systems, as GenAI introduces latency that reduces user control and increases interaction friction.

Answer: B

Explanation

Manipulative content undermines autonomy. Options A, C, and D respect or are unrelated to autonomy.

Question

How does GenAI amplify bias?

- A. By generating content that reflects biases present in training data, at scale.
- B. By reducing accuracy, as lower accuracy correlates with higher bias and increased error rates.
- C. By slowing training, as longer training correlates with higher bias and overfitting to minority classes.
- D. By increasing costs, as higher costs correlate with higher bias and increased computational complexity.

Answer: A

Explanation

GenAI can amplify training-data bias at scale. Options B, C, and D are unrelated.

Question

Why is GenAI a challenge for truth and accountability?

- A. It only outputs truth, as GenAI is designed for factual accuracy and grounded generation.
- B. It generates plausible but sometimes incorrect content, making accountability for misinformation difficult.
- C. It has no training data, as GenAI is designed for zero-shot learning and few-shot generalisation.
- D. It never hallucinates, as GenAI is designed for factual accuracy and retrieval-augmented grounding.

Answer: B

Explanation

Hallucinations and plausible output complicate accountability. Options A, C, and D are incorrect.

Question

What consent issues arise with GenAI training data?

- A. None, as GenAI training data is always consented and properly licensed.
- B. Data subjects whose content was used in training may not have consented, raising ethical and legal concerns.
- C. Consent is automatically granted, as GenAI training data is publicly available and web-scraped.
- D. Consent is not required, as GenAI training data is exempt from consent requirements under fair-use doctrines.

Answer: B

Explanation

Training data consent is a critical concern. Options A, C, and D are incorrect.

Question

Why is attribution important in GenAI?

- A. It isn't important, as GenAI outputs are not attributable and are generated de novo.
- B. To trace sources of content and information, supporting transparency and possibly legal rights such as copyright.
- C. To increase cost, as attribution is primarily a cost optimization technique and FinOps unit-economics lever.
- D. To slow inference, as attribution is primarily a latency optimization technique and inference-time acceleration method.

Answer: B

Explanation

Attribution supports transparency and copyright compliance. Options A, C, and D are incorrect.

Question

What is a copyright concern with GenAI?

- A. None, as GenAI outputs are always copyright-clean and free of encumbrances.
- B. Training on copyrighted material without permission and generating content that may infringe copyright.
- C. GenAI is always copyright-clean, as training data is licensed and properly attributed.
- D. Copyright does not apply, as GenAI outputs are not fixed in a tangible medium and are ephemeral.

Answer: B

Explanation

Copyright concerns are significant. Options A, C, and D are incorrect.

Question

Which is a safety concern with GenAI in critical applications?

- A. Faster response, as speed reduces safety in critical applications and increases error propagation.
- B. Generating unsafe or harmful content that could cause physical, psychological, or financial harm.
- C. Lower costs, as cost reduction reduces safety in critical applications and increases technical debt.
- D. Higher accuracy, as accuracy reduces safety in critical applications and increases overconfidence.

Answer: B

Explanation

Harmful content generation is a safety concern. Options A, C, and D are benefits.

Question

Which consumer protection issue arises from GenAI?

- A. Increasing prices only, as GenAI directly causes price increases and inflationary pressure.
- B. Misleading or inaccurate information provided to consumers, potentially leading to harm.
- C. Faster service, as speed reduces consumer protection and increases error propagation.
- D. Lower fees, as fee reduction reduces consumer protection and increases hidden costs.

Answer: B

Explanation

Misleading information is a consumer protection issue. Options A, C, and D are unrelated.

Question

What regulatory concern arises when GenAI is used for financial advice?

- A. Lower costs, as cost reduction reduces regulatory compliance and increases risk-taking.
- B. Potential non-compliance with advice regulations if unvetted or unsuitable advice is provided.
- C. Faster response, as speed reduces regulatory compliance and increases error propagation.
- D. Higher throughput, as throughput increases regulatory compliance and reduces oversight.

Answer: B

Explanation

Regulatory concerns exist around advice suitability. Options A, C, and D are benefits, not concerns.

Question

How can GenAI facilitate fraud?

- A. By making it easier to create convincing phishing content and deepfakes at scale.
- B. By improving accuracy, as accuracy reduces fraud and increases detection rates.
- C. By reducing costs, as cost reduction reduces fraud and increases detection efficiency.
- D. By speeding up transactions, as speed reduces fraud and increases detection velocity.

Answer: A

Explanation

GenAI enables scalable, convincing fraud content. Options B, C, and D are unrelated.

Question

What are the key risks of agentic AI?

- A. Unintended autonomous actions, difficulty in oversight, and cascading effects from chains of decisions.
- B. Faster decisions, as speed reduces oversight requirements and increases error propagation.
- C. Lower cost, as cost reduction reduces oversight requirements and increases technical debt.
- D. Simplicity, as simplicity reduces oversight requirements and increases hidden complexity.

Answer: A

Explanation

Agentic AI risks relate to autonomy and oversight. Options B, C, and D are benefits or inaccurate.

Question

What risk is specific to multi-agent AI systems?

- A. Faster training, as speed reduces coordination risk and increases emergent complexity.
- B. Emergent behaviour arising from interactions between agents, which may be difficult to predict or control.
- C. Lower costs, as cost reduction reduces coordination risk and increases technical debt.
- D. Higher accuracy, as accuracy reduces coordination risk and increases overconfidence.

Answer: B

Explanation

Emergent multi-agent behaviour is a specific risk. Options A, C, and D are benefits.

Question

How does AI increase cyber risk?

- A. By reducing attack surface, as AI simplifies infrastructure and reduces vulnerability exposure.
- B. By enabling more sophisticated, scalable attacks, and by introducing new vulnerabilities (e.g., model extraction, adversarial inputs).
- C. By eliminating threats, as AI replaces traditional security controls and reduces human error.
- D. By lowering costs, as AI reduces security investment needs and increases automation coverage.

Answer: B

Explanation

AI introduces new attack vectors and scales threats. Options A, C, and D are incorrect.

Question

What is AI supply chain risk?

- A. Risk of low prices, as low-cost suppliers may reduce quality and increase technical debt.
- B. Risk arising from dependencies on external AI providers, datasets, and libraries, including security, availability, and continuity concerns.
- C. Lower costs, as cost reduction reduces supply chain risk and increases vendor concentration.
- D. Higher speed, as speed reduces supply chain risk and increases integration complexity.

Answer: B

Explanation

Supply chain risk involves external dependencies. Options A, C, and D are benefits or unrelated.

Question

How should third-party AI risk be managed?

- A. Ignore it, as third-party risk is outside the organisation's control and cannot be mitigated.
- B. Through due diligence, contractual protections, ongoing monitoring, and contingency planning for provider failures.
- C. Only contracts, as contractual protections are sufficient and eliminate residual risk.
- D. Only internal audits, as internal audits cover third-party risk and provide independent assurance.

Answer: B

Explanation

Third-party risk requires due diligence, contracts, monitoring, and contingency planning. Options A, C, and D are insufficient.

Question

Why is AI concentration risk a concern?

- A. Because it reduces costs, as concentration reduces procurement overhead and increases bargaining power.
- B. Because reliance on a small number of AI providers or models creates correlated failure risk across the ecosystem.
- C. Because it increases accuracy, as concentration improves model quality and increases benchmark performance.
- D. Because it improves speed, as concentration reduces integration complexity and increases deployment velocity.

Answer: B

Explanation

Concentration amplifies correlated risk. Options A, C, and D are benefits.

Question

Which is a reputational risk from AI?

- A. Faster delivery, as speed reduces reputational risk and increases customer satisfaction.
- B. Public backlash from biased, unethical, or incompetent AI use, resulting in loss of trust and brand value.
- C. Lower cost, as cost reduction reduces reputational risk and increases competitiveness.
- D. Higher accuracy, as accuracy reduces reputational risk and increases customer trust.

Answer: B

Explanation

Reputational risk stems from public perception of AI failures. Options A, C, and D are benefits.

Question

Which is a legal risk from AI?

- A. Lower cost, as cost reduction reduces legal risk and increases compliance efficiency.
- B. Liability for harm caused by AI decisions, non-compliance with regulations, and IP violations.
- C. Faster deployment, as speed reduces legal risk and increases time-to-market.
- D. Higher accuracy, as accuracy reduces legal risk and increases regulatory compliance.

Answer: B

Explanation

Legal risks include liability, non-compliance, and IP issues. Options A, C, and D are benefits.

Question

What environmental risk does AI present?

- A. None, as AI has no environmental impact and is carbon-neutral by design.
- B. Significant compute-related energy consumption and carbon emissions from training and inference of large models.
- C. Lower energy use, as AI reduces energy consumption and increases efficiency across sectors.
- D. Positive environmental impact only, as AI enables sustainability and reduces carbon footprints.

Answer: B

Explanation

AI compute consumption creates environmental risk. Options A, C, and D are incorrect.

Question

Why is regulatory fragmentation a challenge for AI governance?

- A. Because it speeds up compliance, as fragmentation reduces regulatory burden and increases harmonisation.
- B. Because different jurisdictions impose different rules, complicating compliance for globally operating firms.
- C. Because it reduces cost, as fragmentation reduces compliance overhead and increases efficiency.
- D. Because it simplifies everything, as fragmentation reduces complexity and increases clarity.

Answer: B

Explanation

Fragmented regulation complicates global compliance. Options A, C, and D are incorrect.

Question

Which is the first line's responsibility for AI models?

- A. Independent validation, as the first line validates models independently and provides objective assessment.
- B. Design, develop, own, and operate AI models; monitor performance and own data quality.
- C. Set policy, as the first line sets AI governance policy and establishes risk appetite.
- D. Audit, as the first line audits AI models and provides independent assurance.

Answer: B

Explanation

First line owns AI models and operations. Options A, C, and D belong to second or third lines.

Question

Which is the second line's responsibility for AI models?

- A. Build AI models, as the second line builds models and owns data pipelines.
- B. Independent validation, policy, and challenge of first-line assumptions and controls.
- C. Use models for decisions, as the second line uses models for business decisions and customer interactions.
- D. Market models, as the second line markets AI models and drives adoption.

Answer: B

Explanation

Second line validates, sets policy, and challenges. Options A, C, and D belong to first line.

Question

Which is the third line's responsibility for AI models?

- A. Build models, as the third line builds models and owns data pipelines.
- B. Provide independent assurance on the effectiveness of AI governance and controls.
- C. Set risk appetite, as the third line sets risk appetite and establishes tolerances.
- D. Manage model data, as the third line manages model data and owns data quality.

Answer: B

Explanation

Third line provides independent assurance. Options A, C, and D belong to first or second lines.

Question

Why must AI validation be independent?

- A. To speed up processes, as independence accelerates validation timelines and reduces bureaucratic friction.
- B. To avoid conflicts of interest and ensure objective assessment of model quality and risk.
- C. To reduce costs, as independence reduces validation costs and increases efficiency.
- D. To bypass policies, as independence exempts validation from policies and standard procedures.

Answer: B

Explanation

Independence ensures objective validation. Options A, C, and D are incorrect.

Question

What does effective challenge mean in the AI governance context?

- A. Automatic approval, as effective challenge is a formality and routine step in the approval workflow.
- B. Critical, objective, and informed review of AI models and their assumptions by qualified parties, independent from developers.
- C. Marketing endorsement, as effective challenge is a promotional activity and brand-positioning exercise.
- D. Rubber-stamping, as effective challenge is a compliance formality and documentation exercise.

Answer: B

Explanation

Effective challenge is critical review. Options A, C, and D are inconsistent.

Question

How should AI risk appetite cascade through the organisation?

- A. It shouldn't, as risk appetite is set independently at each level and is not hierarchical.
- B. From board-set appetite to business-unit tolerances and model-level thresholds, all aligned.
- C. Only at business-unit level, as board appetite is not relevant to AI and is too high-level.
- D. Only at IT level, as AI risk is an IT concern and is managed by technology teams.

Answer: B

Explanation

Appetite cascades from board to units to models. Options A, C, and D are too narrow.

Question

Who owns AI risk in an organisation?

- A. The IT department, as AI risk is an IT concern and is managed by technology teams.
- B. Business units that use AI, supported by risk, compliance, and other functions.
- C. Only the CRO, as AI risk is a risk function concern and is managed by risk teams.
- D. Only external vendors, as AI risk is a vendor concern and is managed by third parties.

Answer: B

Explanation

AI risk ownership sits with business users supported by risk functions. Options A, C, and D are incomplete.

Question

What is the purpose of AI escalation paths?

- A. To slow down decisions, as escalation introduces delays and bureaucratic friction.
- B. To ensure that AI incidents, breaches, and issues are reported and addressed at the appropriate level promptly.
- C. To reduce transparency, as escalation limits information flow and increases confidentiality.
- D. To bypass governance, as escalation circumvents controls and standard procedures.

Answer: B

Explanation

Escalation paths ensure timely, appropriate responses. Options A, C, and D are contrary.

Question

What does AI risk aggregation mean?

- A. Averaging individual risks, as aggregation is a simple mean calculation and arithmetic average.
- B. Combining individual AI risks into an enterprise-wide view, accounting for interdependencies and correlations.
- C. Ignoring enterprise risk, as aggregation focuses on individual risks and ignores portfolio effects.
- D. Removing individual risks, as aggregation eliminates individual risk and reduces total exposure.

Answer: B

Explanation

Aggregation considers interdependencies. Options A, C, and D misstate.

Question

Which is a good metric for AI governance success?

- A. Number of meetings, as meeting frequency correlates with governance effectiveness and organisational alignment.
- B. Reduction in AI incidents, improvements in validation timeliness, and increasing compliance with governance standards.
- C. Number of documents, as documentation volume correlates with governance effectiveness and process maturity.
- D. Cost of IT, as IT spend correlates with governance effectiveness and operational efficiency.

Answer: B

Explanation

Governance success metrics track incidents, timeliness, and compliance. Options A, C, and D are not meaningful metrics.

Question

Under GDPR, data subjects can request:

- A. Only access to data, as GDPR grants a single right and is limited in scope.
- B. Access, rectification, erasure, restriction, portability, and objection.
- C. Only rectification, as GDPR grants a single right and is limited in scope.
- D. Nothing, as GDPR grants no rights to data subjects and is purely a compliance burden.

Answer: B

Explanation

GDPR grants multiple rights. Options A, C, and D understate or misstate.

Question

What is the difference between a data controller and a data processor under GDPR?

- A. They are the same, as both roles are interchangeable and functionally equivalent.
- B. Controller determines the purposes and means of processing; processor processes data on behalf of the controller.
- C. Controller only stores data, as the controller's role is limited to storage and retention.
- D. Processor only collects data, as the processor's role is limited to collection and ingestion.

Answer: B

Explanation

Controller decides; processor acts on behalf. Options A, C, and D misstate.

Question

What is the DPO's role under GDPR?

- A. Marketing, as the DPO is responsible for data-driven marketing and customer acquisition.
- B. Independent oversight of data protection compliance, advising on GDPR obligations, and liaising with the supervisory authority.
- C. Sales, as the DPO is responsible for data-driven sales and revenue growth.
- D. IT support, as the DPO is responsible for IT infrastructure and technical operations.

Answer: B

Explanation

The DPO provides independent compliance oversight. Options A, C, and D are unrelated.

Question

What is a DPIA?

- A. A marketing report, as DPIAs are used for brand positioning and competitive differentiation.
- B. A Data Protection Impact Assessment required for processing likely to result in high risk to data subjects, documenting risks and mitigations.
- C. A code review, as DPIAs are technical artefacts and code-quality assessments.
- D. A financial audit, as DPIAs are financial controls and SOX compliance mechanisms.

Answer: B

Explanation

DPIA documents high-risk processing and mitigations. Options A, C, and D misstate.

Question

Which AI systems are classified as high-risk under the EU AI Act?

- A. All chatbots, as chatbots are inherently high-risk and subject to strict obligations.
- B. Those affecting safety, fundamental rights, or critical infrastructure, such as biometrics, employment, credit scoring, and law enforcement.
- C. Only medical devices, as medical devices are the sole high-risk category and are subject to strict obligations.
- D. Only toys, as toys are the sole high-risk category and are subject to strict obligations.

Answer: B

Explanation

High-risk systems affect safety, rights, or critical infrastructure. Options A, C, and D are narrower or incorrect.

Question

What transparency obligations does the EU AI Act impose?

- A. None, as the AI Act does not address transparency and is limited to safety obligations.
- B. Notifying users when interacting with AI systems, disclosing deepfakes, and providing information about AI-generated content.
- C. Only financial disclosure, as the AI Act is limited to financial transparency and investor reporting.
- D. Only marketing disclosure, as the AI Act is limited to marketing transparency and advertising standards.

Answer: B

Explanation

The AI Act imposes transparency obligations. Options A, C, and D are incorrect.

Question

What does the EU AI Act require of general-purpose AI (GPAI) models?

- A. Nothing, as GPAI models are exempt from the AI Act and are subject to no obligations.
- B. Documentation, transparency on training data, copyright compliance, and safety measures, with stricter obligations for models with systemic risk.
- C. Only marketing disclosure, as GPAI obligations are limited to marketing transparency and advertising standards.
- D. Only to register, as GPAI obligations are limited to registration and notification requirements.

Answer: B

Explanation

GPAI models have documentation, transparency, and safety obligations. Options A, C, and D understate or misstate.

Question

How is the EU AI Act being implemented over time?

- A. All at once, as the AI Act has no transitional periods and applies immediately.
- B. In phases: prohibitions first, then high-risk rules, then remaining provisions, with transitional periods for existing systems.
- C. Only in 2030, as the AI Act is delayed indefinitely and is subject to ongoing negotiation.
- D. Never, as the AI Act has been withdrawn and is no longer applicable.

Answer: B

Explanation

The AI Act phases in obligations. Options A, C, and D are inaccurate.

Question

What rights does CCPA grant California consumers?

- A. No rights, as CCPA grants no consumer rights and is purely a compliance burden.
- B. Right to know, delete, opt out of sale, correct, and non-discrimination.
- C. Only access, as CCPA grants a single right and is limited in scope.
- D. Only deletion, as CCPA grants a single right and is limited in scope.

Answer: B

Explanation

CCPA grants multiple rights. Options A, C, and D are incomplete.

Question

What is the purpose of China's Personal Information Protection Law (PIPL)?

- A. To ban personal data, as PIPL prohibits all personal data processing and is highly restrictive.
- B. To protect personal data, requiring consent, data minimisation, transparency, and specific rules for cross-border transfer.
- C. Only for government data, as PIPL applies solely to government data and public-sector processing.
- D. Only for healthcare data, as PIPL applies solely to healthcare data and medical records.

Answer: B

Explanation

PIPL protects personal data with consent, minimisation, and transfer rules. Options A, C, and D are inaccurate.

Question

What evidence should model validation produce?

- A. None, as validation is an informal process and does not require documentation.
- B. Documented assessment of conceptual soundness, data appropriateness, implementation, performance, and limitations, with findings and recommendations.
- C. Only accuracy metrics, as validation is limited to performance measurement and benchmark comparisons.
- D. Only source code, as validation is limited to code review and static analysis.

Answer: B

Explanation

Validation produces documented, comprehensive evidence. Options A, C, and D are insufficient.

Question

Why does validation include benchmarking against challenger models?

- A. To confuse stakeholders, as benchmarking introduces unnecessary complexity and obscures decision-making.
- B. To compare performance against alternatives, identifying whether the model is the best choice for its purpose.
- C. To reduce cost, as benchmarking reduces validation costs and increases efficiency.
- D. To speed up training, as benchmarking accelerates model development and hyperparameter search.

Answer: B

Explanation

Benchmarking compares models for suitability. Options A, C, and D are incorrect.

Question

What does sensitivity analysis in model validation examine?

- A. Only cost, as sensitivity analysis is limited to financial considerations and FinOps unit economics.
- B. How model outputs respond to changes in input assumptions, parameters, or data, evaluating robustness.
- C. Only speed, as sensitivity analysis is limited to latency considerations and inference-time optimisation.
- D. Only storage, as sensitivity analysis is limited to infrastructure considerations and data-lake tiering.

Answer: B

Explanation

Sensitivity analysis tests robustness to changes. Options A, C, and D are unrelated.

Question

What is the purpose of stress testing AI models?

- A. To make them faster, as stress testing optimizes performance and reduces latency.
- B. To evaluate their behaviour under extreme, adverse, or unlikely scenarios.
- C. To reduce cost, as stress testing reduces infrastructure costs and increases efficiency.
- D. To simplify them, as stress testing reduces model complexity and increases interpretability.

Answer: B

Explanation

Stress tests evaluate behaviour under extremes. Options A, C, and D are unrelated.

Question

What does backtesting involve for AI models?

- A. Forward-looking simulation only, as backtesting is a forecasting technique and predictive-modelling approach.
- B. Testing the model against historical data to evaluate how it would have performed.
- C. Marketing analysis, as backtesting is a promotional activity and brand-positioning exercise.
- D. Cost analysis, as backtesting is a financial control and SOX compliance mechanism.

Answer: B

Explanation

Backtesting evaluates historical performance. Options A, C, and D are unrelated.

Question

Outcome analysis in model validation is used to:

- A. Reduce costs, as outcome analysis is a cost optimization technique and FinOps unit-economics lever.
- B. Compare model predictions against realised outcomes to detect bias, degradation, or incorrect assumptions.
- C. Speed up training, as outcome analysis accelerates model development and hyperparameter search.
- D. Increase accuracy, as outcome analysis directly improves model accuracy and benchmark performance.

Answer: B

Explanation

Outcome analysis compares predictions to reality. Options A, C, and D are unrelated.

Question

What evidence should the use test provide?

- A. None, as the use test is an informal process and does not require documentation.
- B. Qualitative evidence that the model is used appropriately, understood by users, and aligned with business decisions.
- C. Only accuracy metrics, as the use test is limited to performance measurement and benchmark comparisons.
- D. Only cost data, as the use test is limited to financial considerations and FinOps unit economics.

Answer: B

Explanation

The use test produces qualitative evidence. Options A, C, and D are insufficient.

Question

How should model use restrictions be enforced?

- A. Through documentation only, as documentation is sufficient to enforce restrictions and ensure compliance.
- B. Through documentation, access controls, system limits, and monitoring.
- C. By ignoring them, as use restrictions are optional and impede innovation.
- D. By removing them, as use restrictions impede innovation and slow deployment.

Answer: B

Explanation

Use restrictions require documentation, controls, limits, and monitoring. Options A, C, and D are insufficient or contrary.

Question

What documentation should accompany model validation?

- A. Nothing, as validation is an informal process and does not require documentation.
- B. A validation report with methodology, findings, limitations, and recommendations, plus supporting analyses.
- C. Only the model code, as validation is limited to code review and static analysis.
- D. Only marketing material, as validation is a promotional activity and brand-positioning exercise.

Answer: B

Explanation

Validation documentation is comprehensive. Options A, C, and D are insufficient or inappropriate.

Question

Which best summarises the guiding principle for responsible AI governance?

- A. Maximise complexity, as complexity improves governance outcomes and increases robustness.
- B. Balance performance with governance, validation, monitoring, fairness, explainability, and ethical alignment, consistent with regulatory and business expectations.
- C. Deploy without validation, as validation slows innovation and increases time-to-market.
- D. Ignore risks, as risk-taking drives innovation and increases competitive advantage.

Answer: B

Explanation

Responsible AI balances performance and governance across all dimensions. Options A, C, and D contradict the principle.