

# GARP Risk & AI — Q551–Q650

Unofficial Exam Prep

Unofficial

October 5, 2026

**Question**

Which is a common category in AI risk taxonomies?

- A. Technical risk, encompassing adversarial robustness, distributional shift, data-provenance integrity, and reproducibility variance.
- B. Ethical risk, encompassing allocative harm, representational harm, transparency deficits, and human-oversight sufficiency.
- C. Operational risk, encompassing third-party concentration, model-drift telemetry, process-failure modes, and resource-exhaustion scenarios.
- D. All of the above are common categories in AI risk taxonomies, alongside legal, regulatory, and reputational dimensions.

**Answer:** D

**Explanation**

AI risk taxonomies span technical, ethical, operational, and other categories; all are common.

**Question**

Which is the most relevant model risk for an AI-based credit scoring system?

- A. The colour scheme and layout of the user interface may not align with brand guidelines, leading to inconsistent customer experience and reduced engagement metrics.
- B. Potential unfairness, inaccuracy, and adverse regulatory consequences from biased or poorly performing predictions, including proxy discrimination, explainability gaps, and ECOA/FCRA non-compliance.
- C. The storage cost of retaining historical training data may exceed budgeted cloud infrastructure allocations over multi-year horizons, impacting FinOps unit economics.
- D. Marketing effectiveness may decline if the model's output is not integrated with campaign management platforms, customer segmentation tools, and CRM workflows.

**Answer: B**

**Explanation**

Credit scoring model risk includes fairness, accuracy, and regulatory risk. Options A, C, and D are secondary or irrelevant.

**Question**

For an AI fraud detection model, which is the most critical model risk?

- A. Overfitting to historical fraud patterns, missing novel fraud schemes, and failing to adapt to adversarial concept drift under non-stationary distributions.
- B. Storage cost of retaining transaction-level features and model artefacts across multiple retraining cycles, impacting data-lake tiering strategies.
- C. User interface design that may not clearly present fraud alerts to analysts, slowing investigation workflows and increasing mean-time-to-resolution.
- D. Training time may exceed operational windows, delaying model refresh cycles and reducing responsiveness to emerging fraud typologies.

**Answer: A**

**Explanation**

Overfitting to past patterns can miss novel fraud. Options B, C, and D are less critical.

## Question

Which is a typical risk for an AI-enhanced market risk model?

- A. Training on non-stationary data leading to unstable predictions and regime-change failures under structural breaks.
- B. Data storage costs may escalate if tick-level market data is retained for extended backtesting periods, impacting data-lake economics.
- C. Logo design may not comply with brand guidelines, creating inconsistent visual identity across risk reports and dashboards.
- D. Employee satisfaction may decline if model outputs are not integrated with existing workflow tools and trading-desk platforms.

Answer: A

## Explanation

Non-stationarity of financial markets is a key challenge. Options B, C, and D are unrelated.

## Question

Which is a recognised fairness metric used in lending AI?

- A. Demographic parity, measured as the ratio of positive prediction rates between protected and non-protected groups.
- B. Storage cost parity, measured as the ratio of data storage costs between protected and non-protected groups.
- C. Latency parity, measured as the ratio of inference times between protected and non-protected groups.
- D. Colour parity, measured as the ratio of UI colour contrast ratios between protected and non-protected groups.

Answer: A

## Explanation

Demographic parity is a standard fairness metric. Options B, C, and D are not.

**Question**

Which is a key risk of algorithmic trading models?

- A. Flash-crash amplification, unintended feedback loops, and market impact from correlated strategies under stressed liquidity conditions.
- B. Storage capacity may be insufficient to retain tick-level data for regulatory reporting and backtesting, impacting MiFID II compliance.
- C. Logo design may not comply with exchange branding requirements, creating regulatory filing issues and venue-certification delays.
- D. Employee benefits may not be competitive, leading to talent attrition in quantitative teams and loss of institutional knowledge.

**Answer: A**

**Explanation**

Algorithmic trading risks include flash crashes and correlated behaviour. Options B, C, and D are unrelated.

**Question**

Which is the most significant risk from deploying a GenAI chatbot in banking?

- A. Generating incorrect or non-compliant advice to customers, creating legal and reputational risk under consumer-protection regimes.
- B. Faster response times may reduce customer wait times but could increase volume of unverified interactions and hallucination exposure.
- C. Cost savings may improve margins but could reduce investment in human oversight and quality assurance.
- D. Customer satisfaction may increase but could mask underlying model deficiencies and bias.

**Answer:** A

**Explanation**

Advice quality and compliance are central concerns. Options B, C, and D are benefits.

**Question**

Which is a key governance concern when using alternative data in financial AI?

- A. Legal right to use the data, privacy, and potential for discriminatory impact through proxy variables and disparate-impact mechanisms.
- B. Storage cost may exceed budgeted cloud infrastructure allocations over multi-year horizons, impacting FinOps unit economics.
- C. Model speed may decline if alternative data requires extensive preprocessing and feature engineering pipelines.
- D. Logo design may not comply with brand guidelines, creating inconsistent visual identity across risk reports.

**Answer: A**

**Explanation**

Legal rights, privacy, and fairness are key governance concerns. Options B, C, and D are secondary.

## Question

What is A/B testing in the AI deployment context?

- A. Randomly assigning users to two model versions and comparing outcomes under controlled experimental conditions.
- B. Retiring both models, as A/B testing is a decommissioning technique and archival strategy.
- C. Doubling training data, as A/B testing is a data augmentation technique and synthetic-data generation method.
- D. Storing data twice, as A/B testing is a storage redundancy technique and disaster-recovery strategy.

Answer: A

## Explanation

A/B testing compares model versions via controlled experiments. Options B, C, and D are unrelated.

## Question

An AI model begins producing biased outputs. Which is the most likely root cause?

- A. Data drift or newly biased data, as the model's input distribution has shifted under non-stationary conditions.
- B. Server failure, as hardware faults can introduce bias into model outputs through corrupted tensor operations.
- C. Power outage, as interrupted training can introduce bias into model weights through checkpoint corruption.
- D. UI redesign, as changes to the user interface can introduce bias into model inputs through altered data-capture flows.

Answer: A

## Explanation

Bias in outputs typically stems from data or model development. Options B, C, and D are unrelated.

**Question**

What is red teaming in AI?

- A. A marketing exercise, as red teaming is used for brand positioning and competitive differentiation.
- B. Adversarial testing by a dedicated team to uncover vulnerabilities and unintended behaviours before deployment.
- C. A financial audit, as red teaming is used for financial controls testing and SOX compliance.
- D. A user survey, as red teaming is used for customer feedback collection and satisfaction measurement.

**Answer: B**

**Explanation**

Red teaming proactively tests for vulnerabilities. Options A, C, and D misstate.

**Question**

Which is a common gap in AI governance frameworks?

- A. Extensive controls, as frameworks often over-control low-risk use cases and create unnecessary friction.
- B. Missing governance for GenAI, alternative data, or third-party models.
- C. Complete coverage, as frameworks often over-cover all risk dimensions and create unnecessary complexity.
- D. Perfect alignment, as frameworks often over-align with business objectives and create unnecessary constraints.

**Answer: B**

**Explanation**

Common gaps include GenAI, alternative data, and vendor models. Options A, C, and D are not gaps.

**Question**

What is the purpose of an AI governance roadmap?

- A. A marketing document, as roadmaps are used for brand positioning and competitive differentiation.
- B. A phased plan outlining how the organisation will build, enhance, and sustain AI governance capabilities over time.
- C. A single policy, as roadmaps are policy documents and compliance artefacts.
- D. A code repository, as roadmaps are technical artefacts and infrastructure-as-code modules.

**Answer: B**

**Explanation**

A roadmap is a phased plan. Options A, C, and D misstate.

**Question**

Which is an appropriate KPI for AI governance?

- A. Number of employees, as headcount correlates with governance effectiveness and organisational capacity.
- B. Percentage of models validated on schedule.
- C. Office size, as physical footprint correlates with governance effectiveness and operational scale.
- D. Marketing spend, as brand awareness correlates with governance effectiveness and market positioning.

**Answer: B**

**Explanation**

Validation timeliness is a governance KPI. Options A, C, and D are unrelated.

**Question**

How can GenAI undermine human autonomy?

- A. By providing choices, as GenAI expands the option set available to users and increases decision-space dimensionality.
- B. By generating persuasive, tailored content that manipulates users without their awareness.
- C. By explaining decisions, as GenAI provides transparency into decision processes and interpretability artefacts.
- D. By slowing systems, as GenAI introduces latency that reduces user control and increases interaction friction.

**Answer: B**

**Explanation**

Manipulative content undermines autonomy. Options A, C, and D respect or are unrelated to autonomy.

**Question**

How does GenAI amplify bias?

- A. By generating content that reflects biases present in training data, at scale.
- B. By reducing accuracy, as lower accuracy correlates with higher bias and increased error rates.
- C. By slowing training, as longer training correlates with higher bias and overfitting to minority classes.
- D. By increasing costs, as higher costs correlate with higher bias and increased computational complexity.

**Answer: A**

**Explanation**

GenAI can amplify training-data bias at scale. Options B, C, and D are unrelated.

**Question**

What consent issues arise with GenAI training data?

- A. None, as GenAI training data is always consented and properly licensed.
- B. Data subjects whose content was used in training may not have consented, raising ethical and legal concerns.
- C. Consent is automatically granted, as GenAI training data is publicly available and web-scraped.
- D. Consent is not required, as GenAI training data is exempt from consent requirements under fair-use doctrines.

**Answer: B**

**Explanation**

Training data consent is a critical concern. Options A, C, and D are incorrect.

**Question**

Why is attribution important in GenAI?

- A. It isn't important, as GenAI outputs are not attributable and are generated de novo.
- B. To trace sources of content and information, supporting transparency and possibly legal rights such as copyright.
- C. To increase cost, as attribution is primarily a cost optimization technique and FinOps unit-economics lever.
- D. To slow inference, as attribution is primarily a latency optimization technique and inference-time acceleration method.

**Answer: B**

**Explanation**

Attribution supports transparency and copyright compliance. Options A, C, and D are incorrect.

**Question**

Which is a safety concern with GenAI in critical applications?

- A. Faster response, as speed reduces safety in critical applications and increases error propagation.
- B. Generating unsafe or harmful content that could cause physical, psychological, or financial harm.
- C. Lower costs, as cost reduction reduces safety in critical applications and increases technical debt.
- D. Higher accuracy, as accuracy reduces safety in critical applications and increases overconfidence.

**Answer: B**

**Explanation**

Harmful content generation is a safety concern. Options A, C, and D are benefits.

**Question**

Which consumer protection issue arises from GenAI?

- A. Increasing prices only, as GenAI directly causes price increases and inflationary pressure.
- B. Misleading or inaccurate information provided to consumers, potentially leading to harm.
- C. Faster service, as speed reduces consumer protection and increases error propagation.
- D. Lower fees, as fee reduction reduces consumer protection and increases hidden costs.

**Answer:** B

**Explanation**

Misleading information is a consumer protection issue. Options A, C, and D are unrelated.

**Question**

How can GenAI facilitate fraud?

- A. By making it easier to create convincing phishing content and deepfakes at scale.
- B. By improving accuracy, as accuracy reduces fraud and increases detection rates.
- C. By reducing costs, as cost reduction reduces fraud and increases detection efficiency.
- D. By speeding up transactions, as speed reduces fraud and increases detection velocity.

**Answer: A**

**Explanation**

GenAI enables scalable, convincing fraud content. Options B, C, and D are unrelated.

**Question**

Why is AI concentration risk a concern?

- A. Because it reduces costs, as concentration reduces procurement overhead and increases bargaining power.
- B. Because reliance on a small number of AI providers or models creates correlated failure risk across the ecosystem.
- C. Because it increases accuracy, as concentration improves model quality and increases benchmark performance.
- D. Because it improves speed, as concentration reduces integration complexity and increases deployment velocity.

**Answer: B**

**Explanation**

Concentration amplifies correlated risk. Options A, C, and D are benefits.

**Question**

How should AI risk appetite cascade through the organisation?

- A. It shouldn't, as risk appetite is set independently at each level and is not hierarchical.
- B. From board-set appetite to business-unit tolerances and model-level thresholds, all aligned.
- C. Only at business-unit level, as board appetite is not relevant to AI and is too high-level.
- D. Only at IT level, as AI risk is an IT concern and is managed by technology teams.

**Answer: B**

**Explanation**

Appetite cascades from board to units to models. Options A, C, and D are too narrow.

**Question**

Who owns AI risk in an organisation?

- A. The IT department, as AI risk is an IT concern and is managed by technology teams.
- B. Business units that use AI, supported by risk, compliance, and other functions.
- C. Only the CRO, as AI risk is a risk function concern and is managed by risk teams.
- D. Only external vendors, as AI risk is a vendor concern and is managed by third parties.

**Answer: B**

**Explanation**

AI risk ownership sits with business users supported by risk functions. Options A, C, and D are incomplete.

**Question**

What is the purpose of AI escalation paths?

- A. To slow down decisions, as escalation introduces delays and bureaucratic friction.
- B. To ensure that AI incidents, breaches, and issues are reported and addressed at the appropriate level promptly.
- C. To reduce transparency, as escalation limits information flow and increases confidentiality.
- D. To bypass governance, as escalation circumvents controls and standard procedures.

**Answer:** B

**Explanation**

Escalation paths ensure timely, appropriate responses. Options A, C, and D are contrary.

**Question**

What does AI risk aggregation mean?

- A. Averaging individual risks, as aggregation is a simple mean calculation and arithmetic average.
- B. Combining individual AI risks into an enterprise-wide view, accounting for interdependencies and correlations.
- C. Ignoring enterprise risk, as aggregation focuses on individual risks and ignores portfolio effects.
- D. Removing individual risks, as aggregation eliminates individual risk and reduces total exposure.

**Answer: B**

**Explanation**

Aggregation considers interdependencies. Options A, C, and D misstate.

**Question**

Which is a good metric for AI governance success?

- A. Number of meetings, as meeting frequency correlates with governance effectiveness and organisational alignment.
- B. Reduction in AI incidents, improvements in validation timeliness, and increasing compliance with governance standards.
- C. Number of documents, as documentation volume correlates with governance effectiveness and process maturity.
- D. Cost of IT, as IT spend correlates with governance effectiveness and operational efficiency.

**Answer: B**

**Explanation**

Governance success metrics track incidents, timeliness, and compliance. Options A, C, and D are not meaningful metrics.